

# A SECURE PUBLIC CACHE FOR YARN APPLICATION RESOURCES

CHRIS TREZZO  
@CTREZZO




# AGENDA

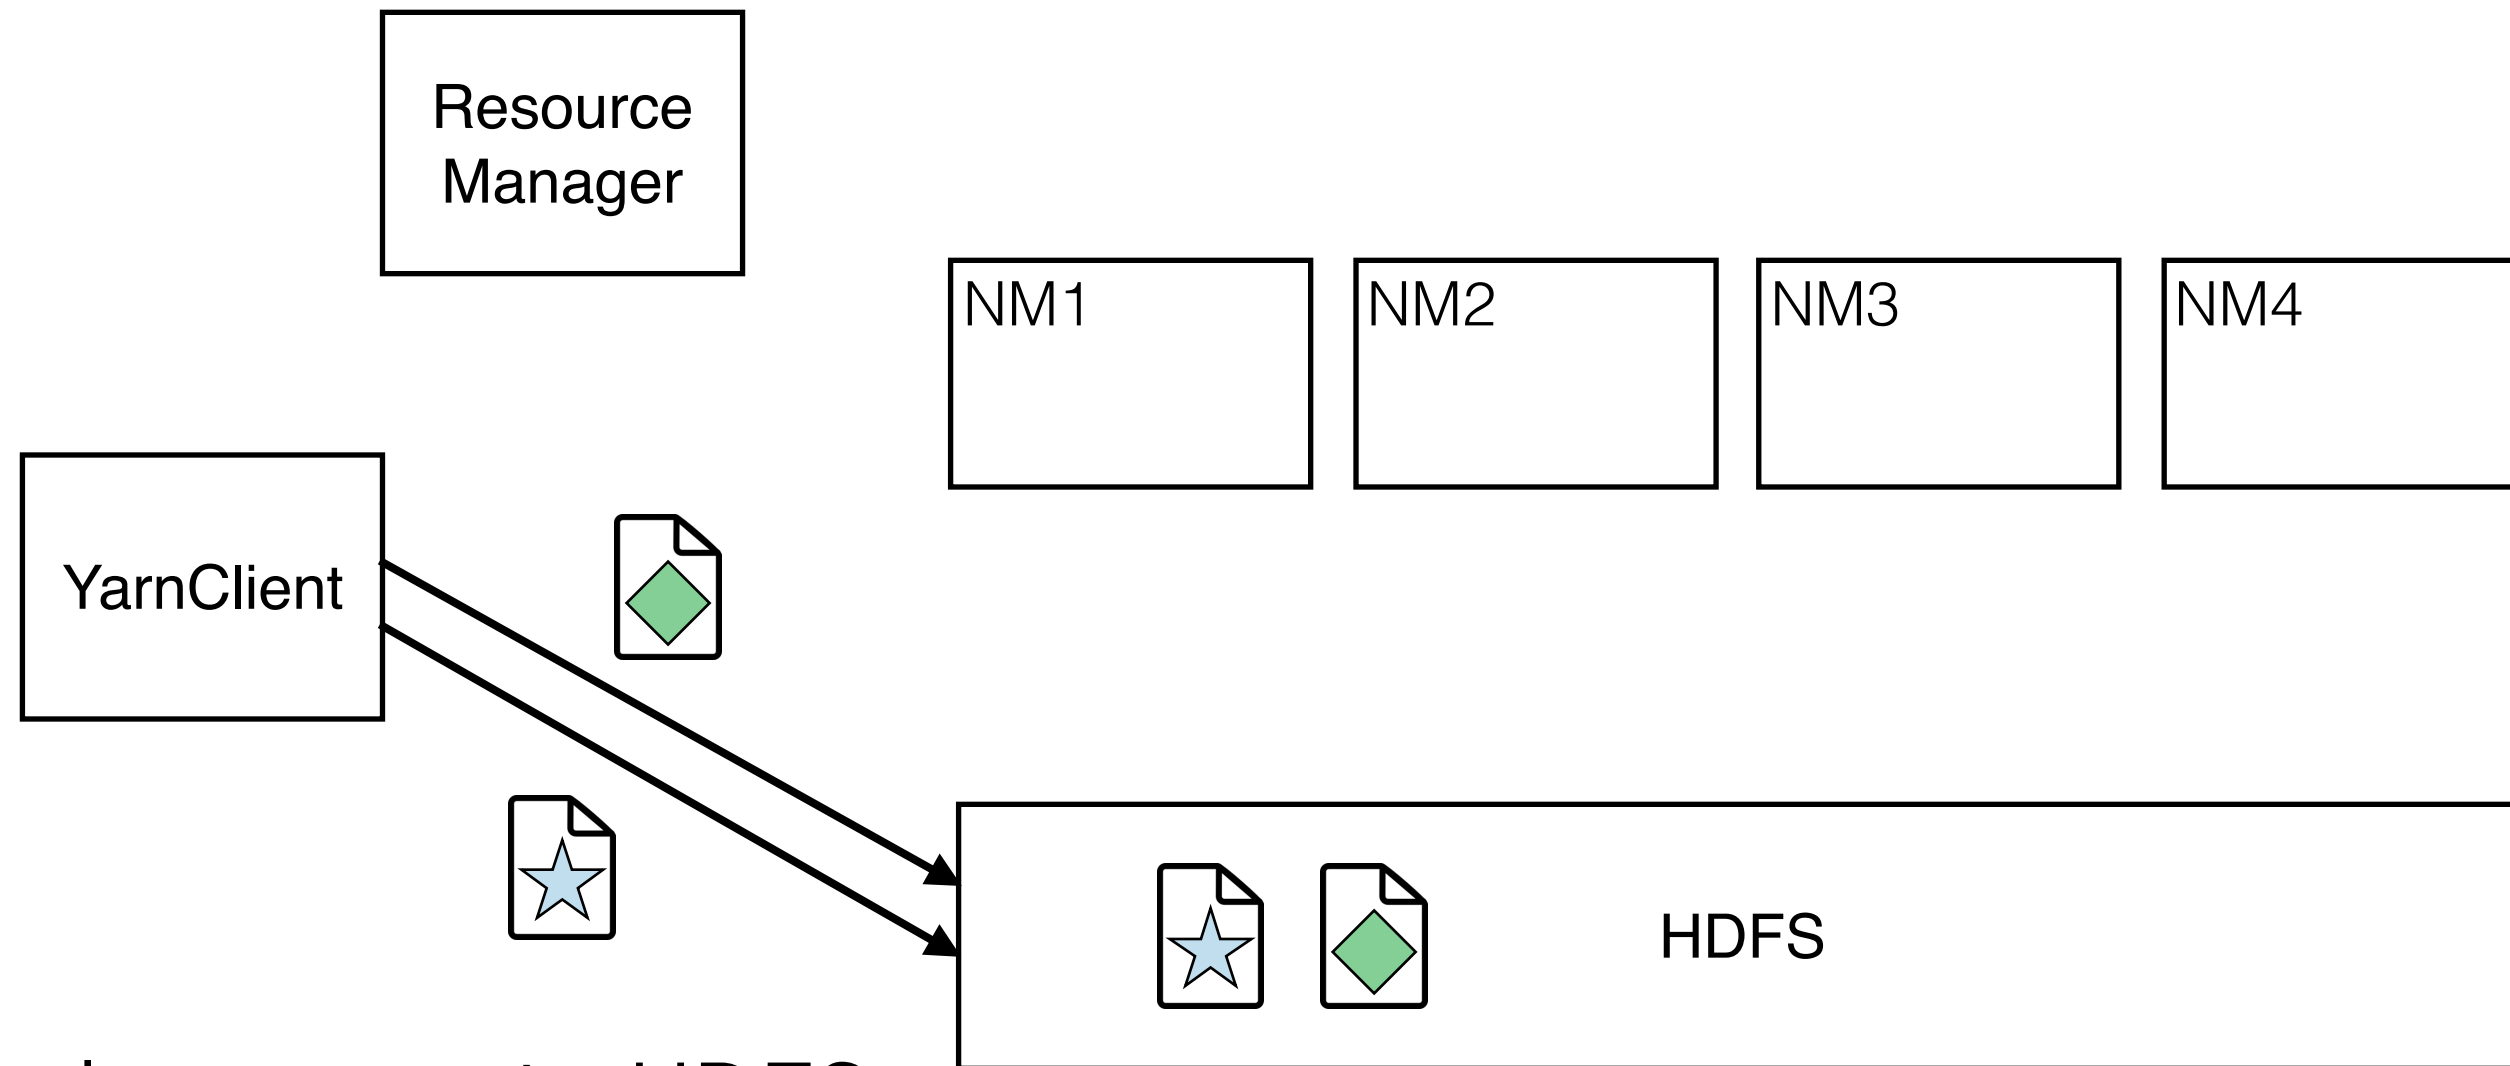
- YARN Localization
- Shared cache design overview
- Adding resources to the cache
- Does it work?
- Anti-patterns that cause churn
- Admin Tips
- Dev Tips



# LOCALIZATION OVERVIEW




# LOCALIZATION OVERVIEW



1. Upload resource to HDFS




# LOCALIZATION OVERVIEW



2. Submit application




# LOCALIZATION OVERVIEW



3. Schedule containers



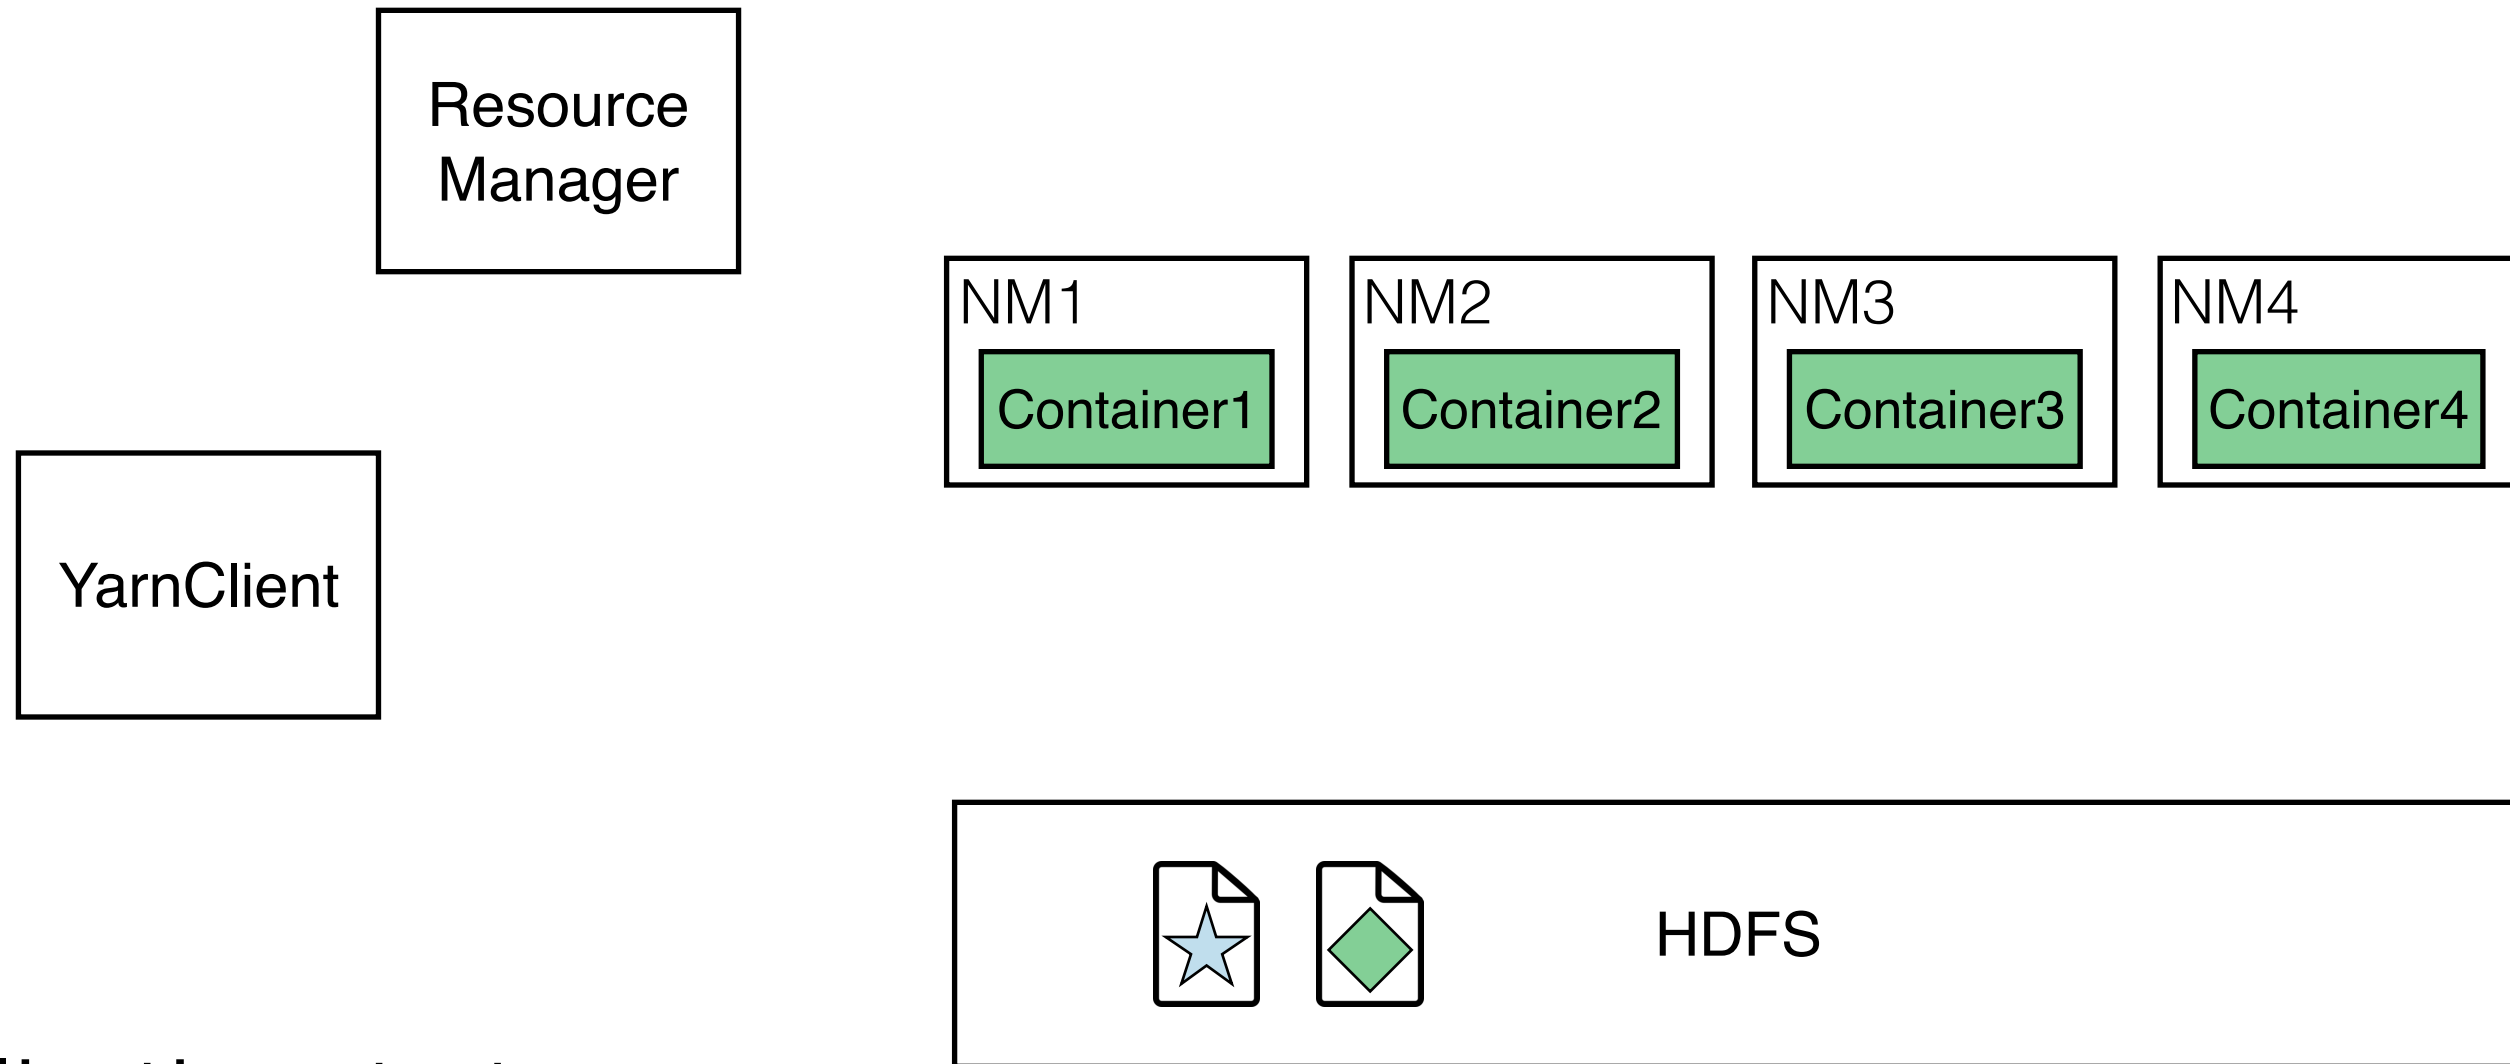
# LOCALIZATION OVERVIEW



4. Copy resources to local disk



# LOCALIZATION OVERVIEW

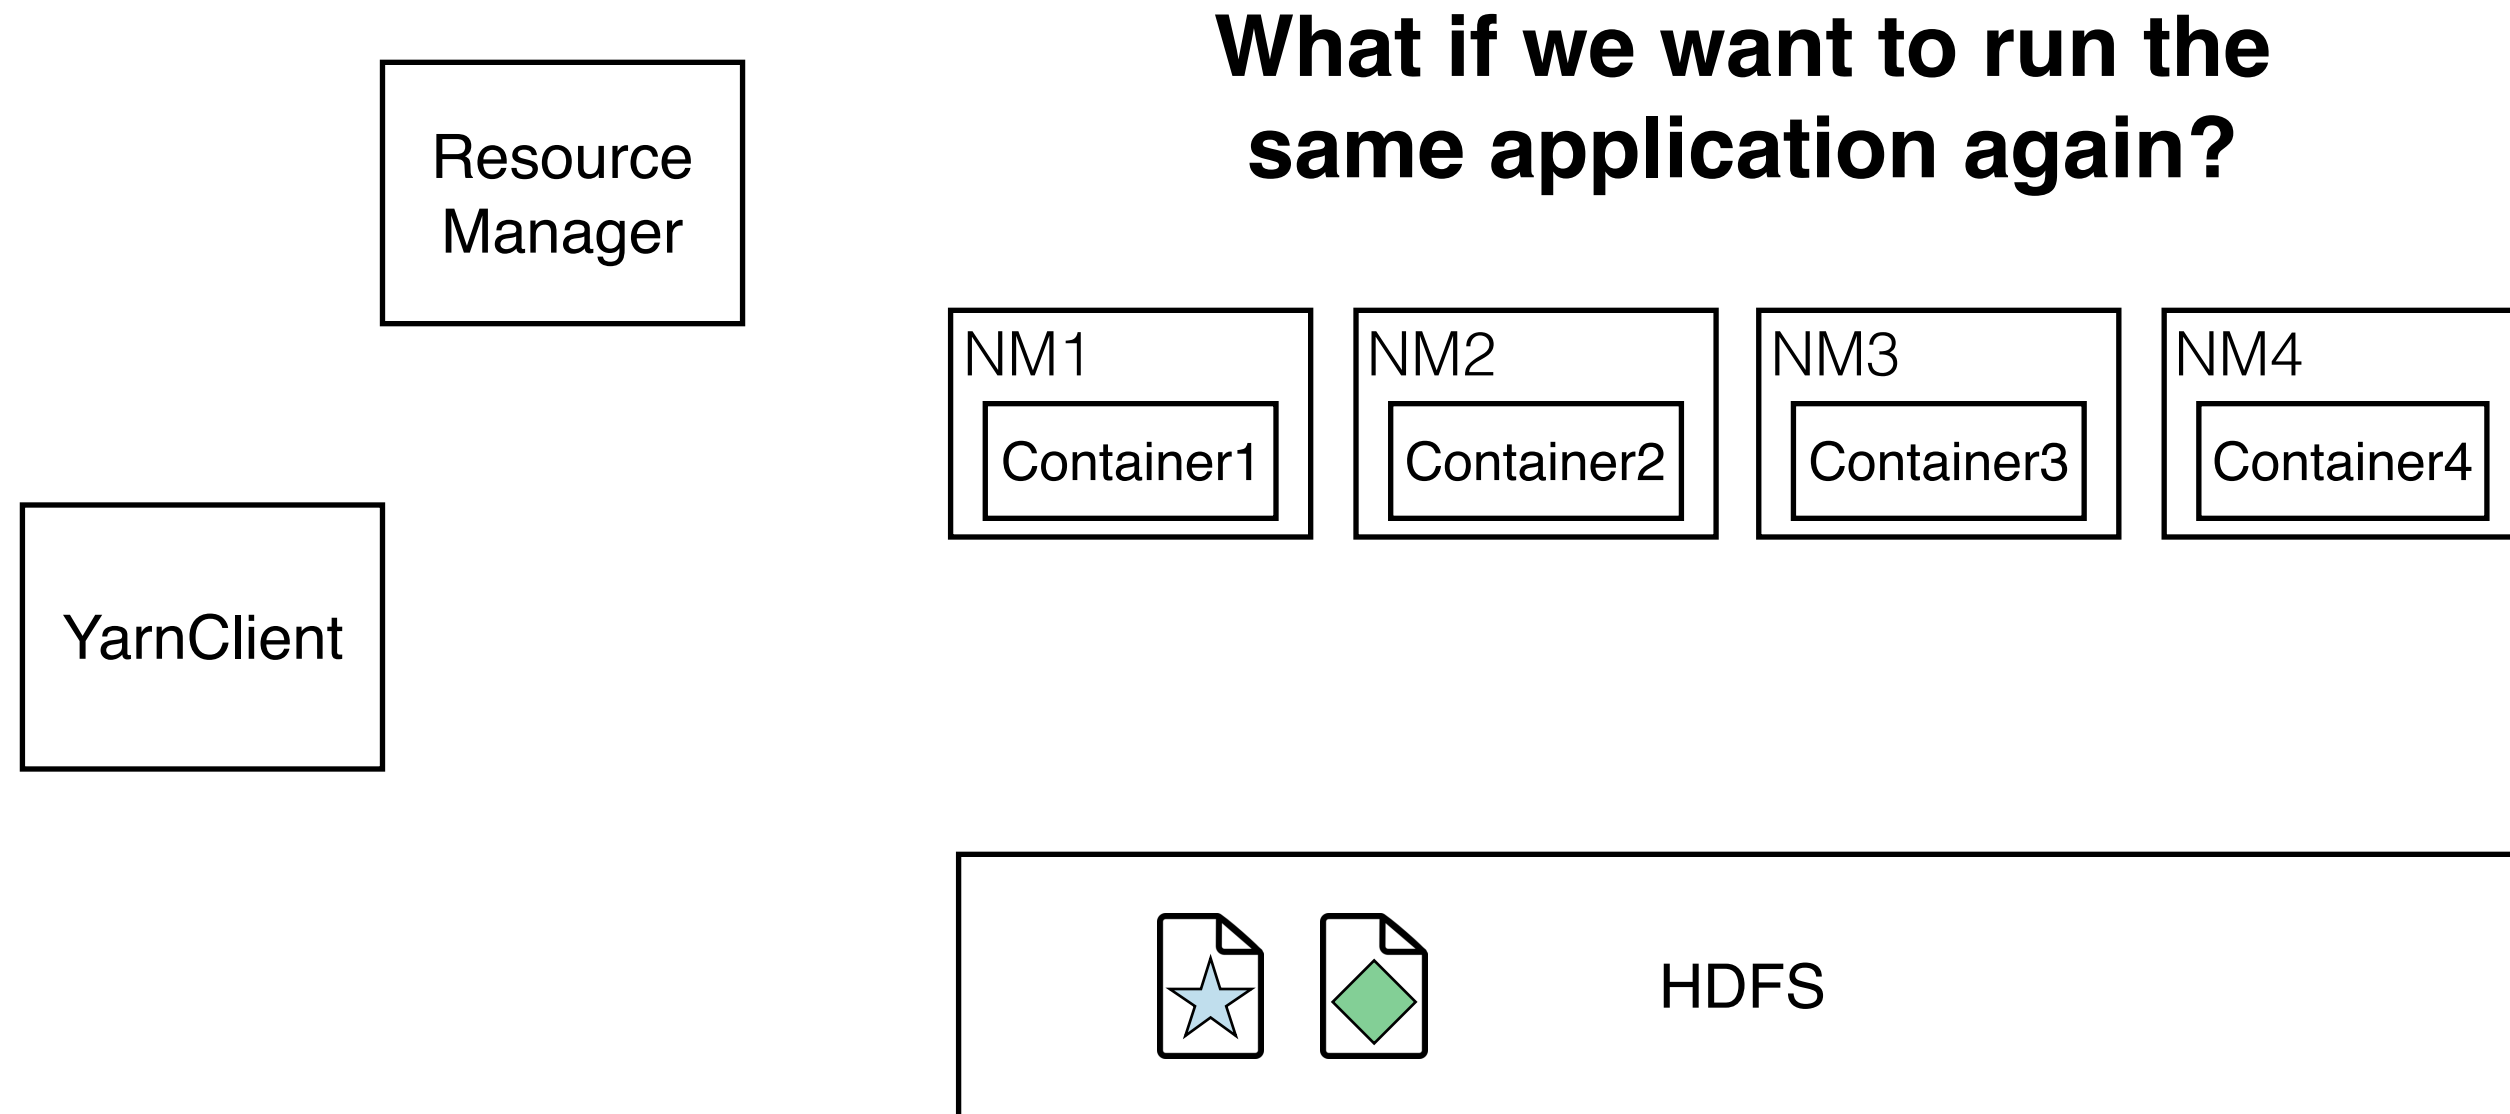


5. Application starts running

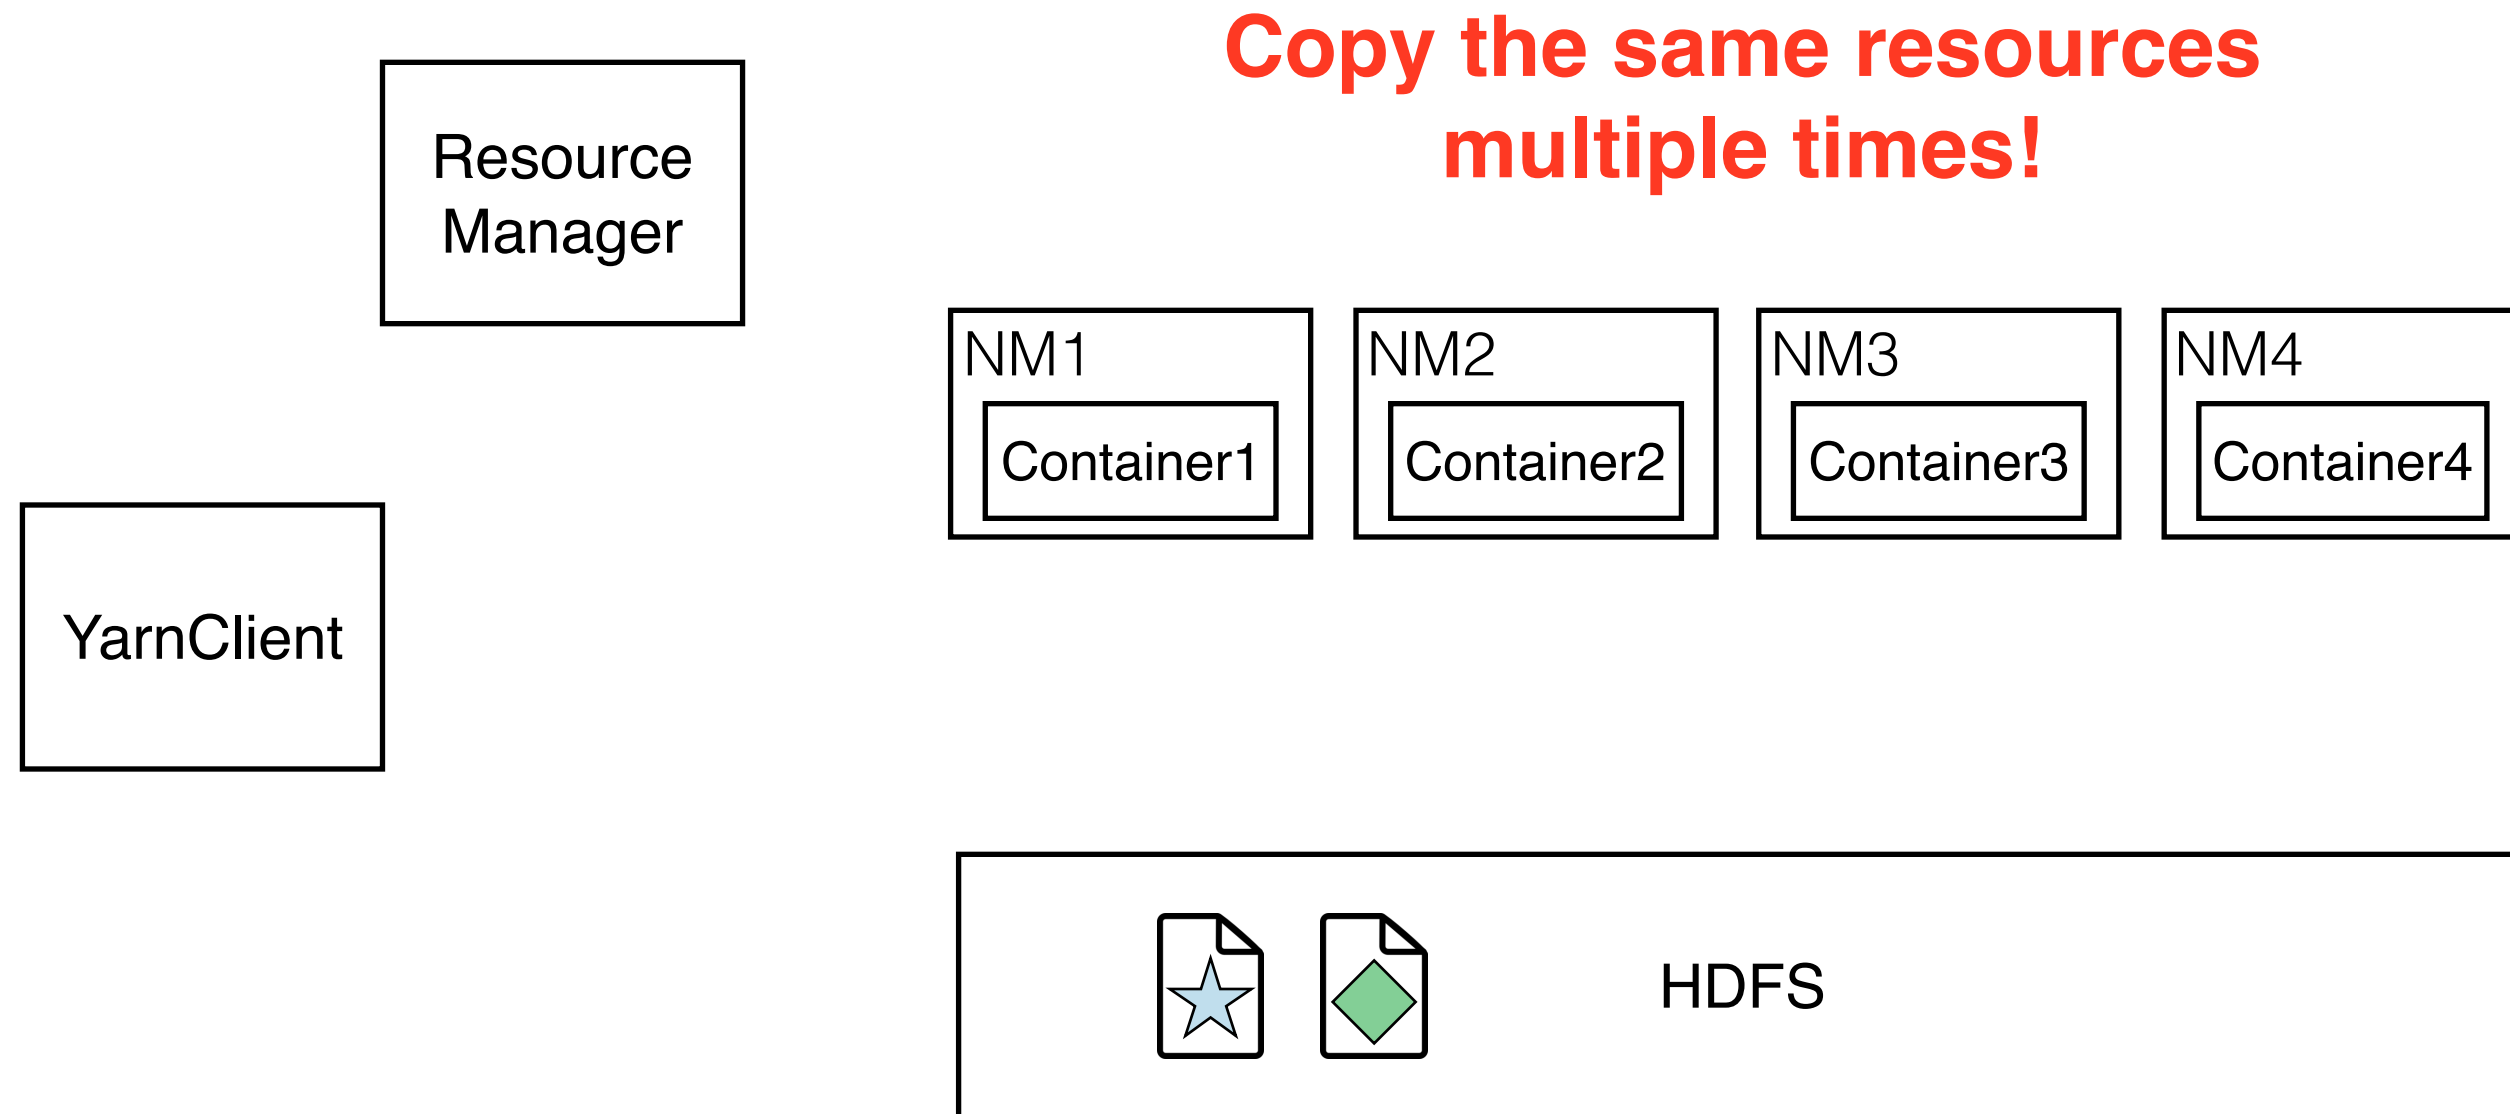




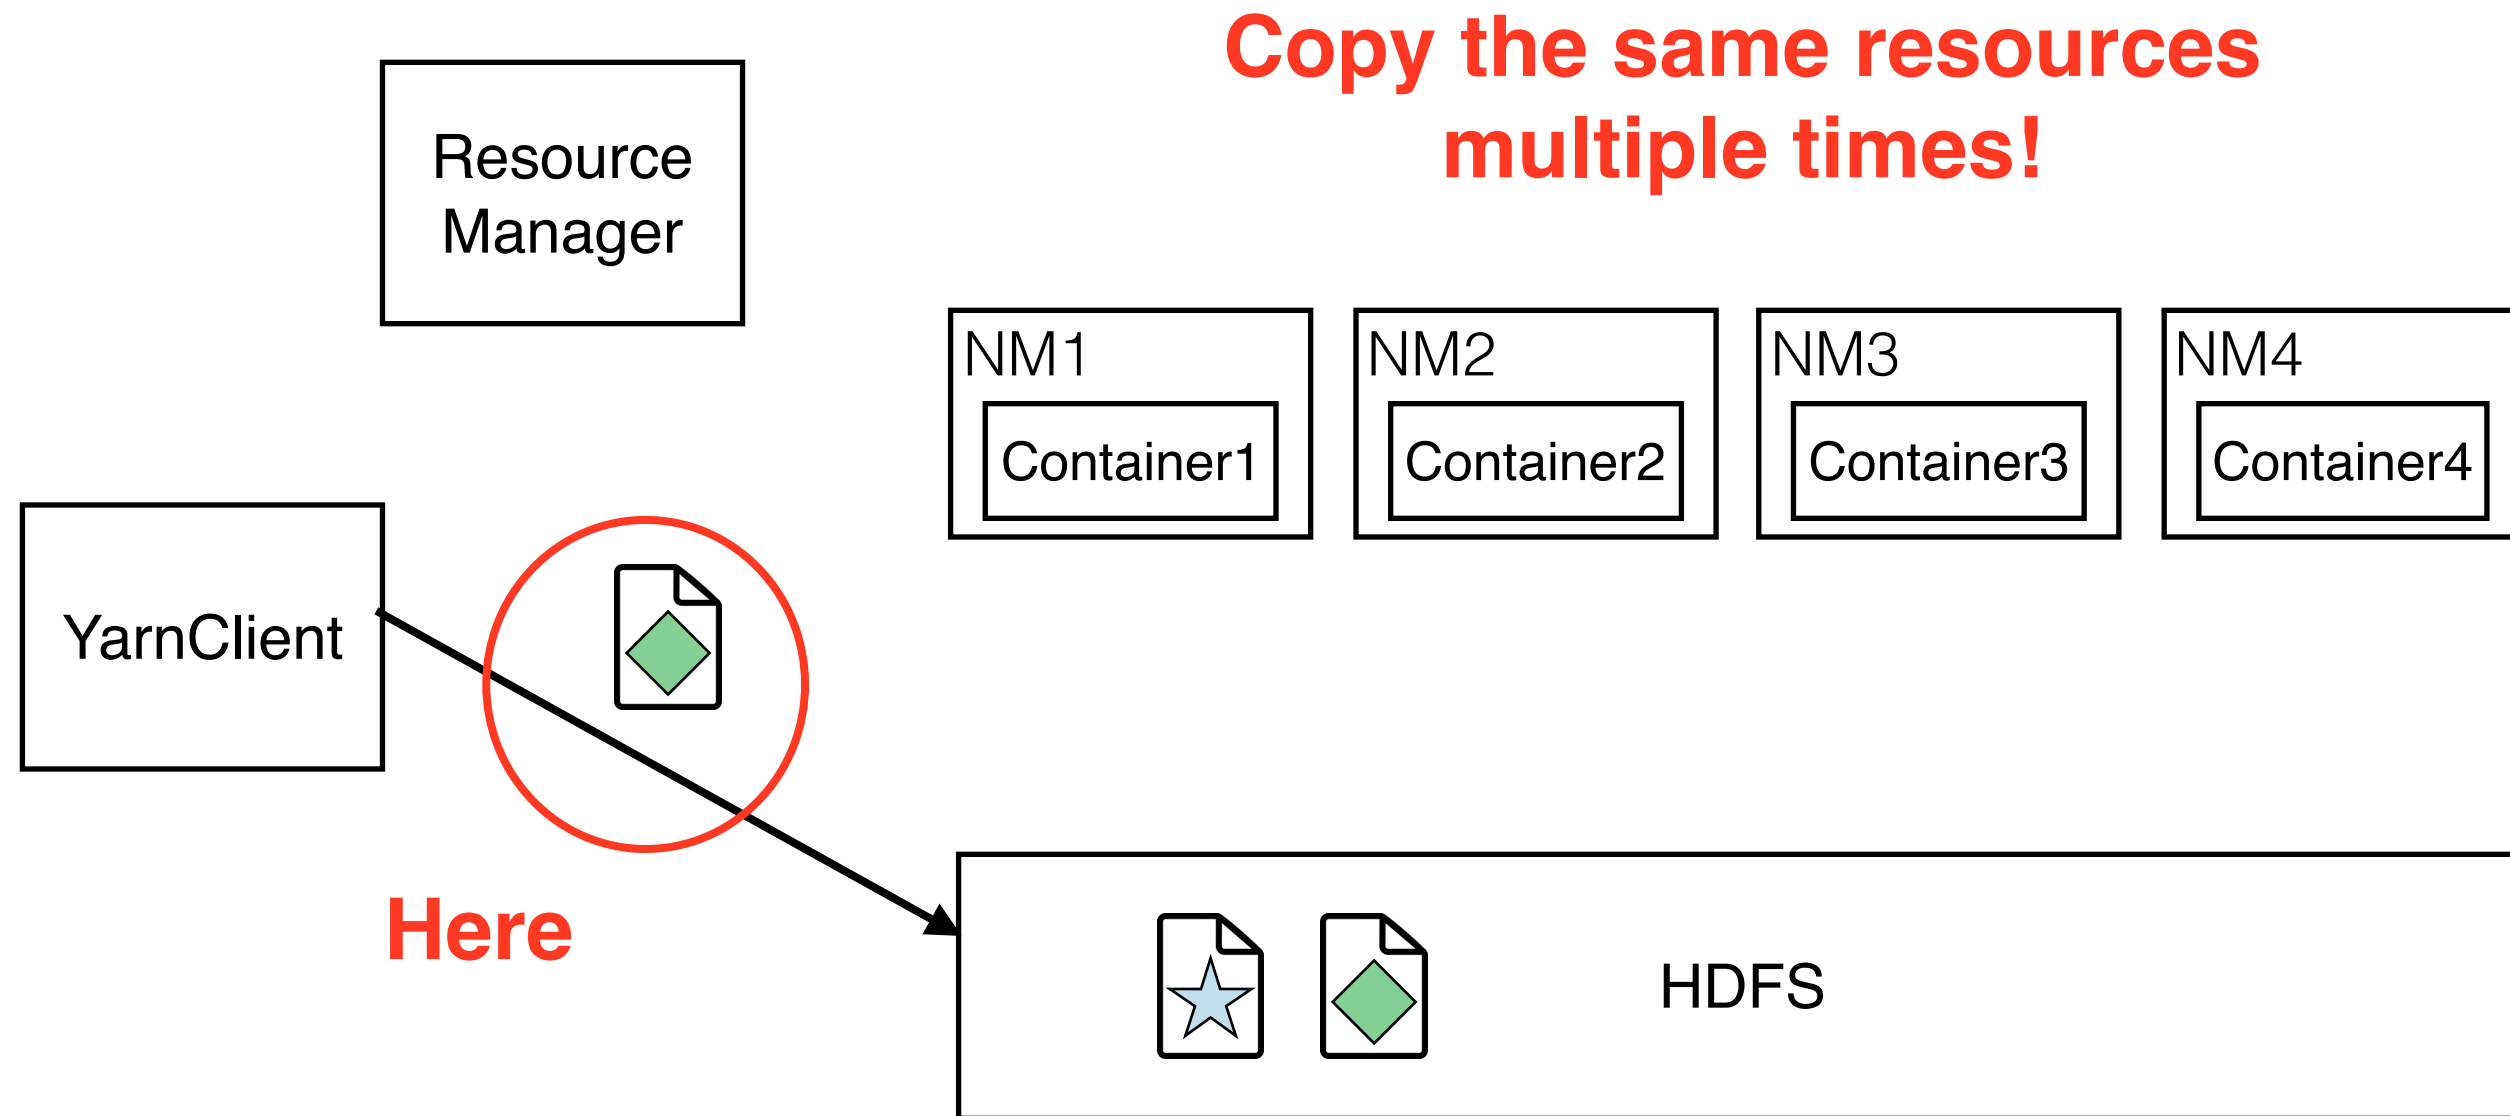
# LOCALIZATION OVERVIEW



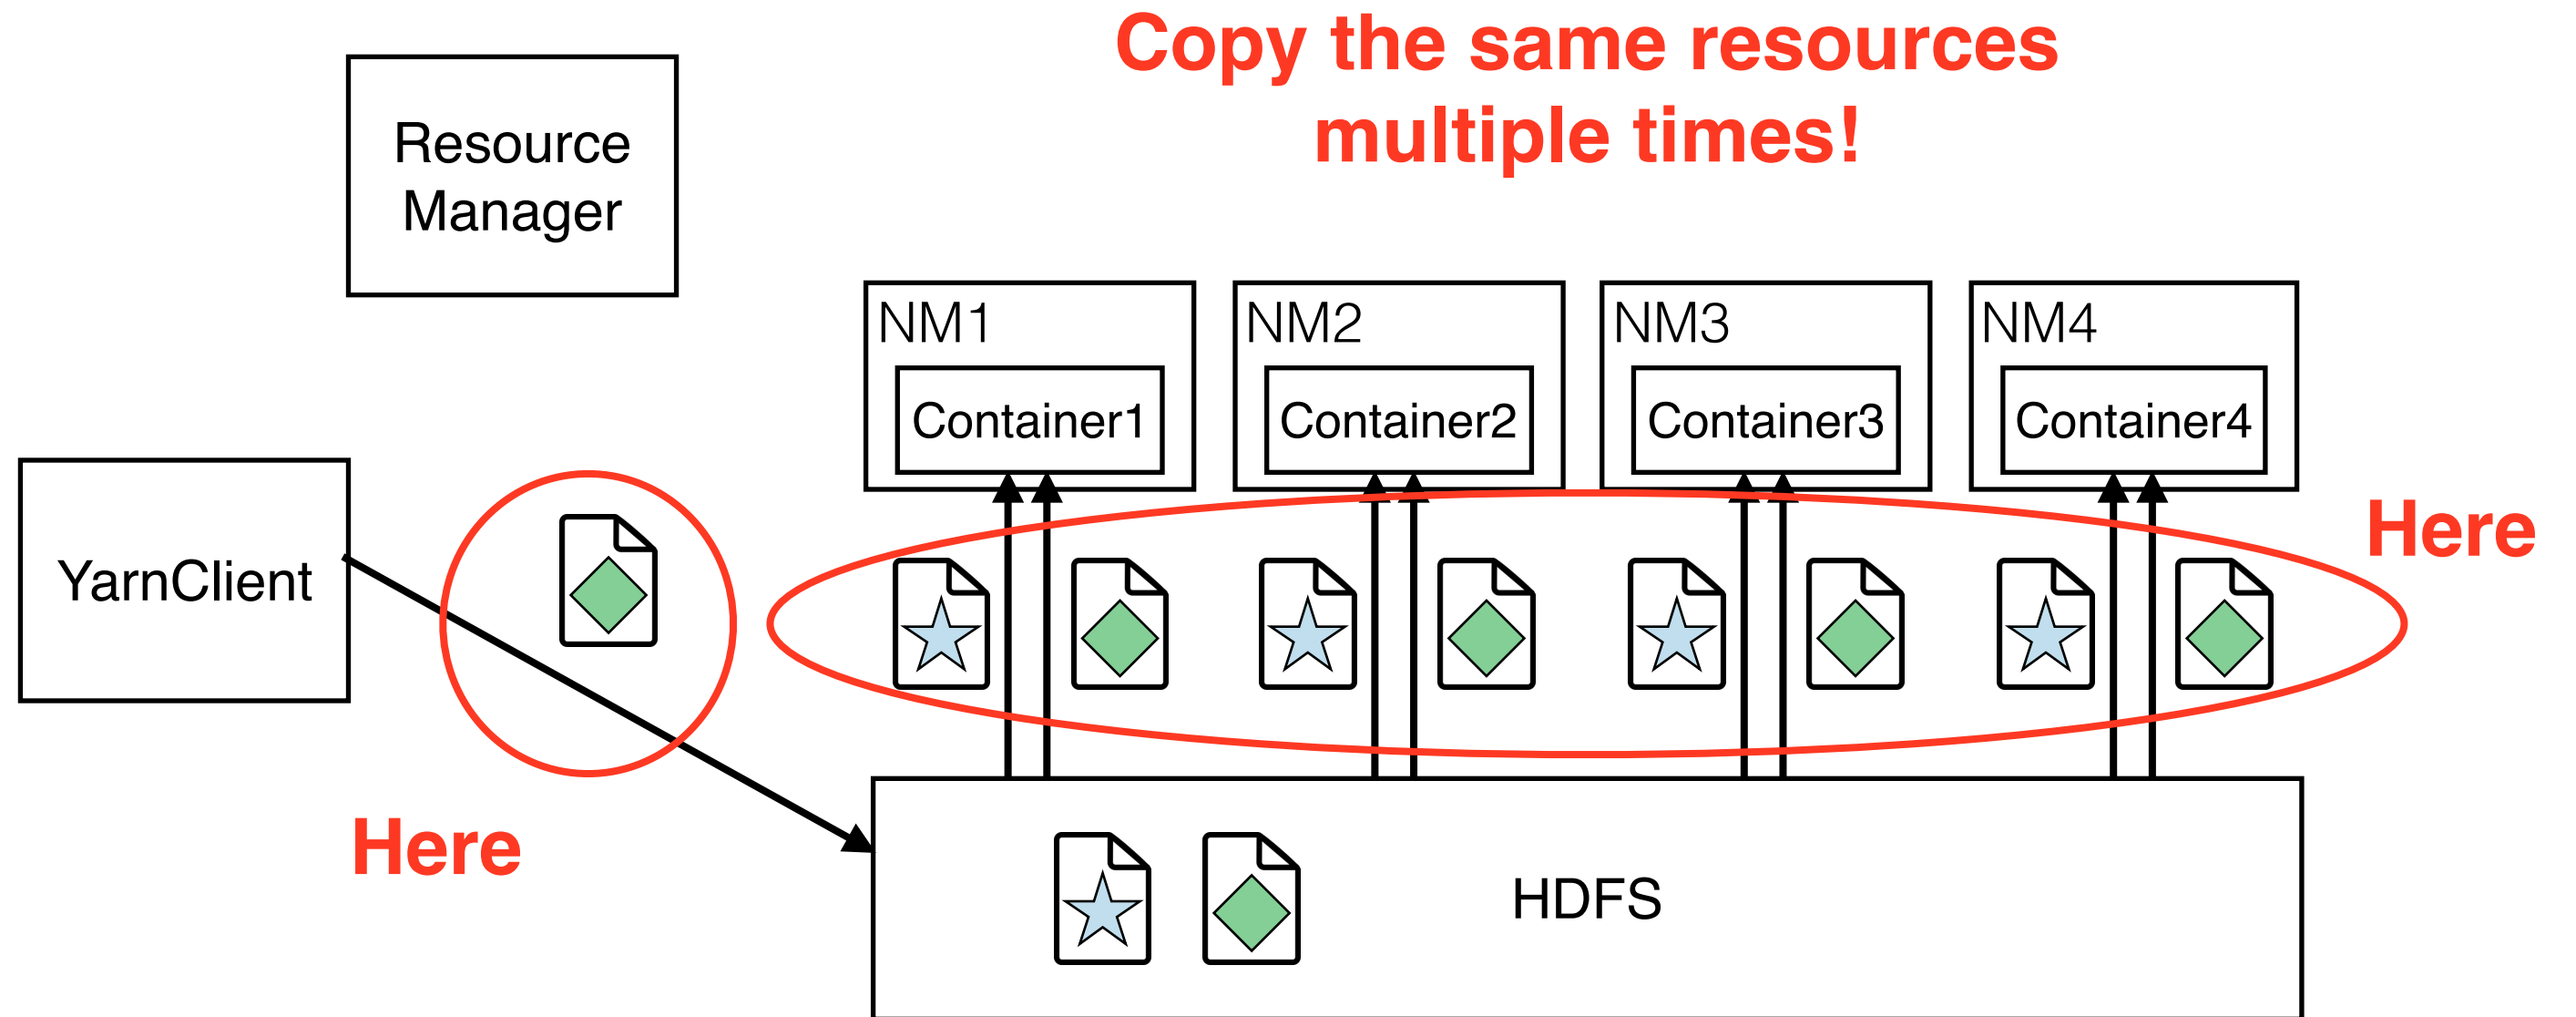
# LOCALIZATION OVERVIEW



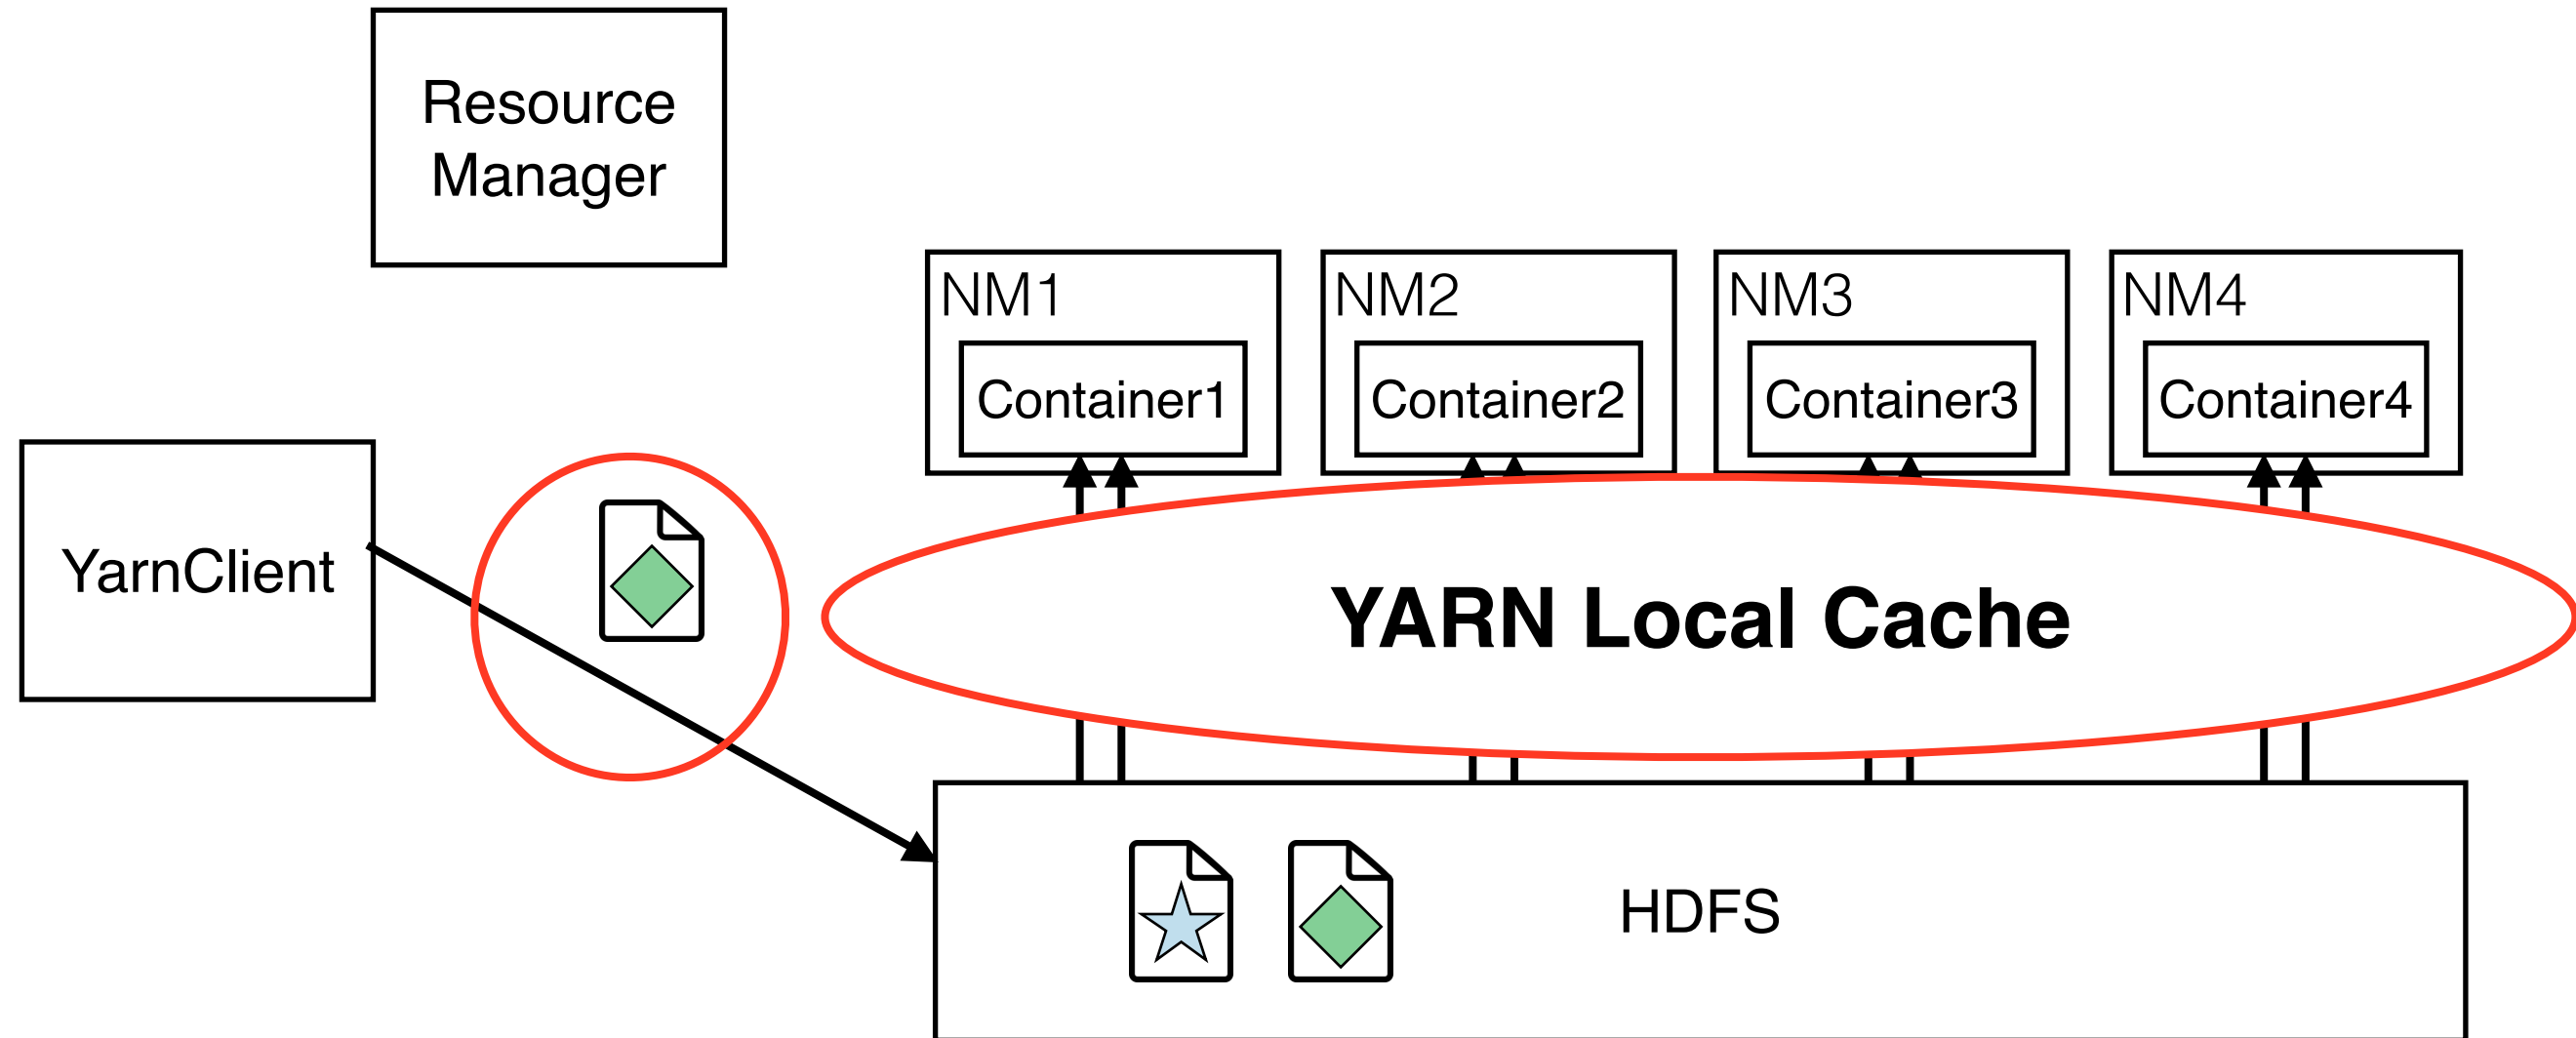
# LOCALIZATION OVERVIEW



# LOCALIZATION OVERVIEW



# LOCALIZATION OVERVIEW

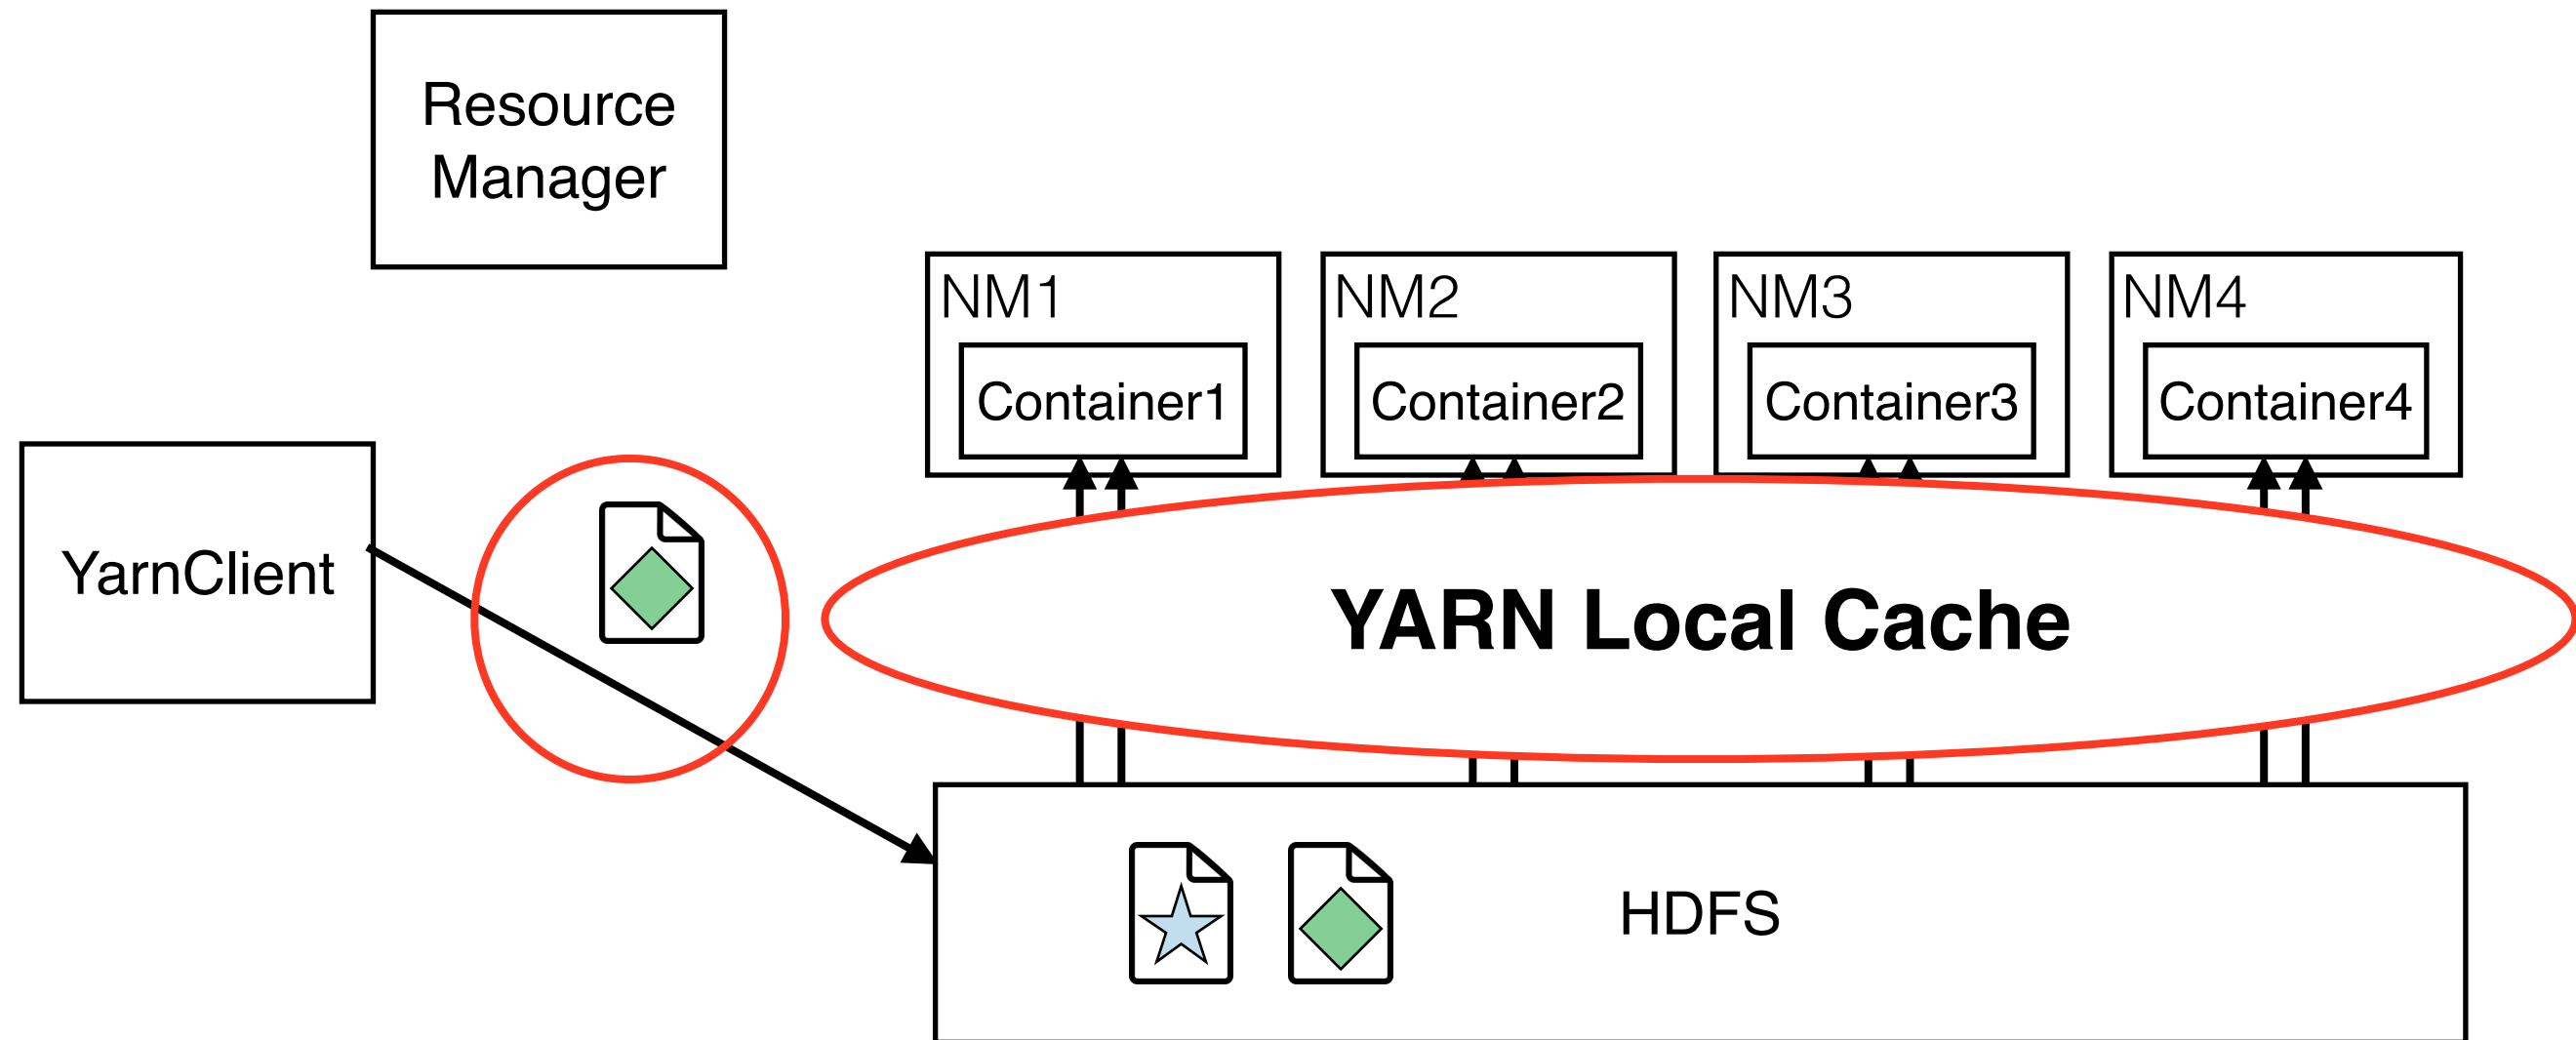


# YARN LOCAL CACHE

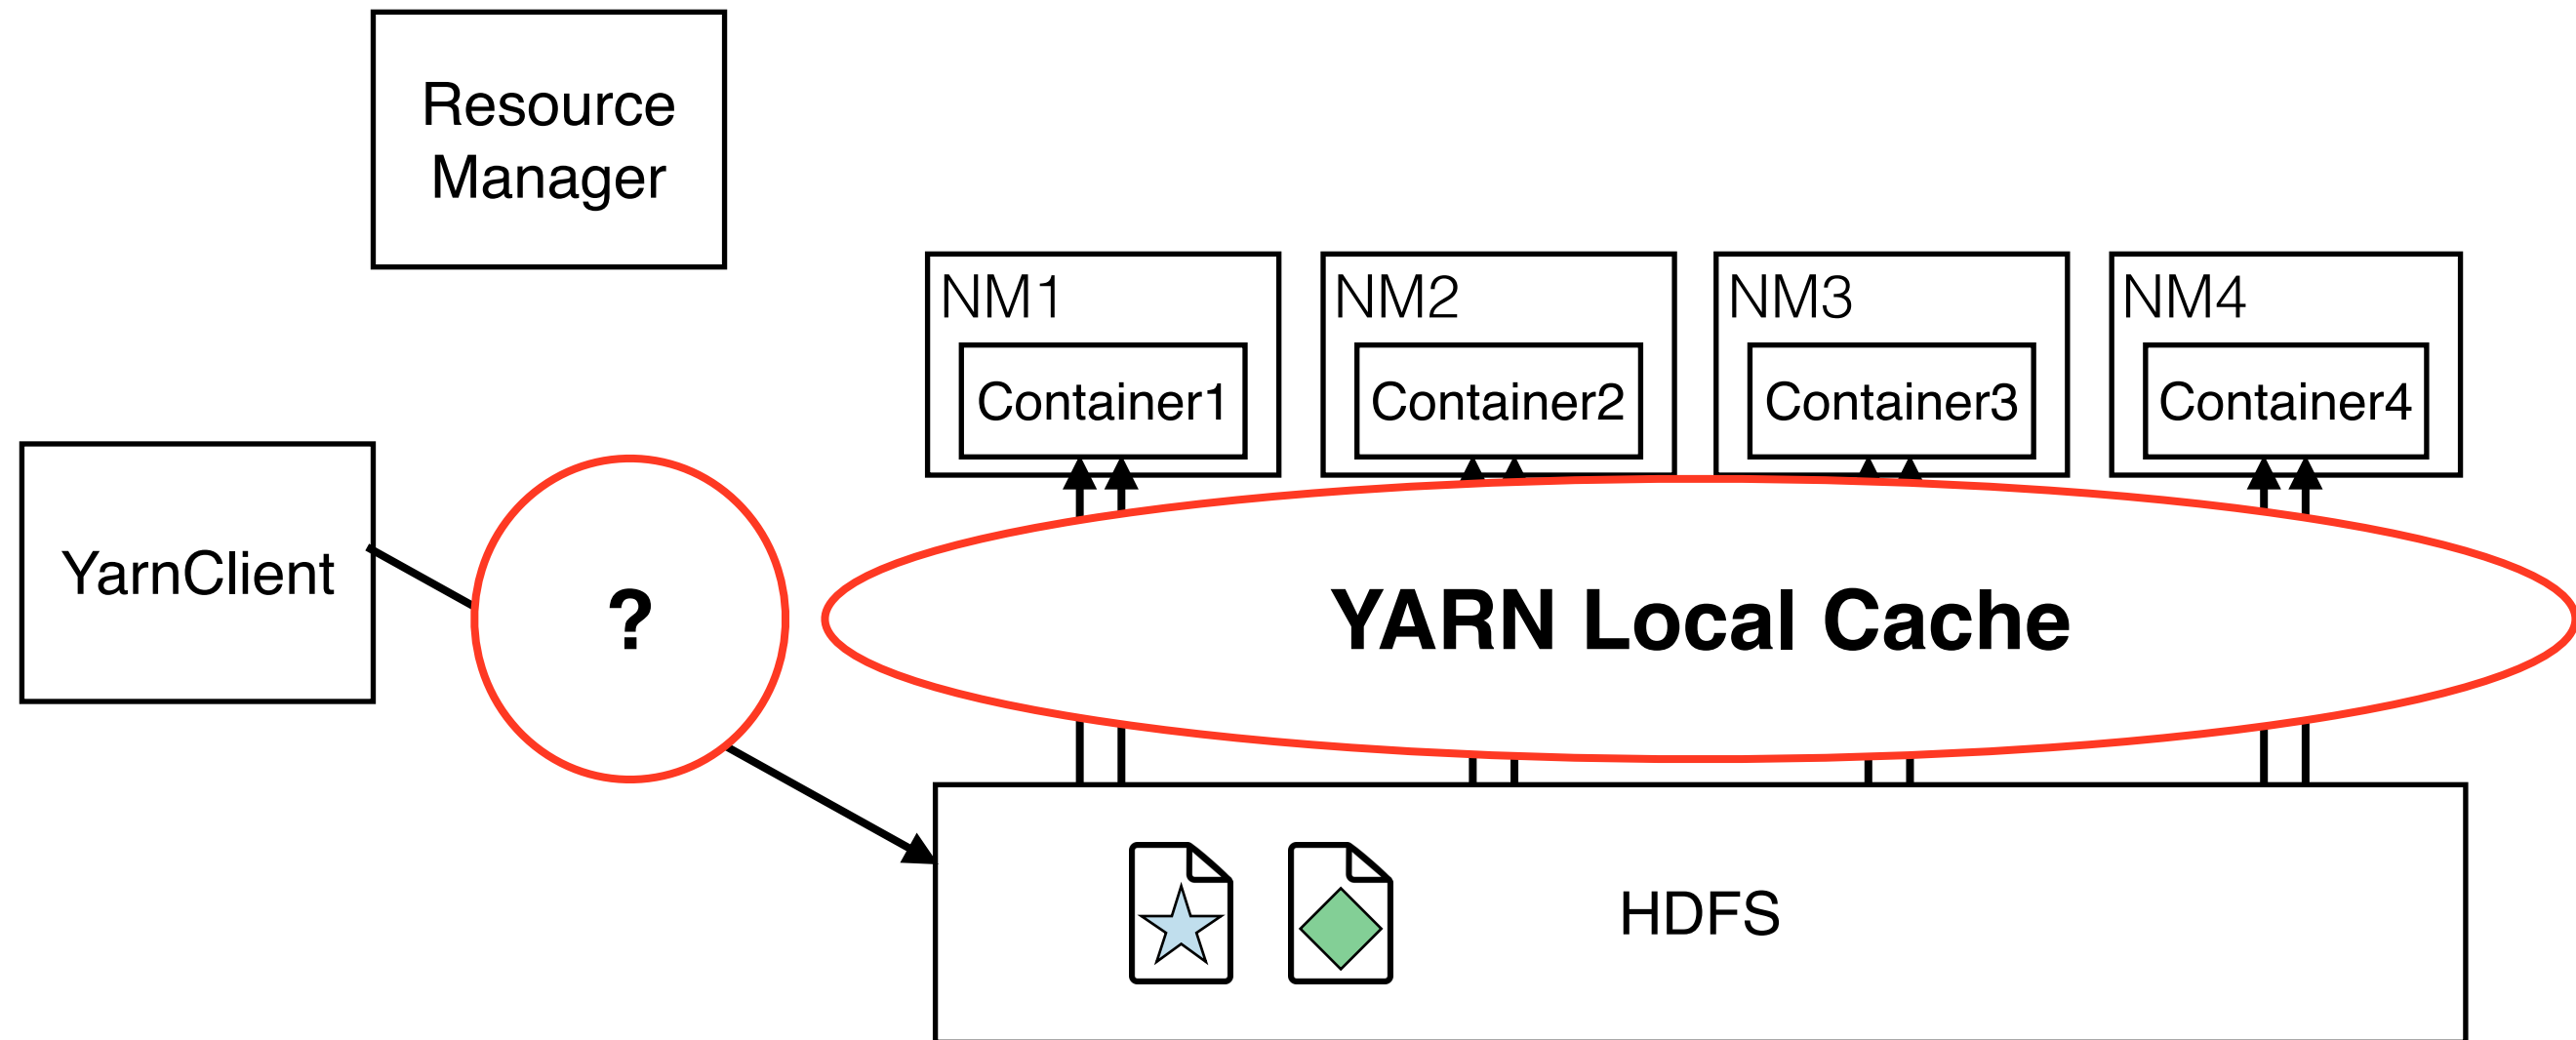
- Caches application resources on Node Manager local disk
- Maintains a target size via an LRU eviction policy
- Three resource visibilities
  - Public - Shared by everyone
  - User - The same user
  - Application - The same application
- Resources identified based on HDFS path



# LOCALIZATION OVERVIEW

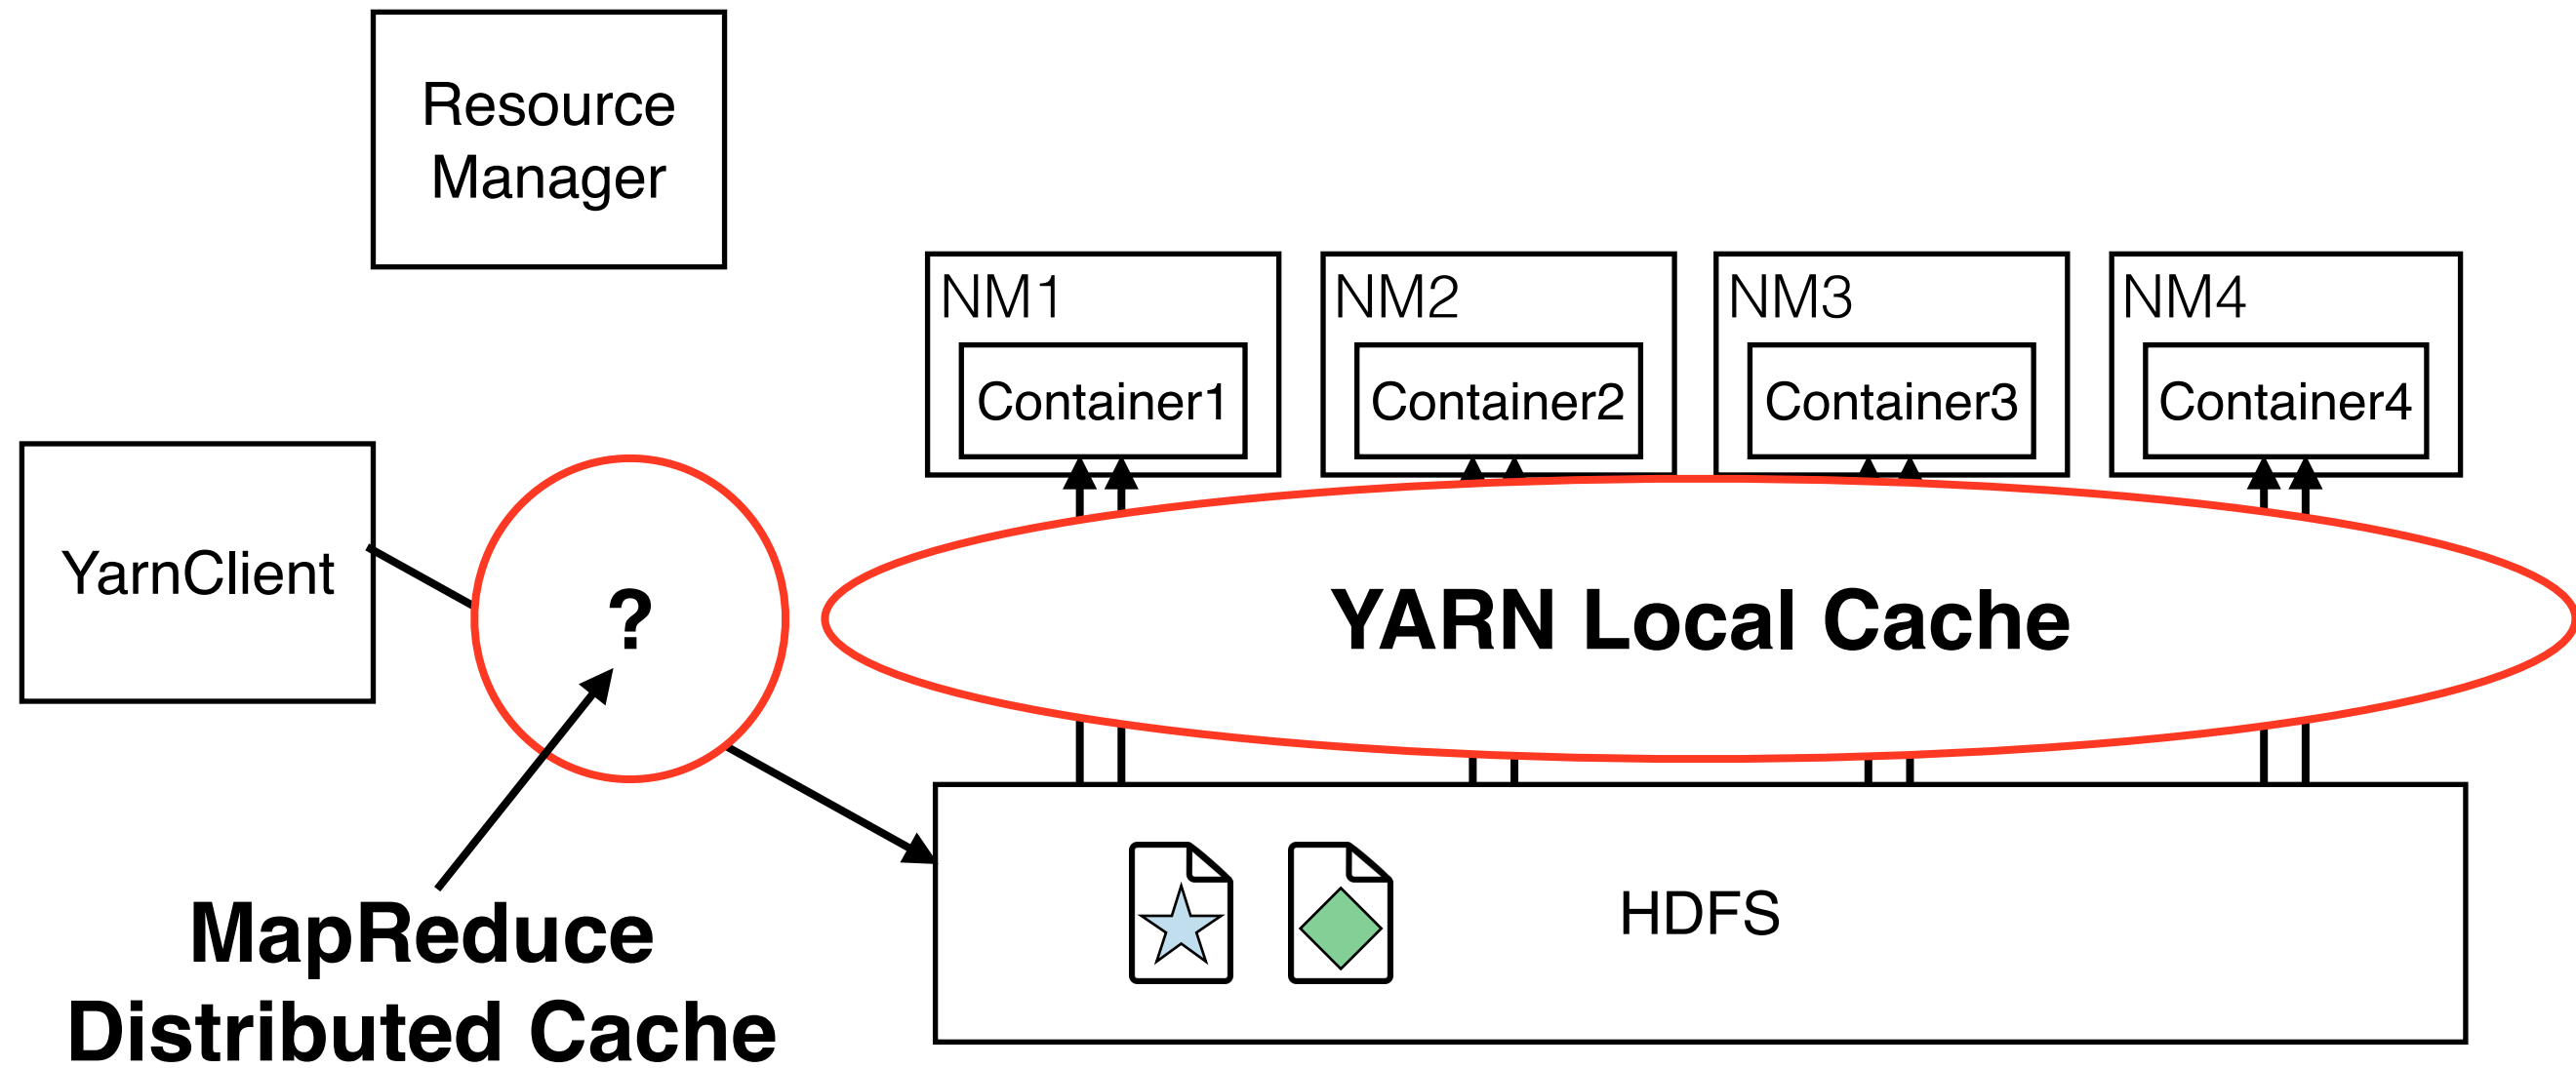


# LOCALIZATION OVERVIEW





# LOCALIZATION OVERVIEW



# MAPREDUCE DISTRIBUTED CACHE

- Caches MapReduce job resources in HDFS
- API for retrieving known paths in HDFS
- Sets YARN local cache visibility based on HDFS file permissions
- When sharing resources it does not manage:
  - Resource location
  - Resource cleanup
  - Resource integrity

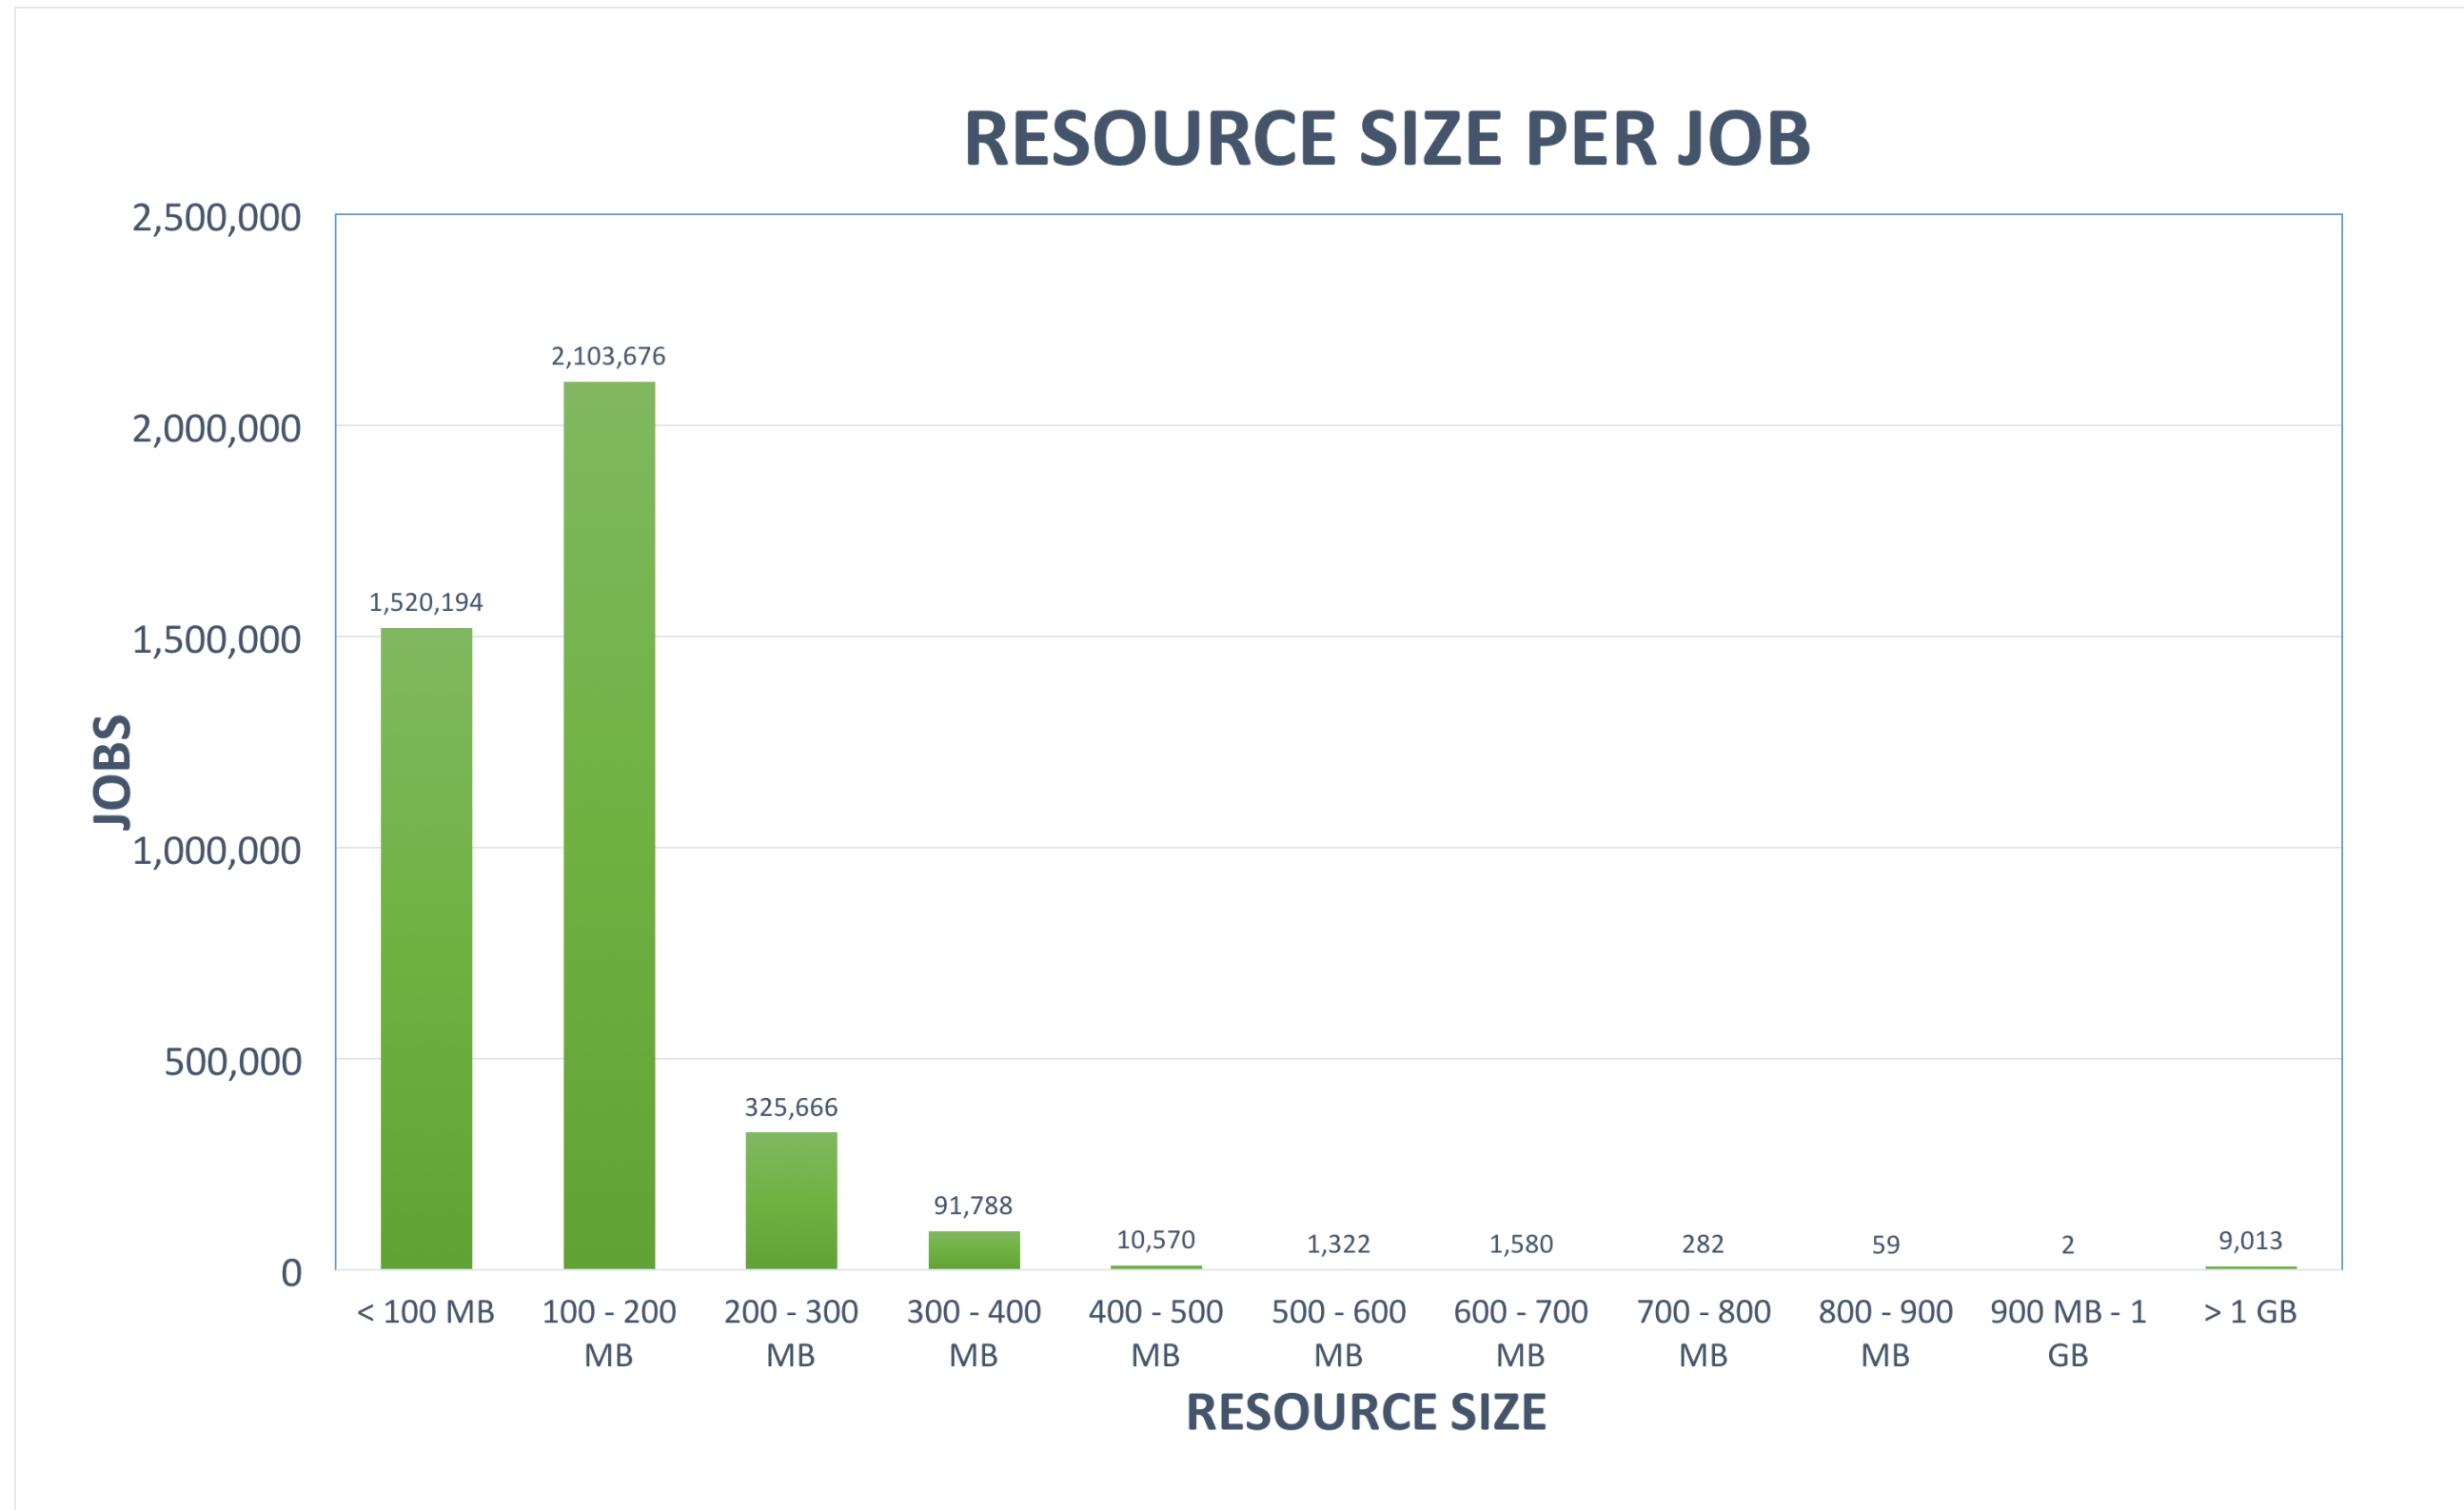


# RESOURCE SHARING IN PRACTICE

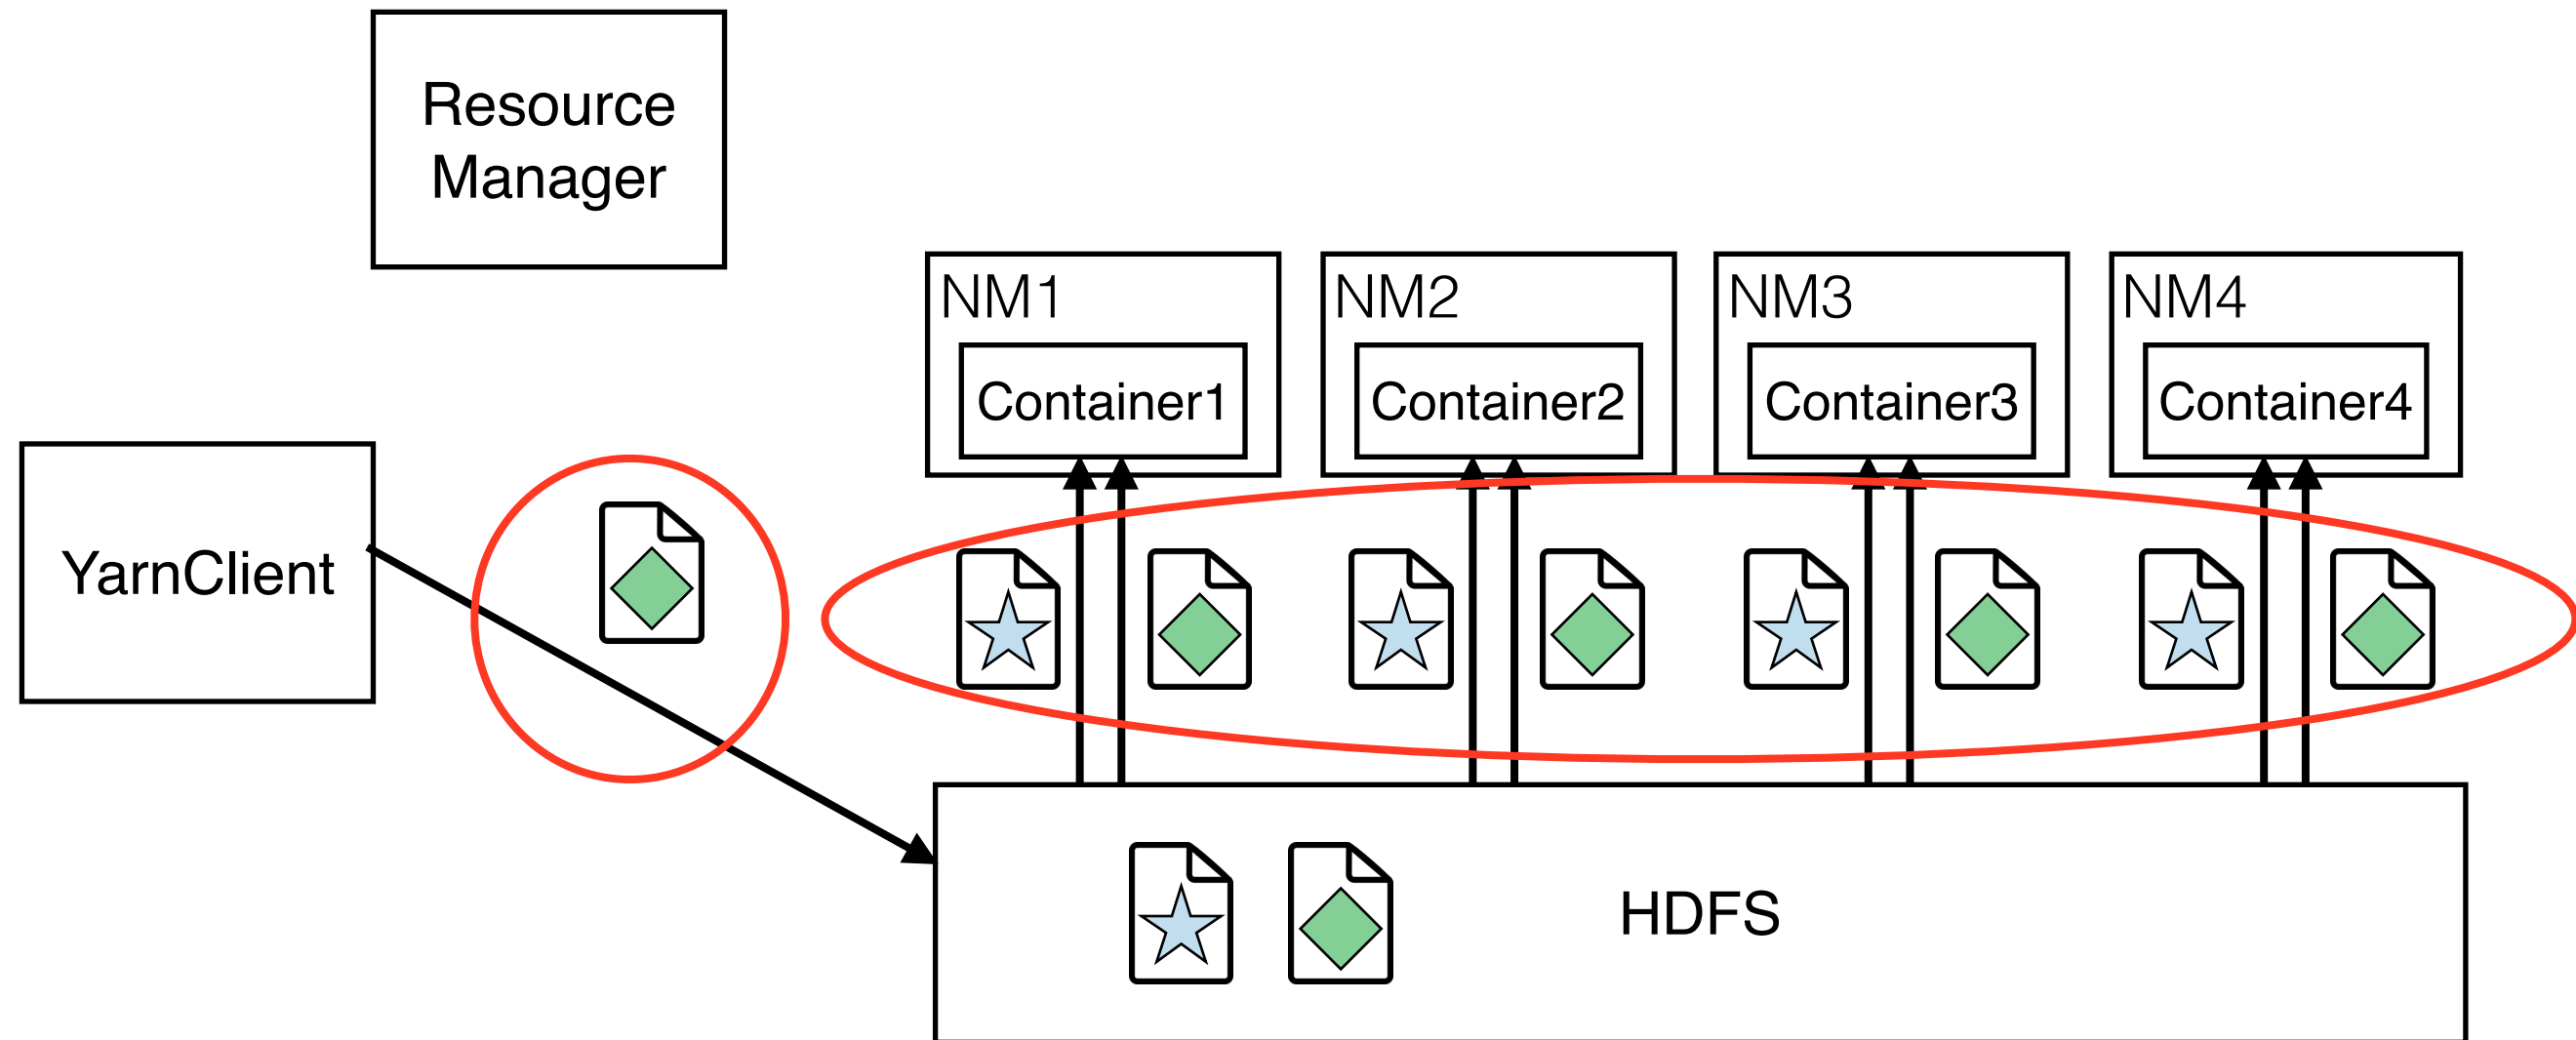
- Little to no resource sharing between applications
  - Lack of coordination at the HDFS level
  - Majority of applications are MapReduce jobs that upload resources into staging directories (i.e. application level visibility)



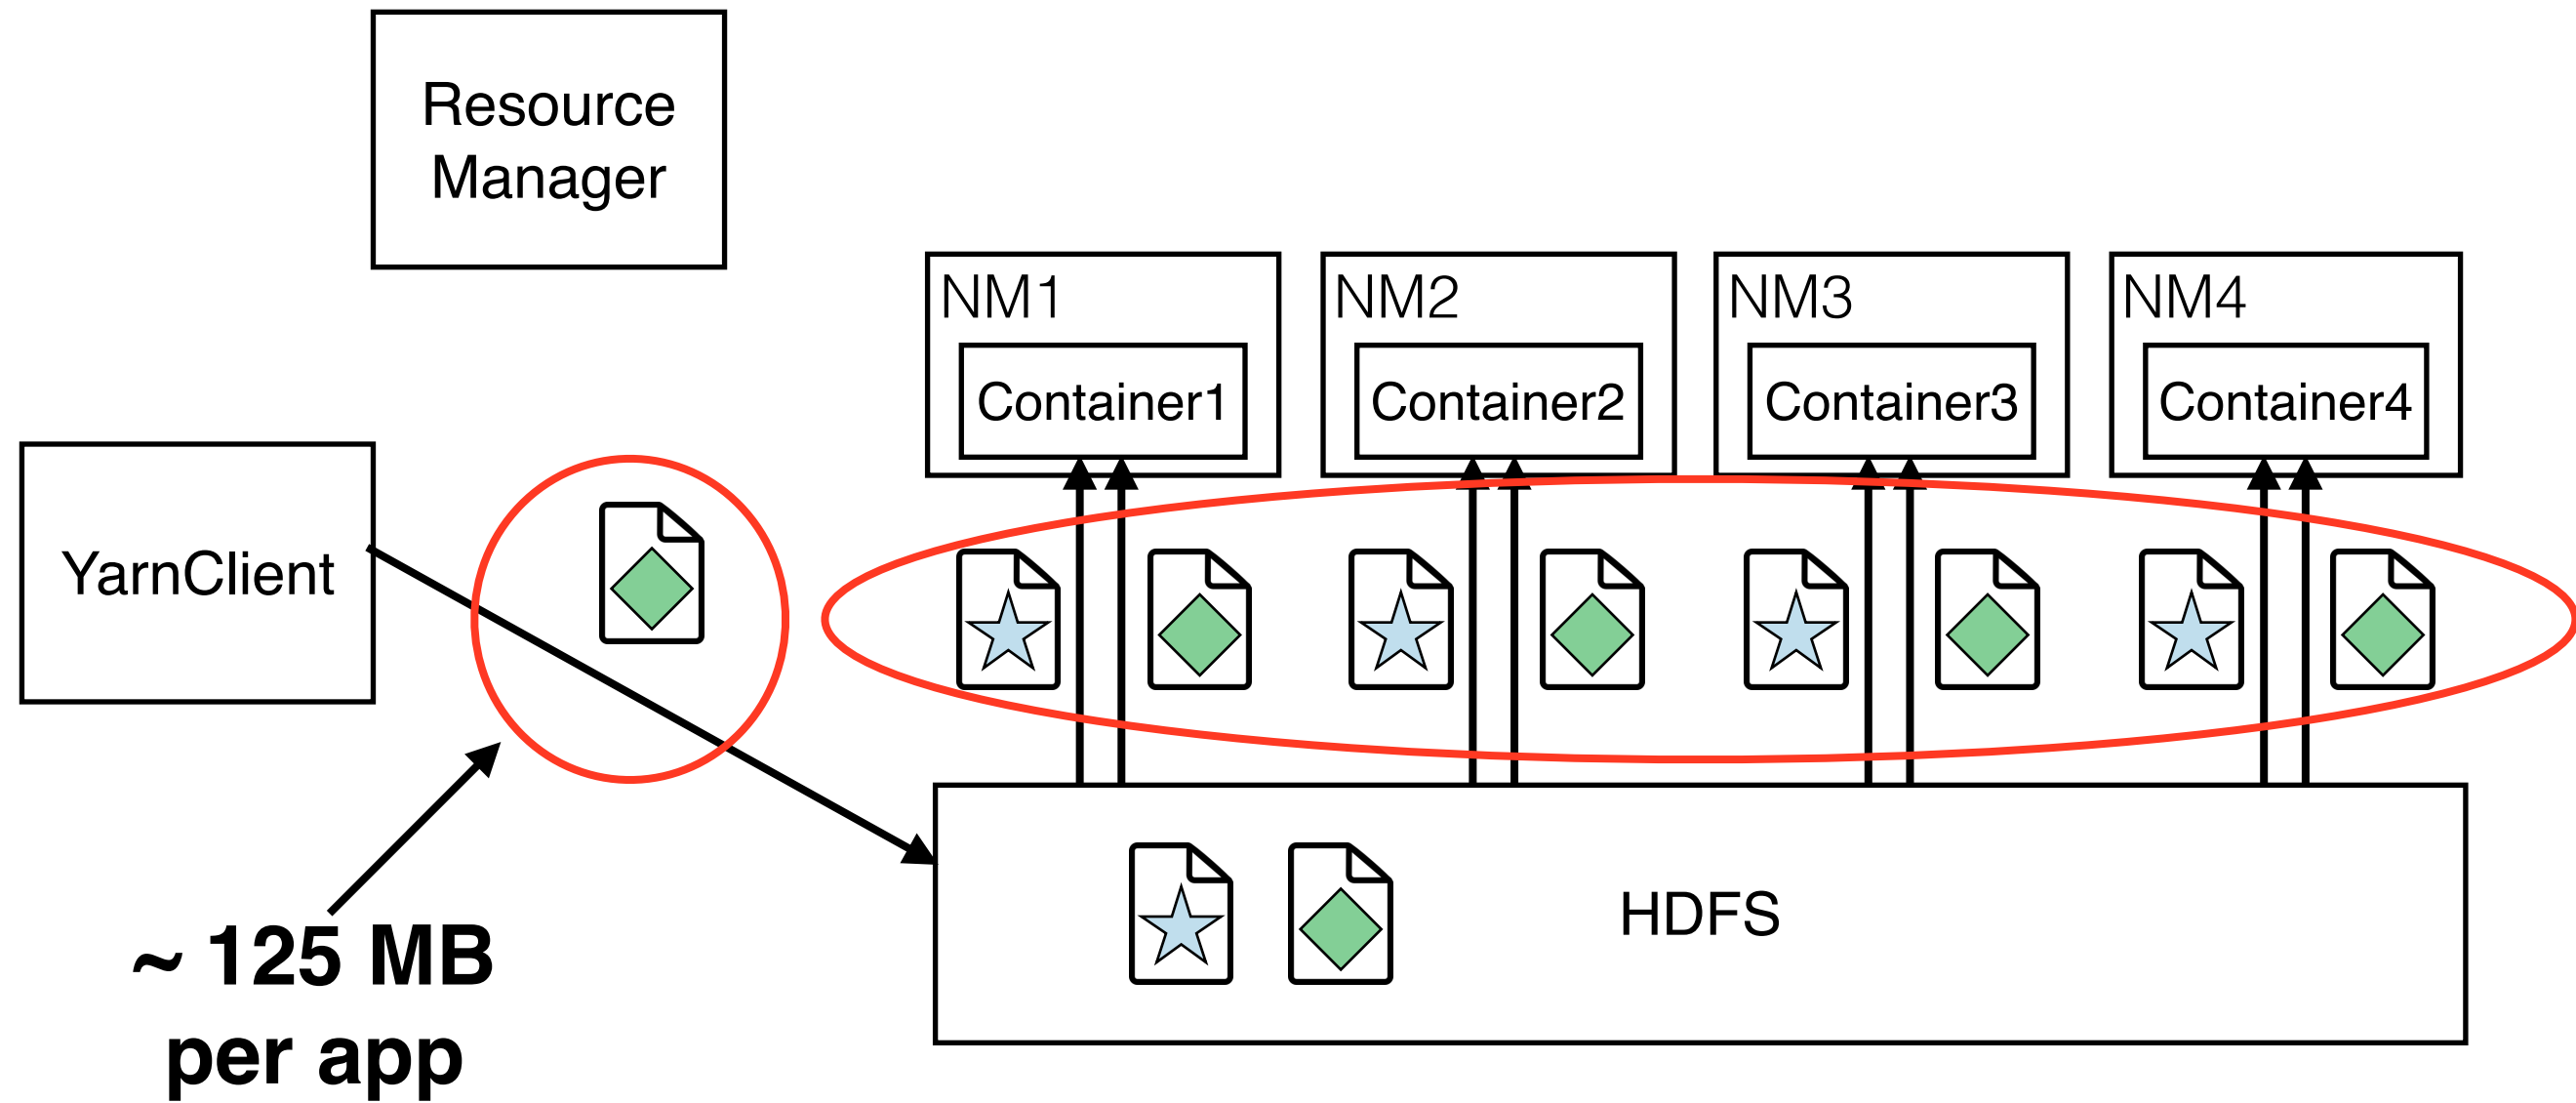
# HOW BIG IS THIS PROBLEM?



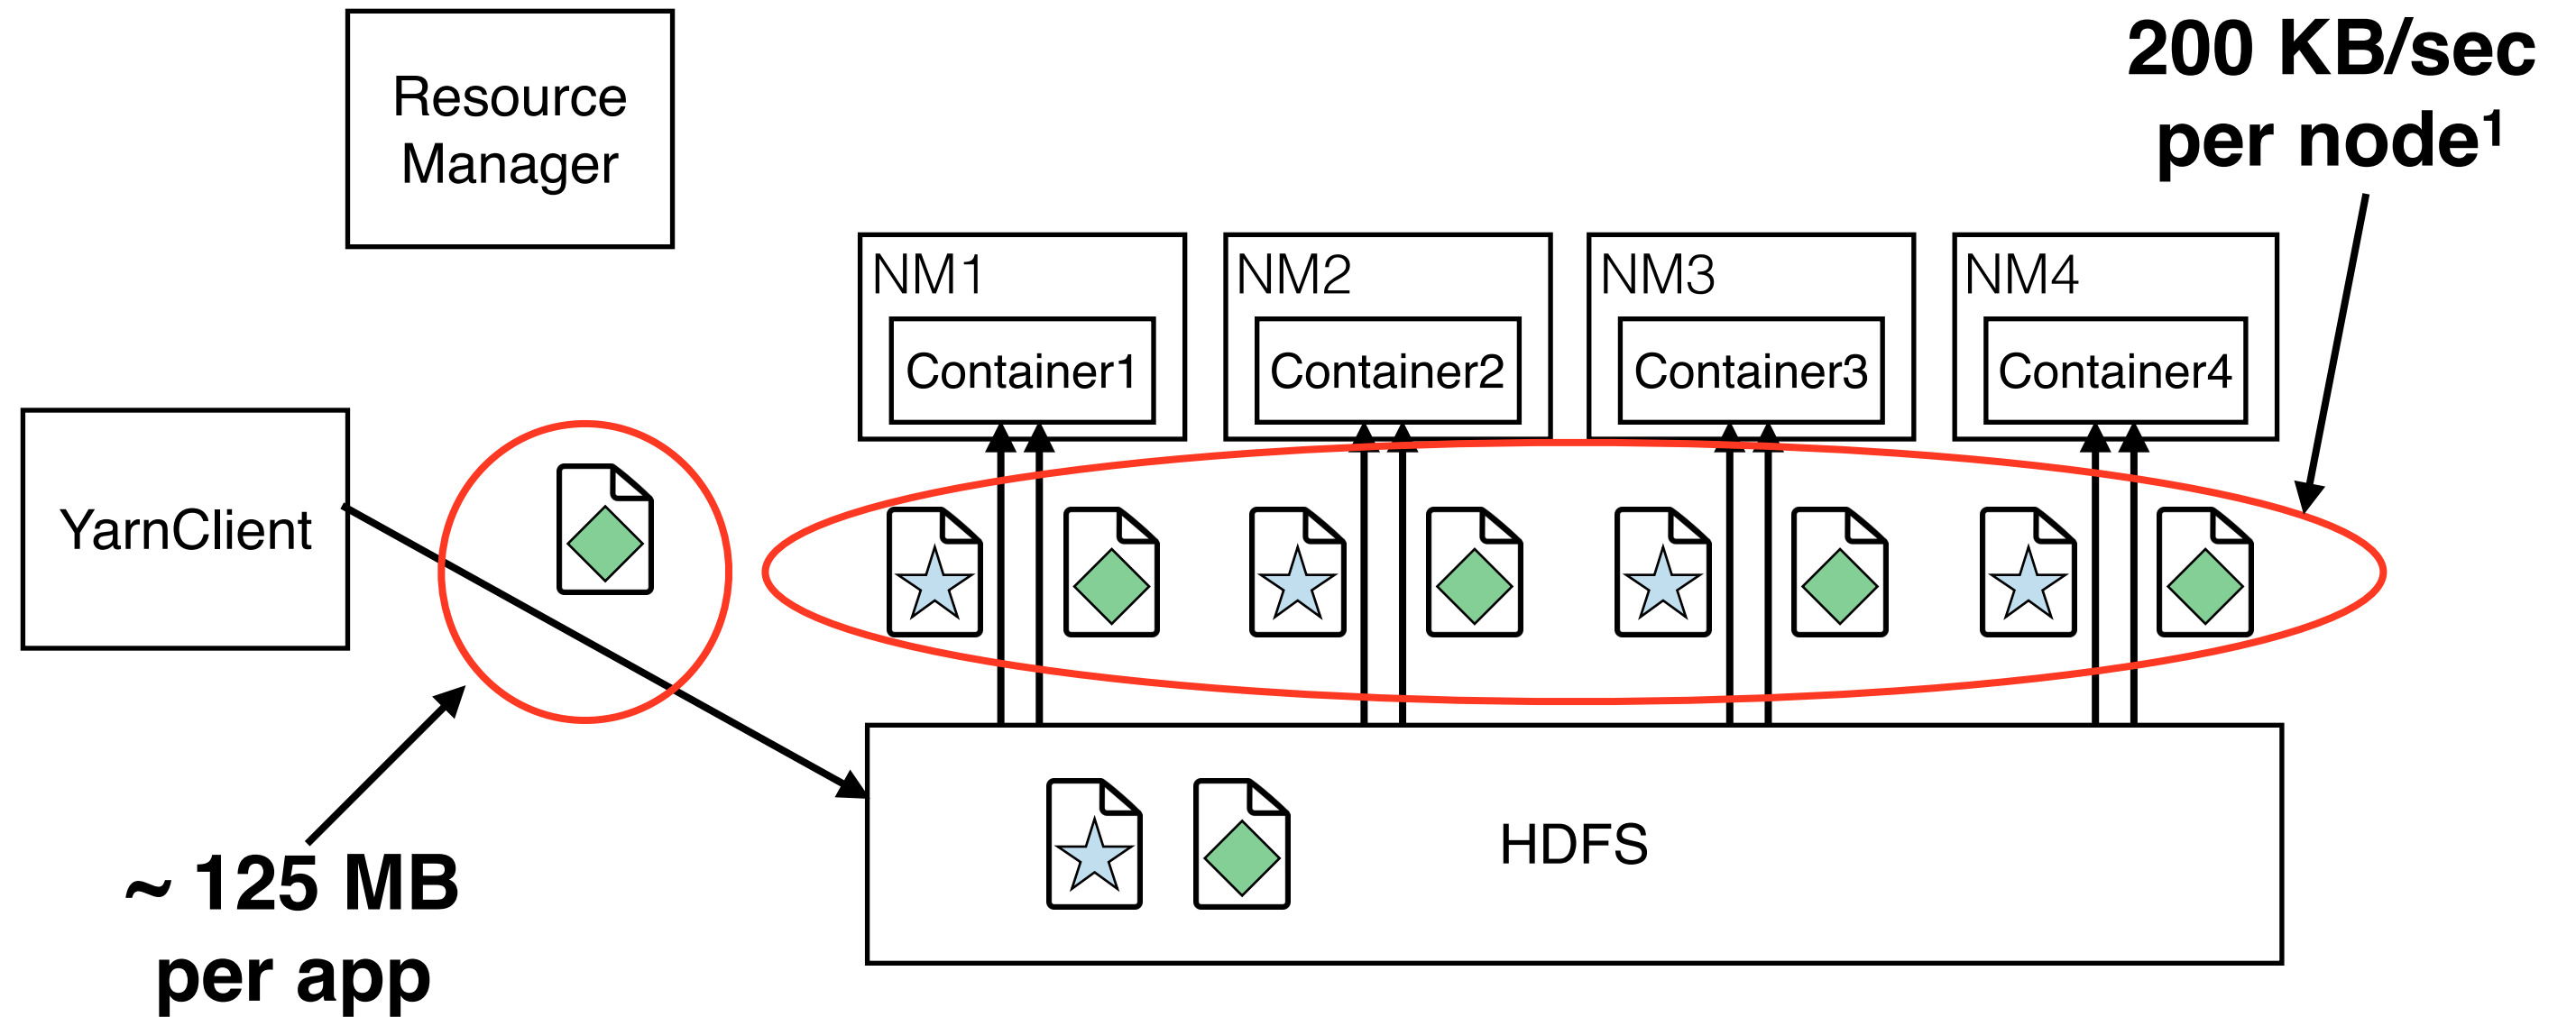
# HOW BIG IS THIS PROBLEM?



# HOW BIG IS THIS PROBLEM?



# HOW BIG IS THIS PROBLEM?



1. Assuming 1.7 containers launched per node per minute



# YARN SHARED CACHE

- YARN-1492
- Currently in production at Twitter for ~ 2 years
- 100's of thousands of applications use it a day
  - MapReduce jobs (i.e. Scalding/Cascading)
- Working towards open source release
  - The YARN feature is in 2.7 (Beta)
  - Full MapReduce support coming (MAPREDUCE-5951)





# DESIGN GOALS

- Scalable
  - Accommodate large number of cache entries
    - Thousands of cached resources
  - Have minimal load on Namenode and Resource Manager
  - Handle spikes in cache size gracefully



# DESIGN GOALS

- Secure
  - Identify resources in the cache based on their contents, not storage location
  - Trust that if the cache says it has a resource, it actually has the resource



# DESIGN GOALS

- Fault tolerant
  - YARN applications continue running without shared cache
  - Shared cache should tolerate restarts
    - Either persist or recreate cache resource meta data



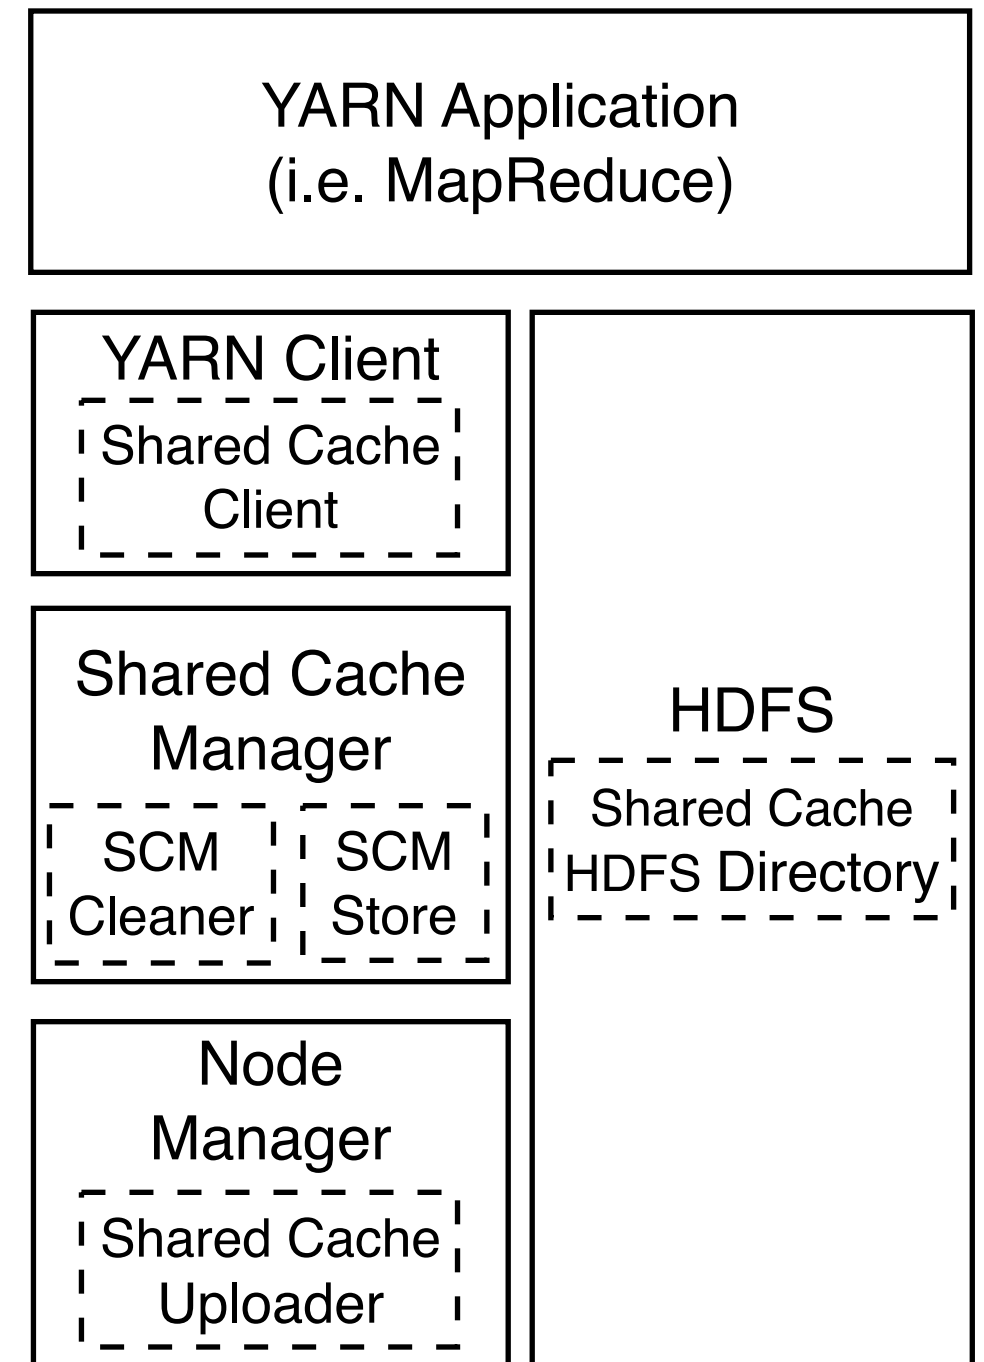
# DESIGN GOALS

- Transparent
  - YARN application developer perspective
    - Management of the cache (adding/deleting resources)
  - MapReduce job developer perspective
    - Jobs can use the cache with no change



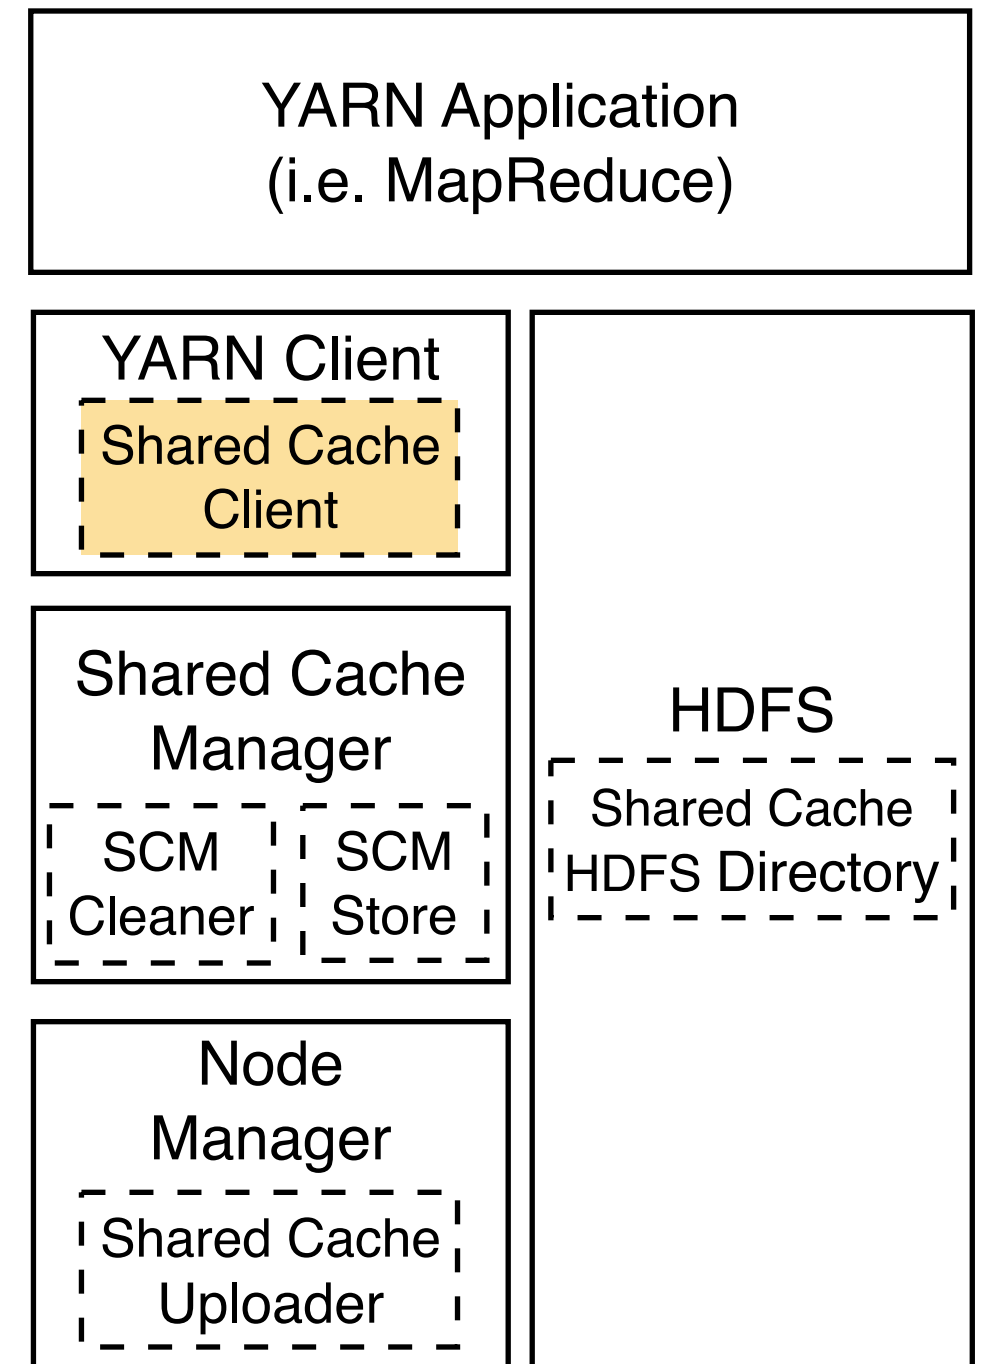
# DESIGN OVERVIEW

- Shared Cache Client
- Shared Cache HDFS Directory
- Shared Cache Manager
- Shared Cache Uploader



# SHARED CACHE CLIENT

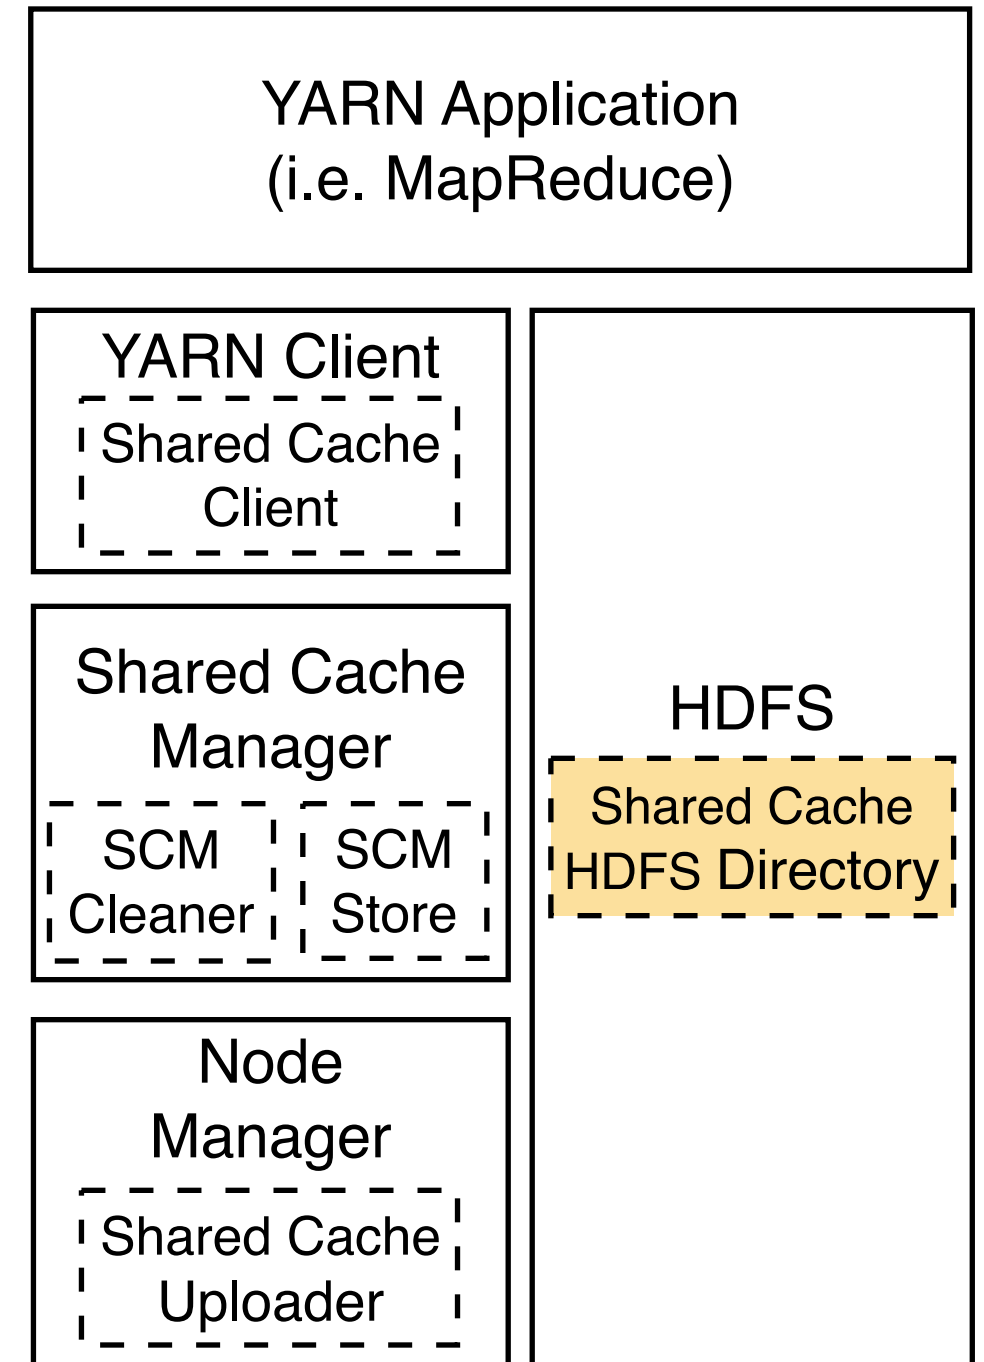
- Interacts with shared cache manager
- Responsible for
  - Computing checksum of resources
  - Claiming application resources
- Public API:  
`Path use(String checksum, String appId);`
- Return Value
  - HDFS path to resource in cache
  - Null if not in the shared cache



# HDFS DIRECTORY

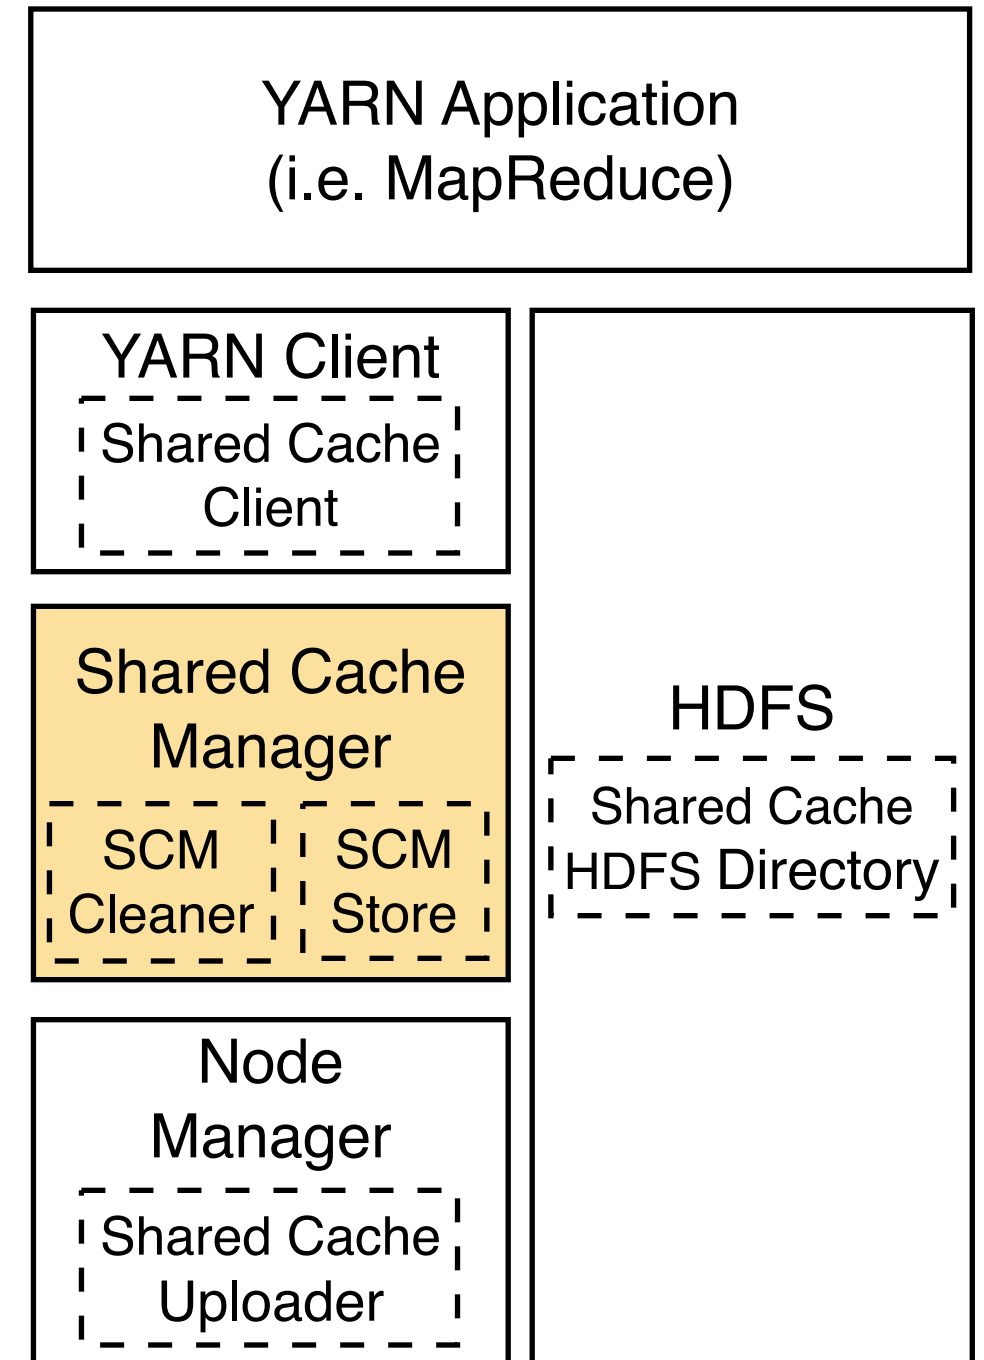
- Stores shared cache resources
- Protected by HDFS permissions
  - Globally readable
  - Writing restricted to trusted user
- Only modified by uploader and cleaner
- Directory structure:

```
/sharedcache/a/8/9/a896857d078/foo.jar  
/sharedcache/5/0/f/50f11b09f87/bar.jar  
/sharedcache/a/6/7/a678cb1aa8f/job.jar
```



# SHARED CACHE MANAGER

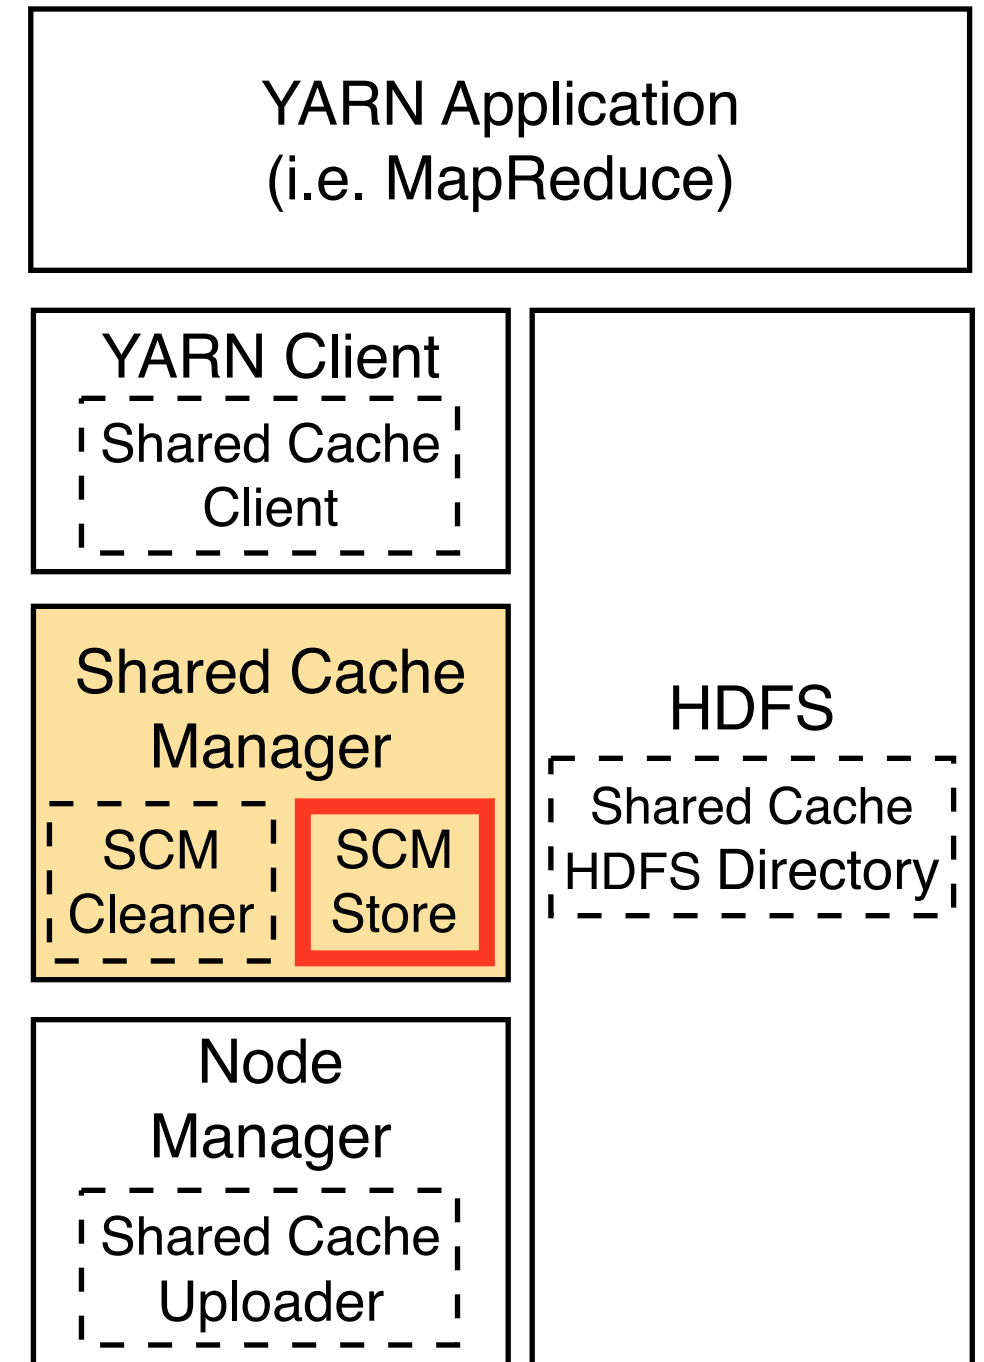
- Serves requests from clients
- Manages resources in the shared cache
  - Metadata
  - Persisted resources in HDFS





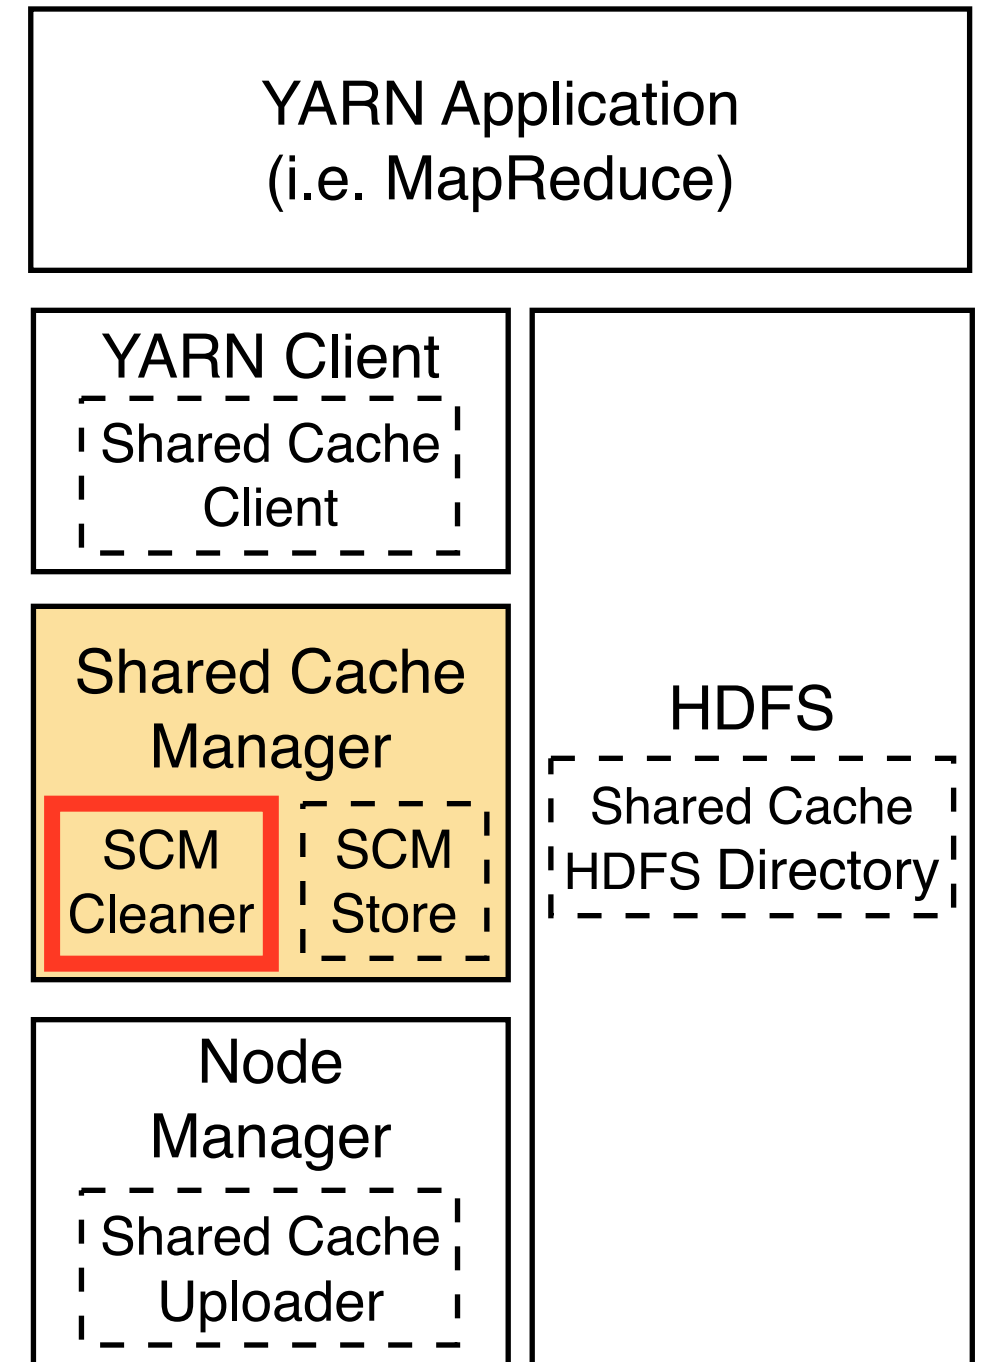
# SHARED CACHE MANAGER

- Store
  - Maintains cache metadata
    - Resources in the shared cache
    - Applications using resources
  - Implementation is pluggable
    - Currently uses in-memory store
    - Recreates state after restart



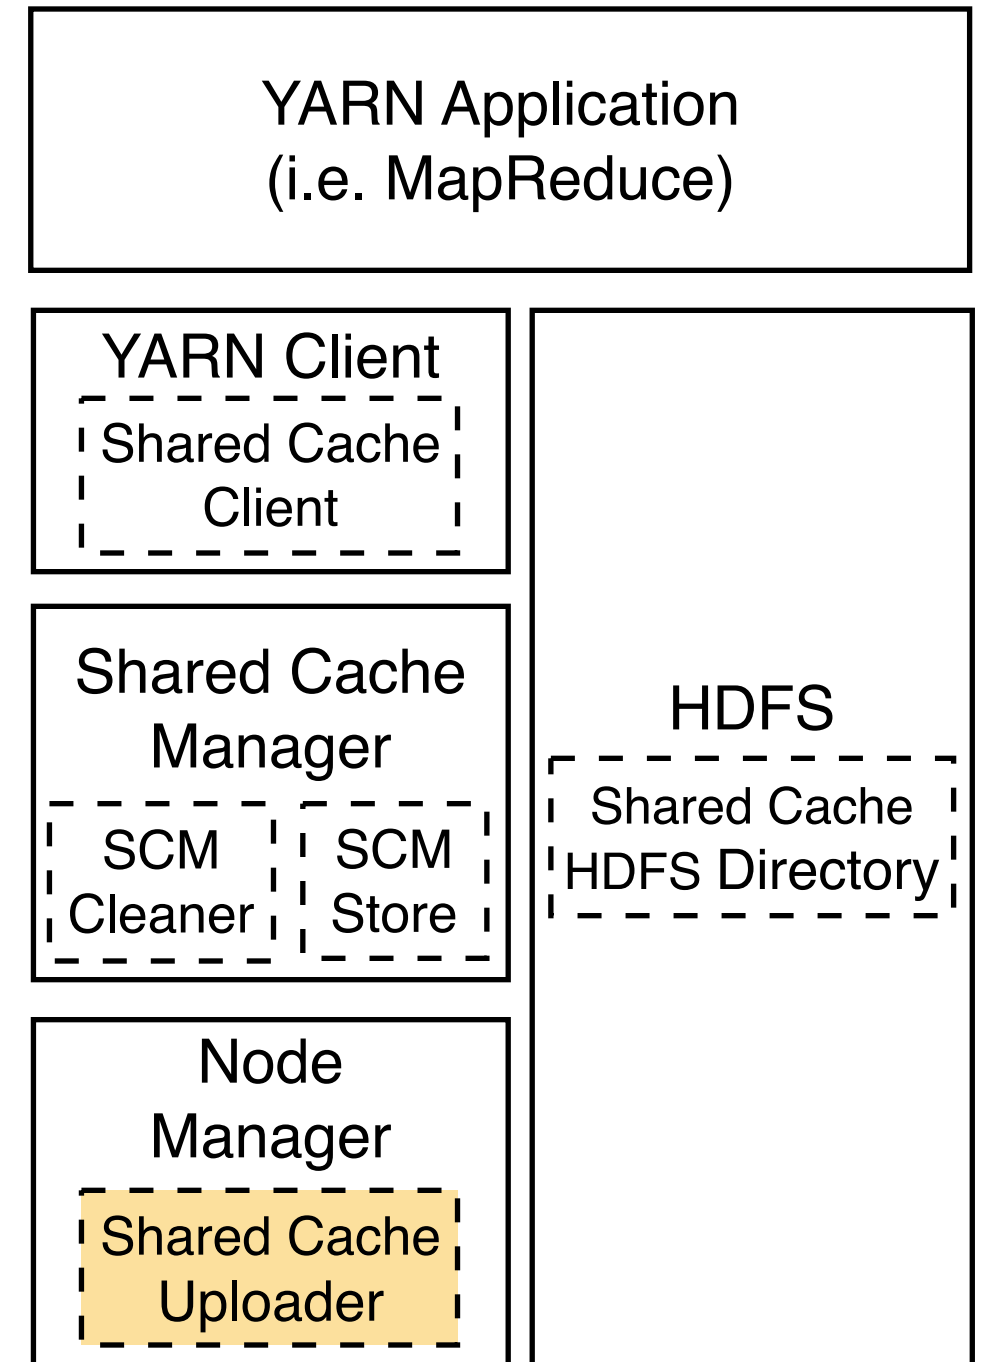
# SHARED CACHE MANAGER

- Cleaner
  - Maintains persisted resources in HDFS
  - Runs periodically (default: once a day)
- Evicts resource if both of the following hold:
  - Resource has exceeded stale period (default: 1 week)
  - There are no live applications using the resource

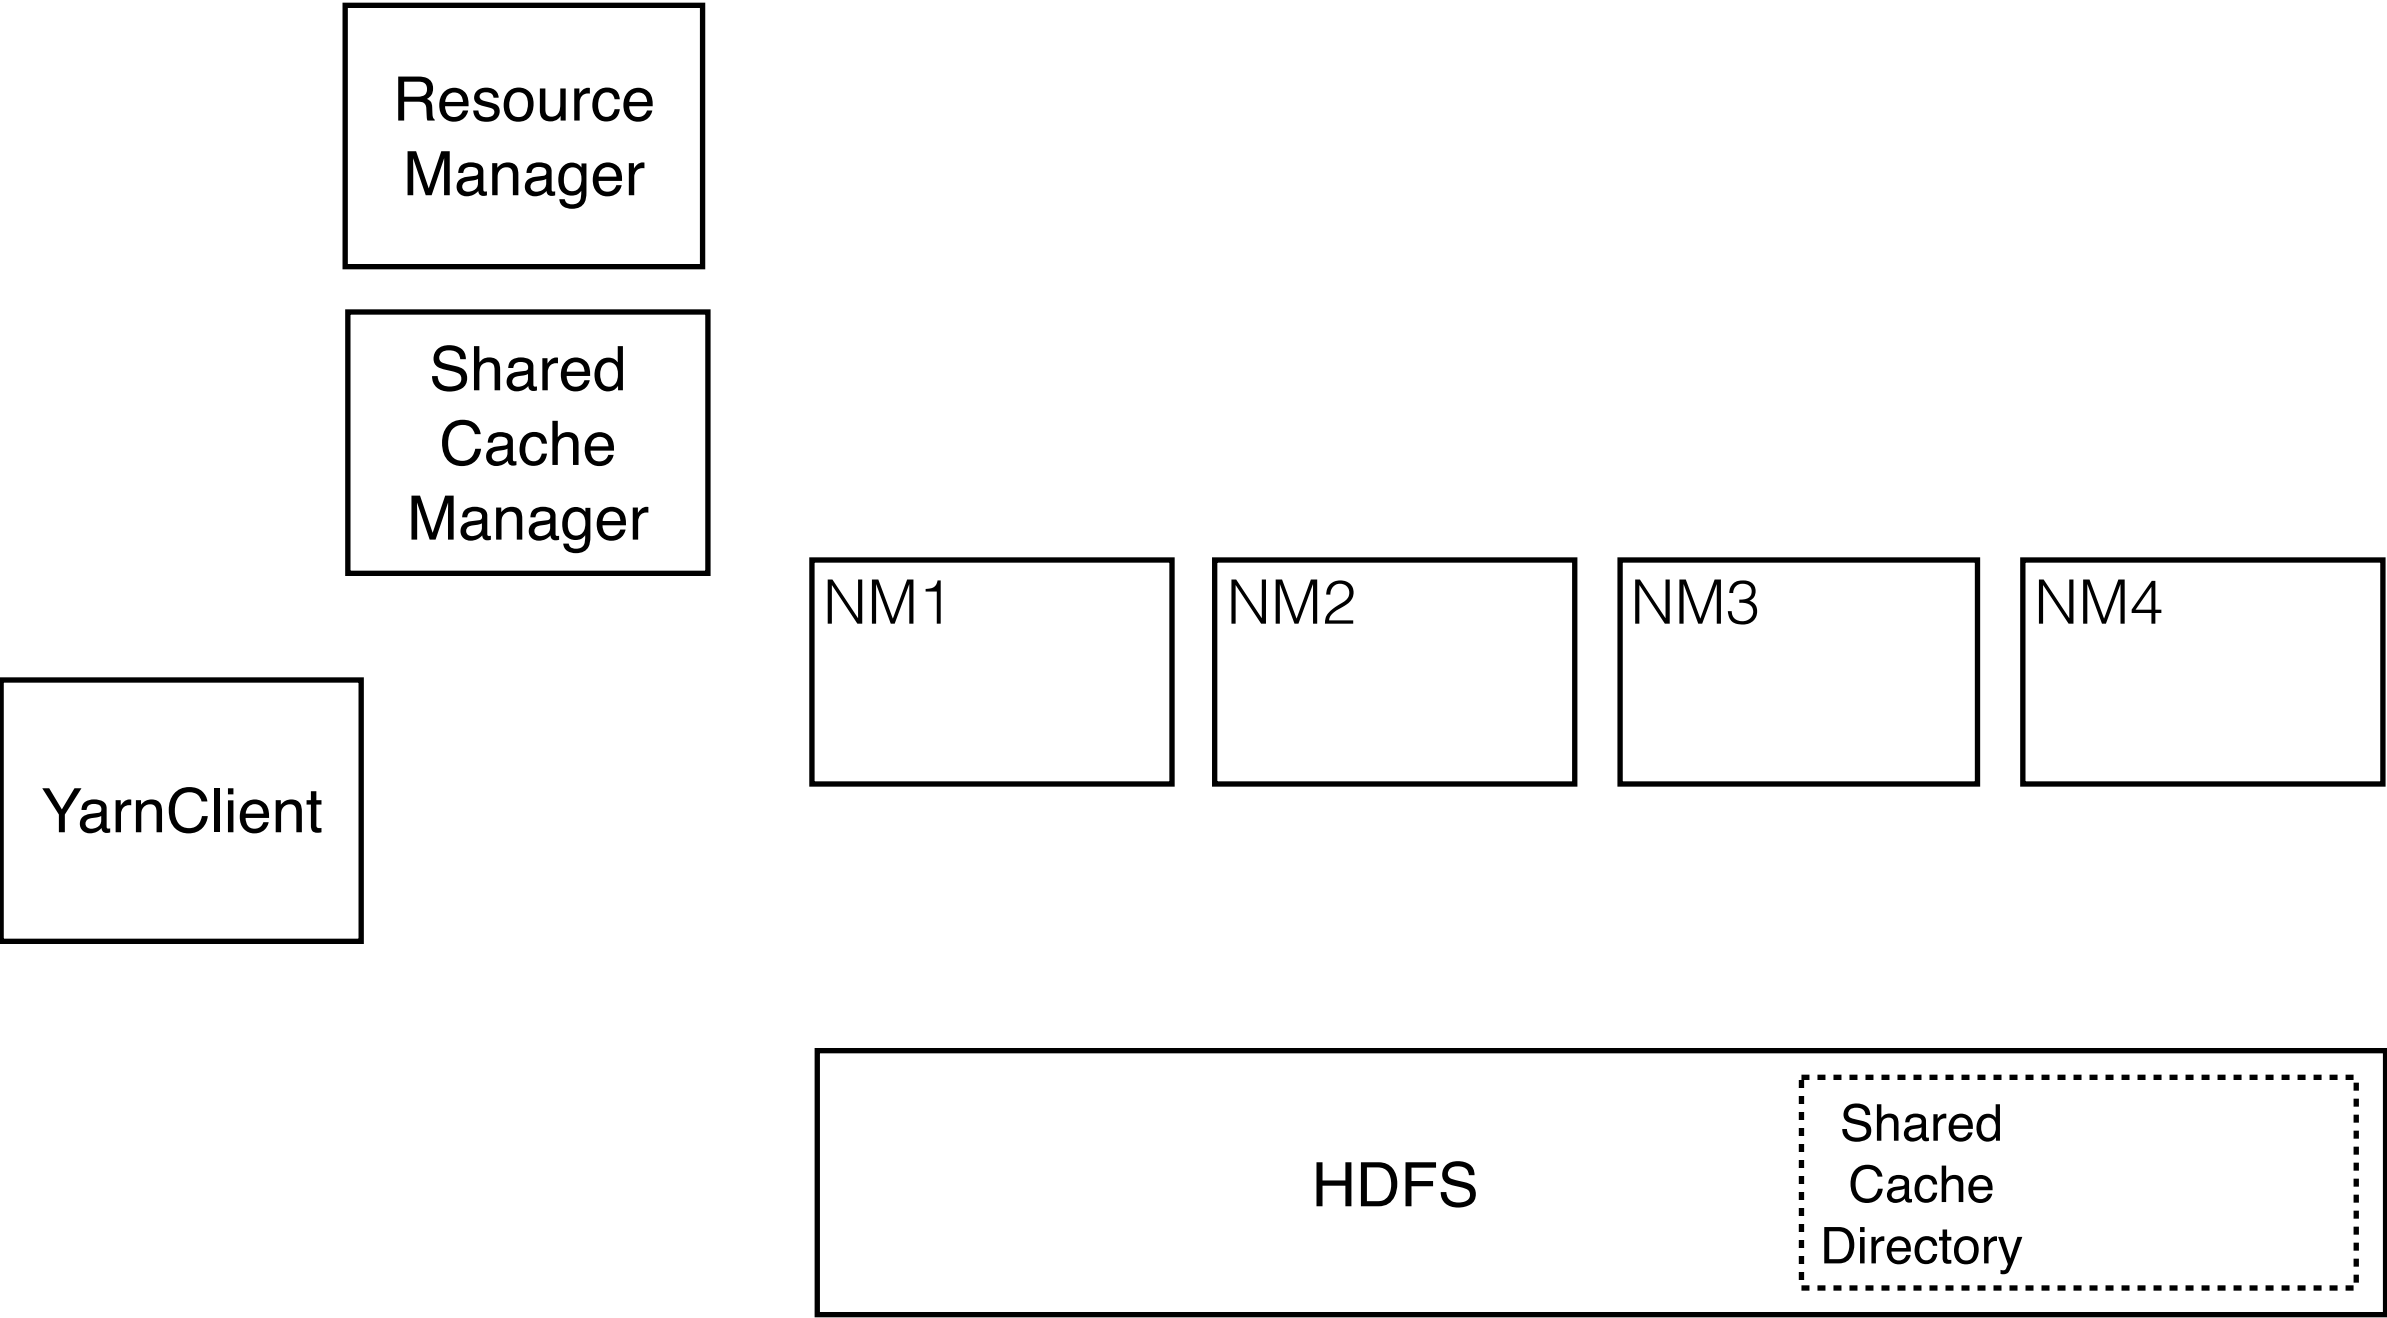


# SHARED CACHE UPLOADER

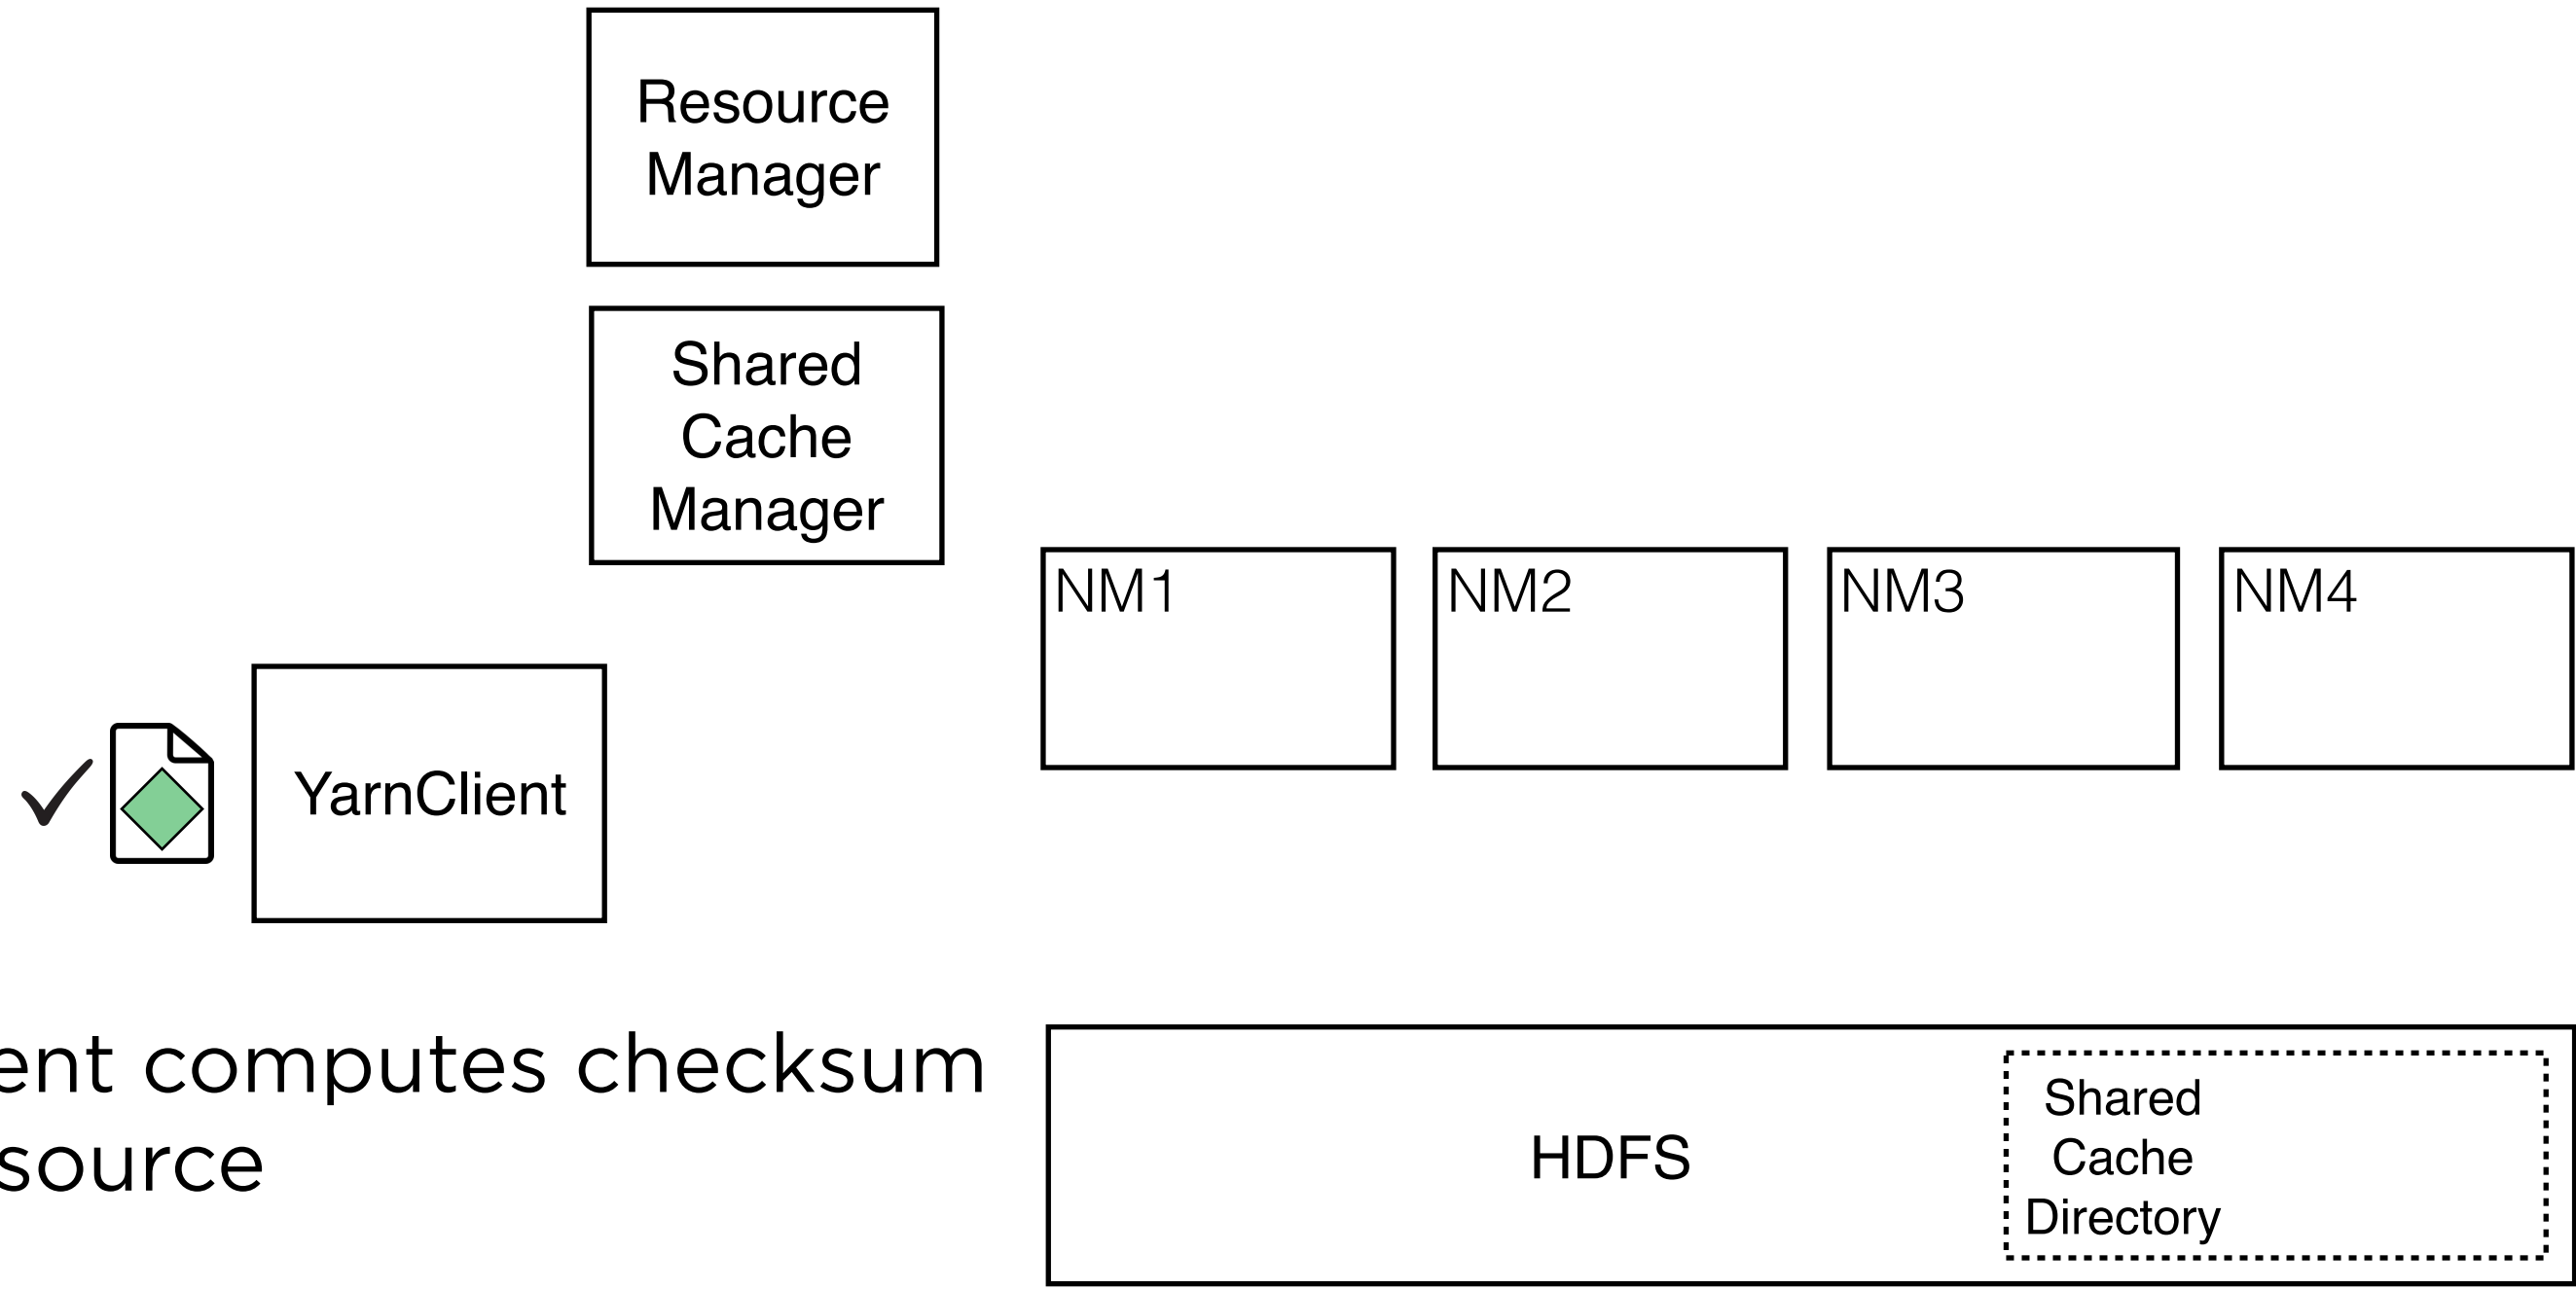
- Runs on Node Manager
- Adds resources to the shared cache
  - Verifies resource checksum
  - Uploads resource to HDFS
  - Notifies the shared cache manager
- Asynchronous from container launch
- Best-effort: failure does not impact running applications



# ADDING TO THE CACHE



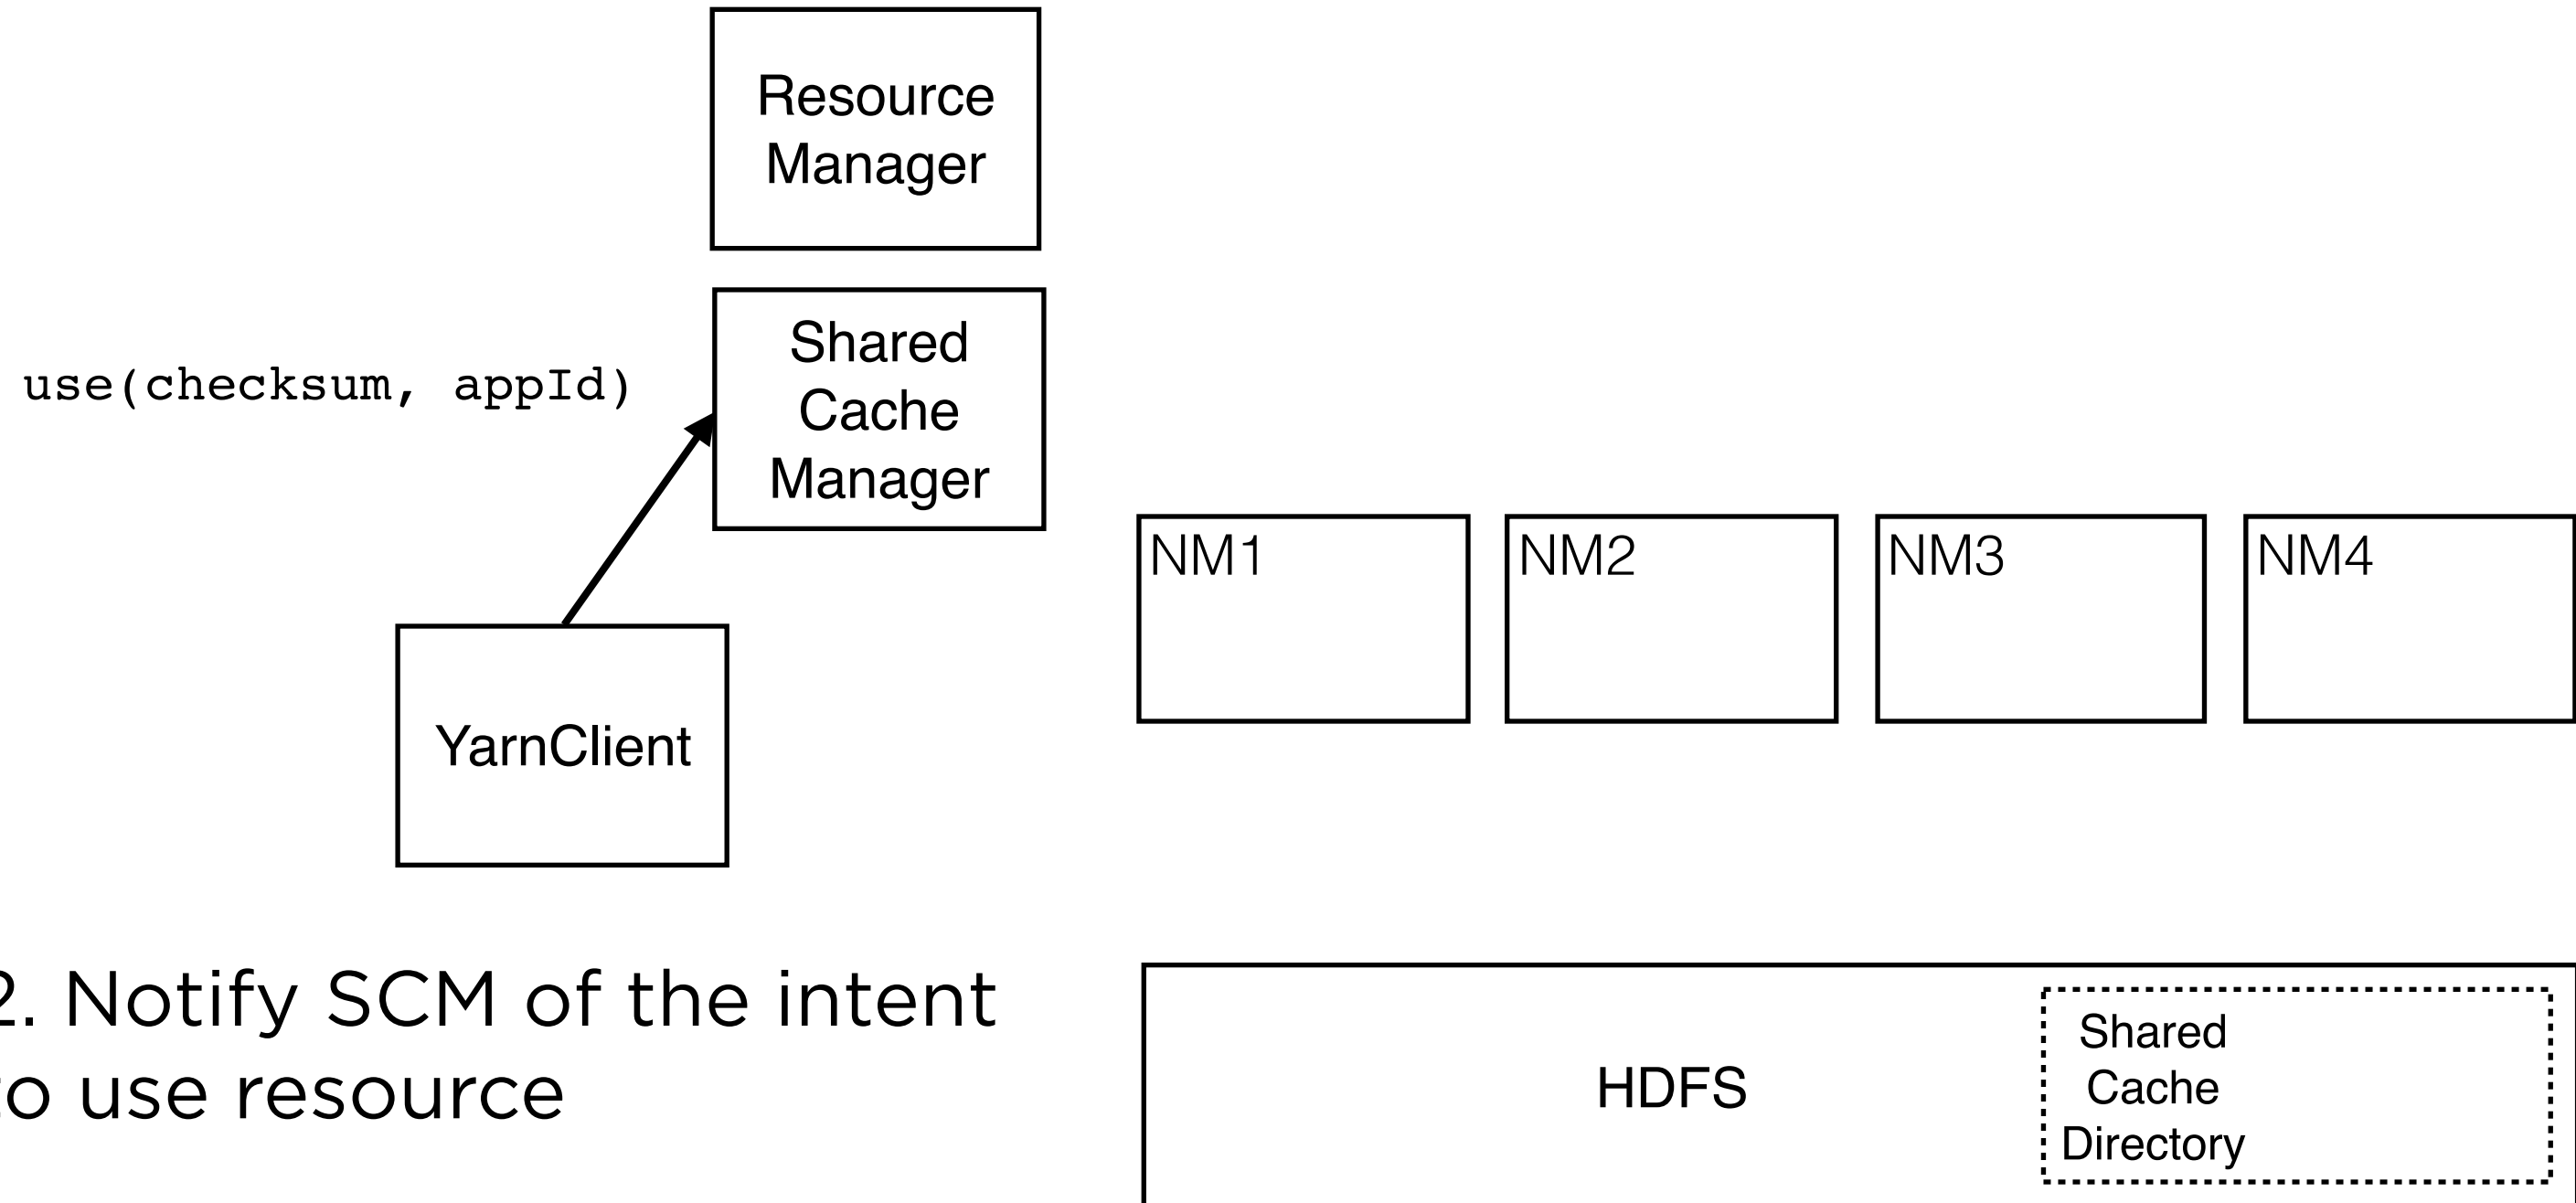
# ADDING TO THE CACHE



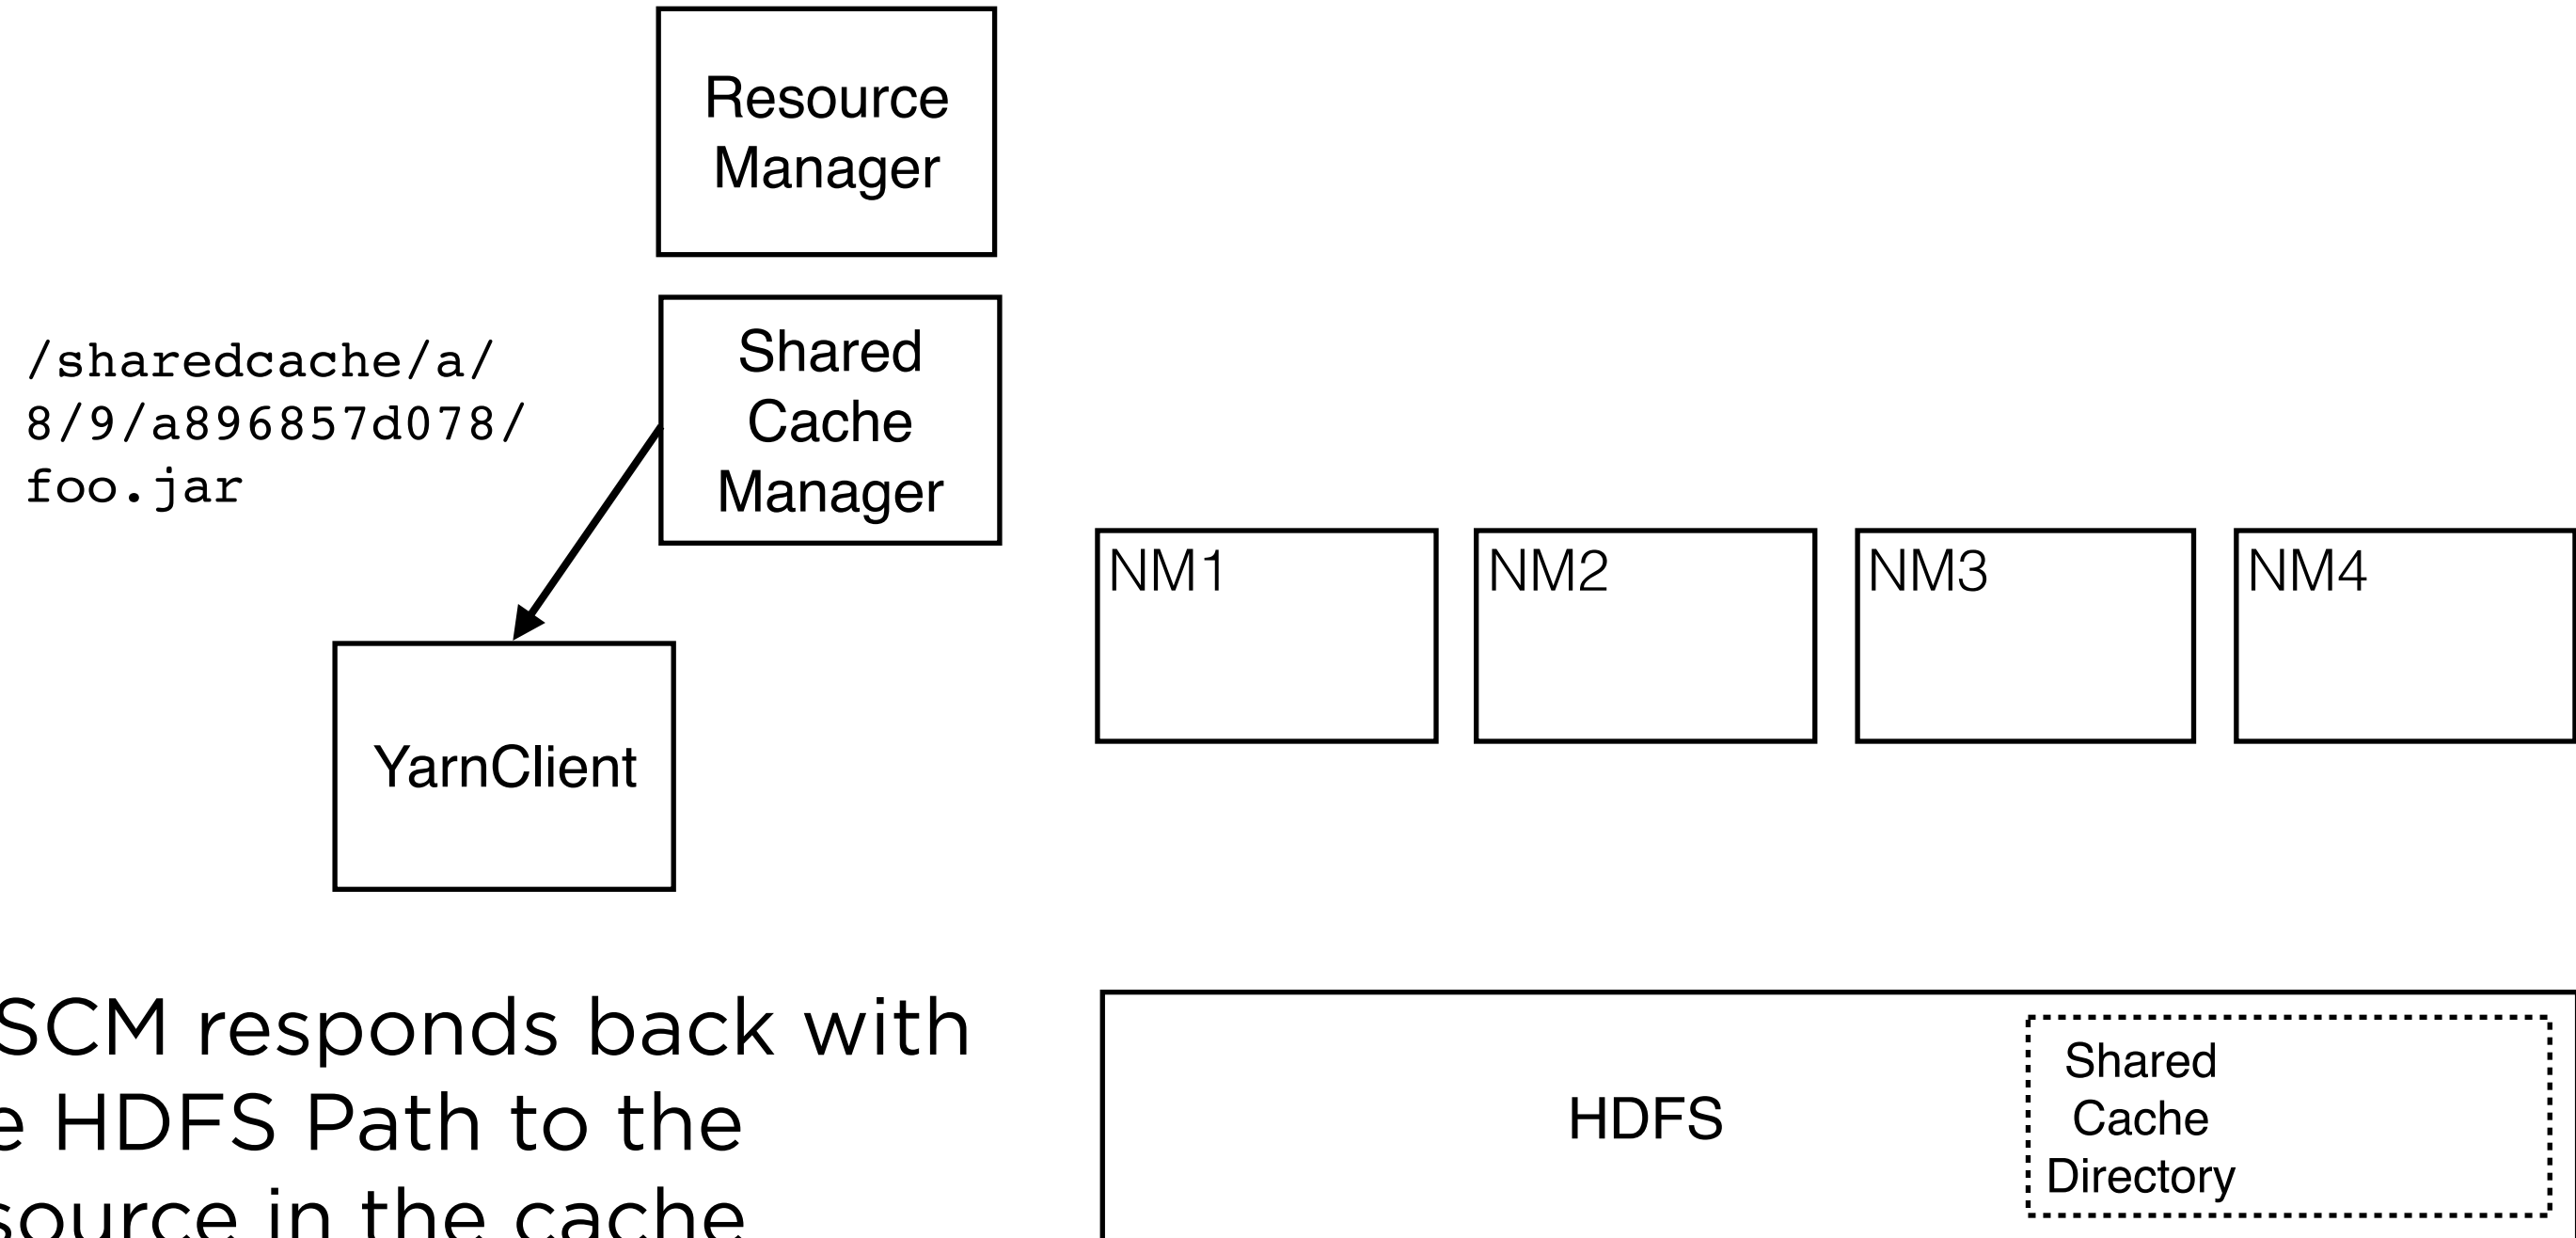
1. Client computes checksum of resource



# ADDING TO THE CACHE



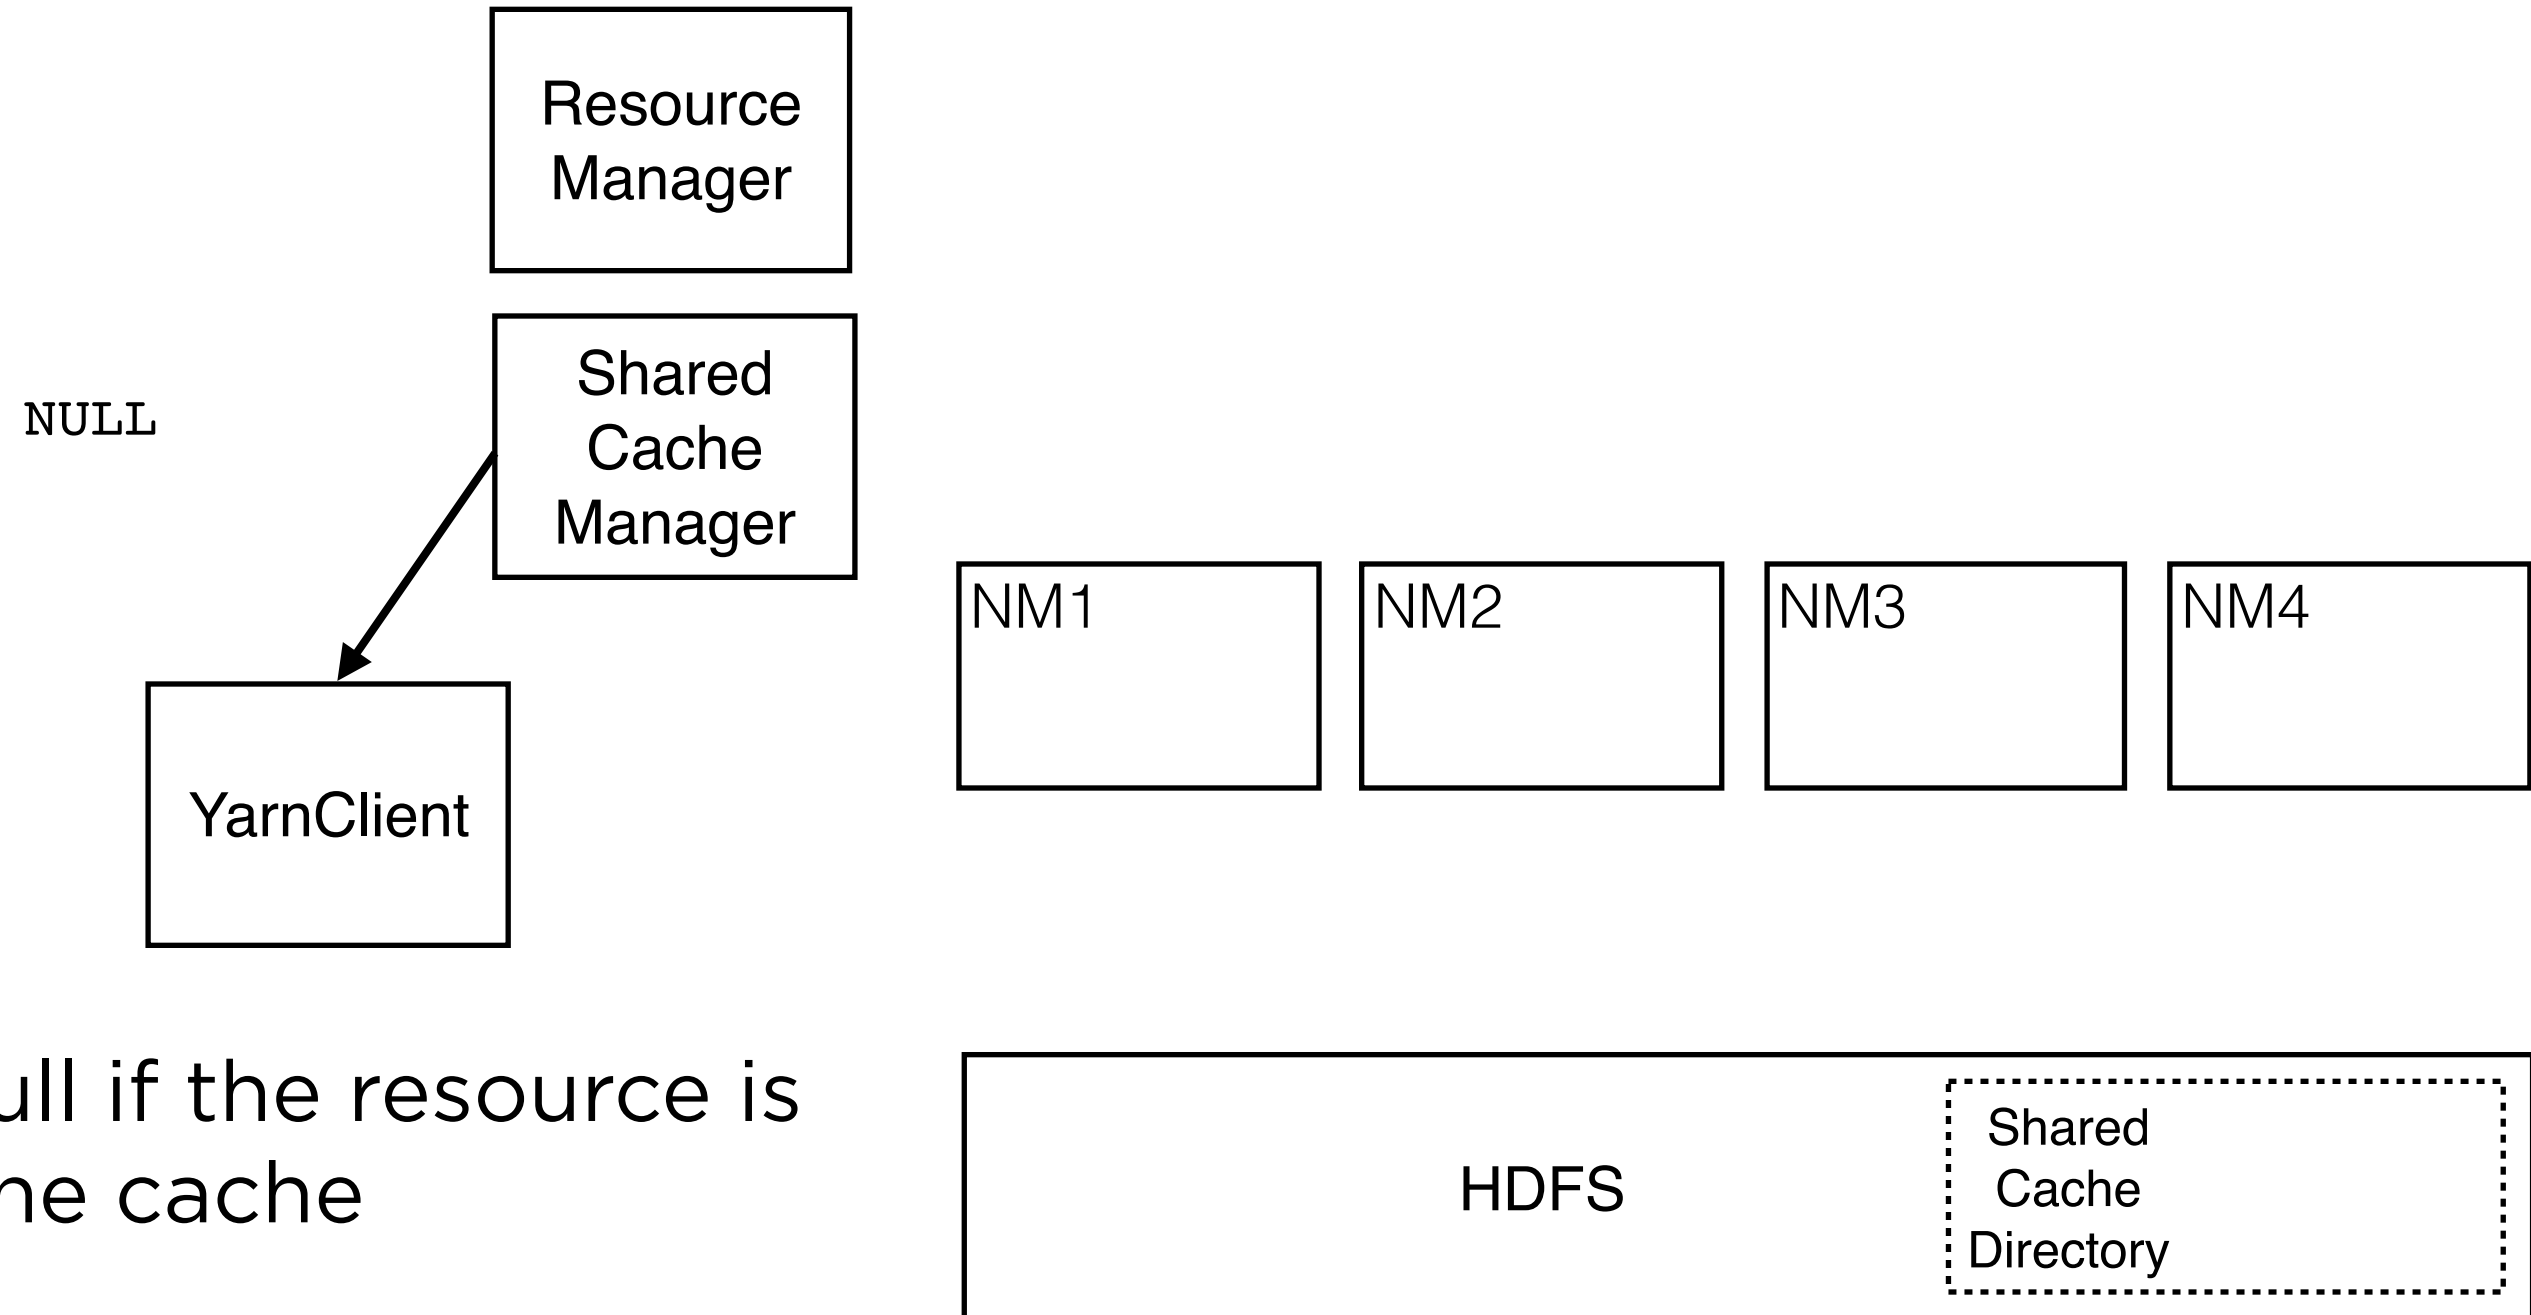
# ADDING TO THE CACHE



3. SCM responds back with the HDFS Path to the resource in the cache



# ADDING TO THE CACHE

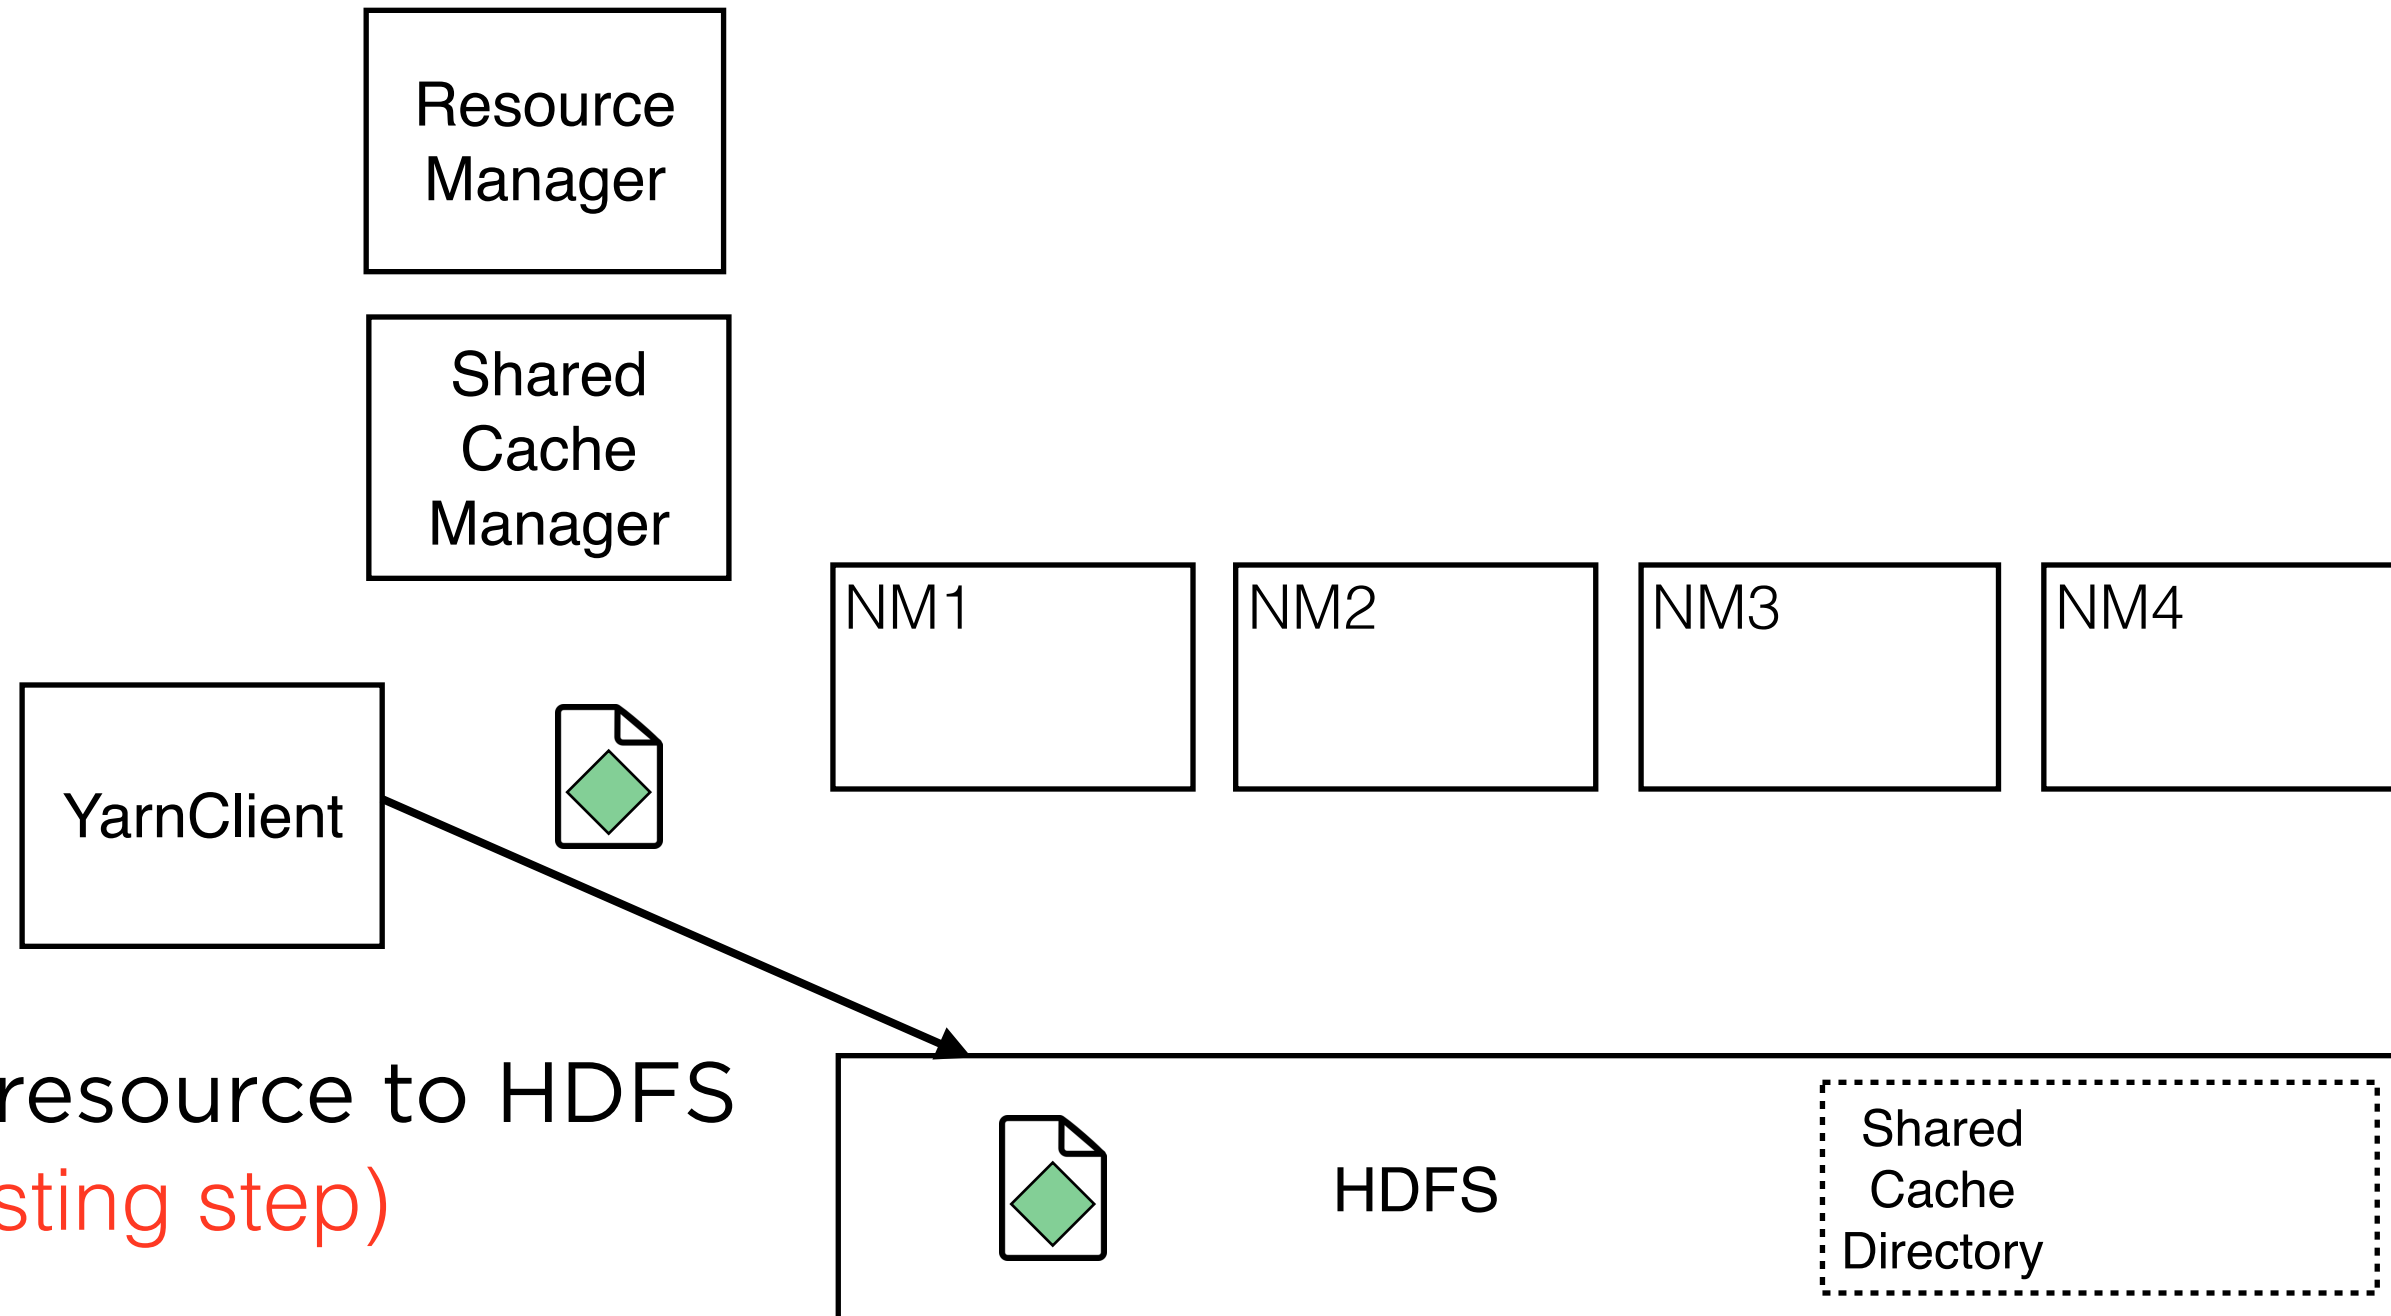


3. Or Null if the resource is not in the cache






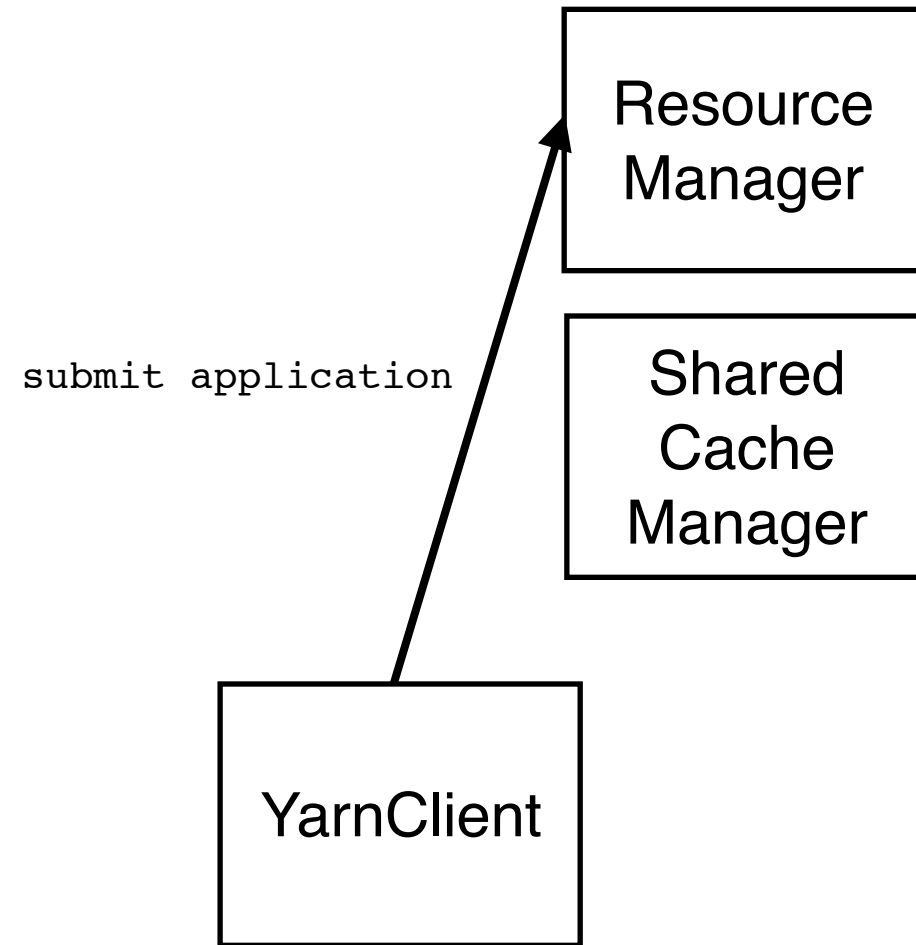
# ADDING TO THE CACHE



4. Upload resource to HDFS  
(Existing step)



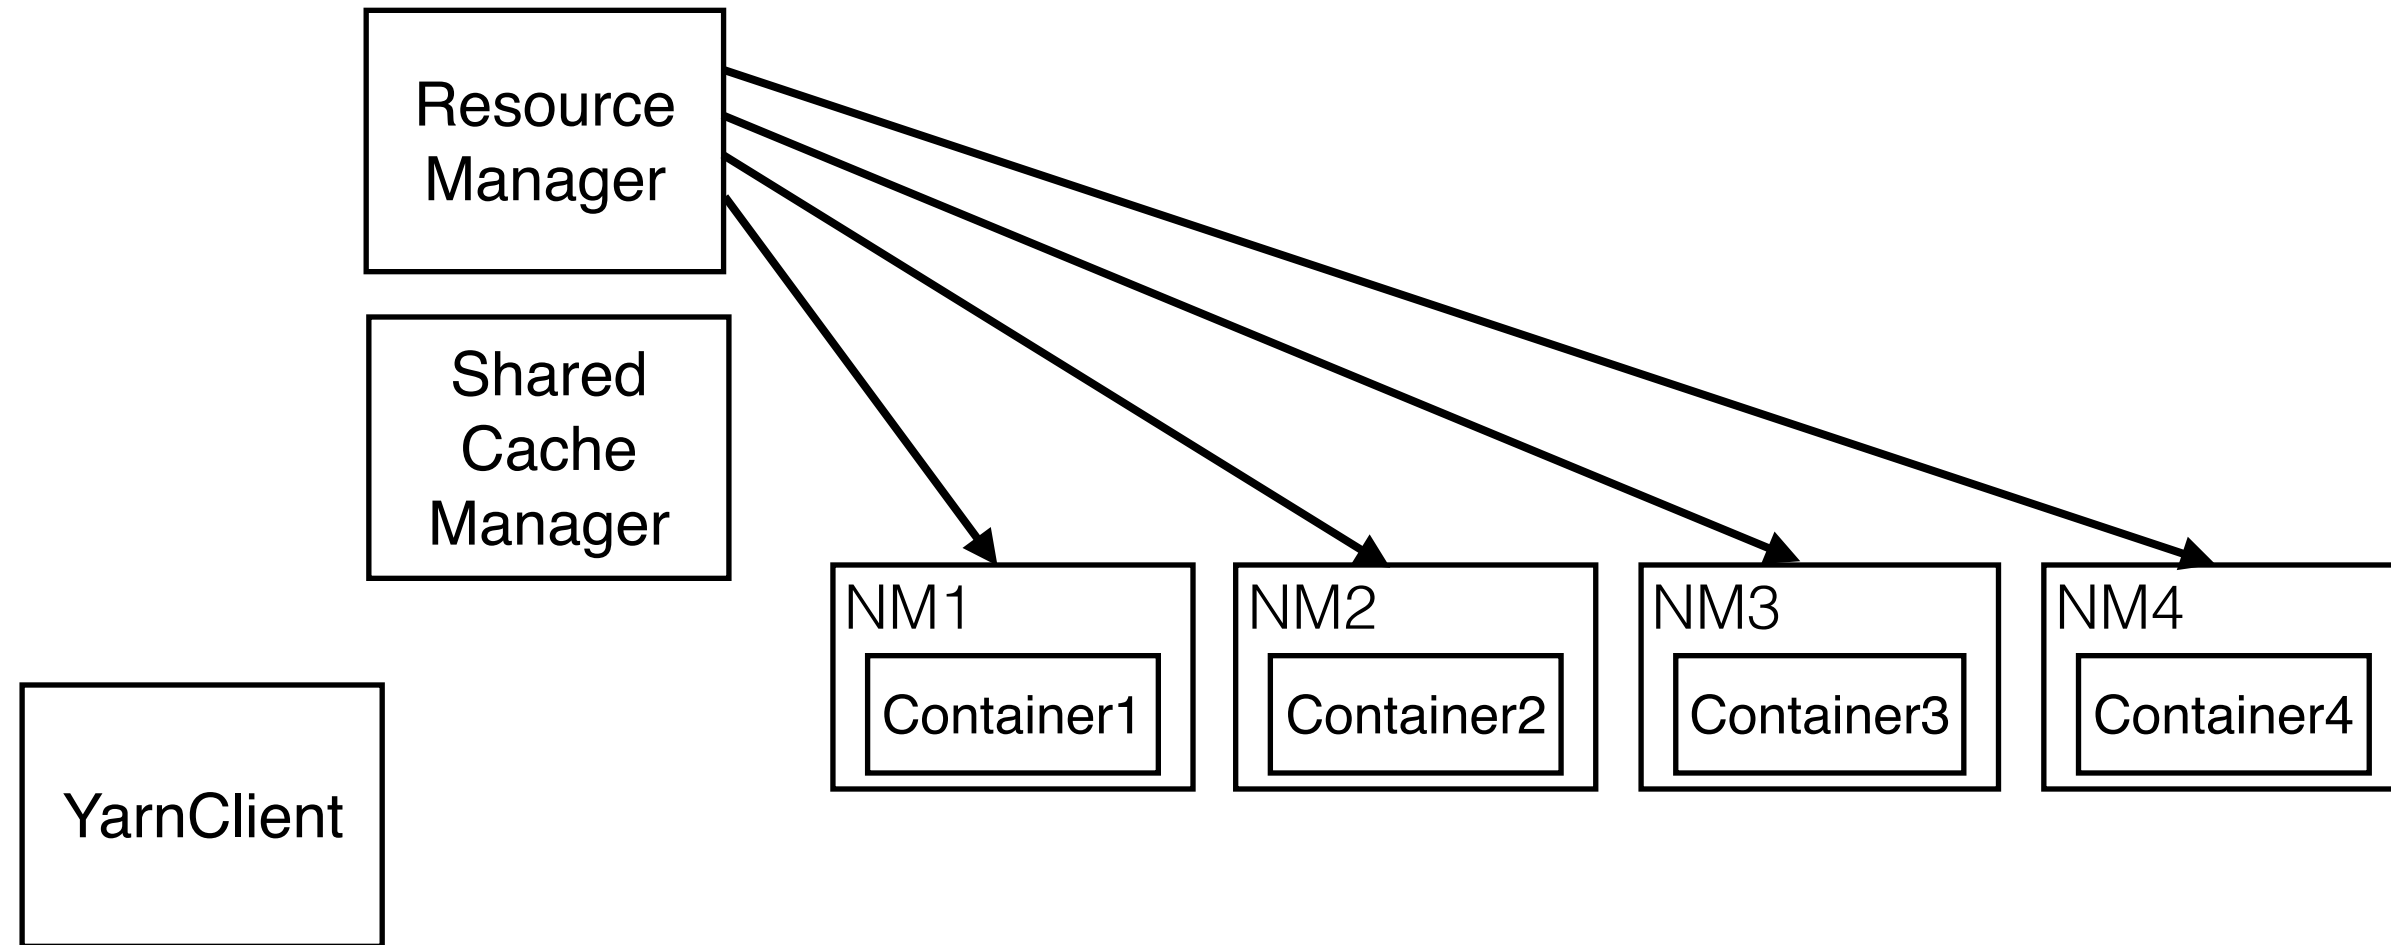
# ADDING TO THE CACHE



5. Set resource to be added to cache (via LocalResource API) and submit app



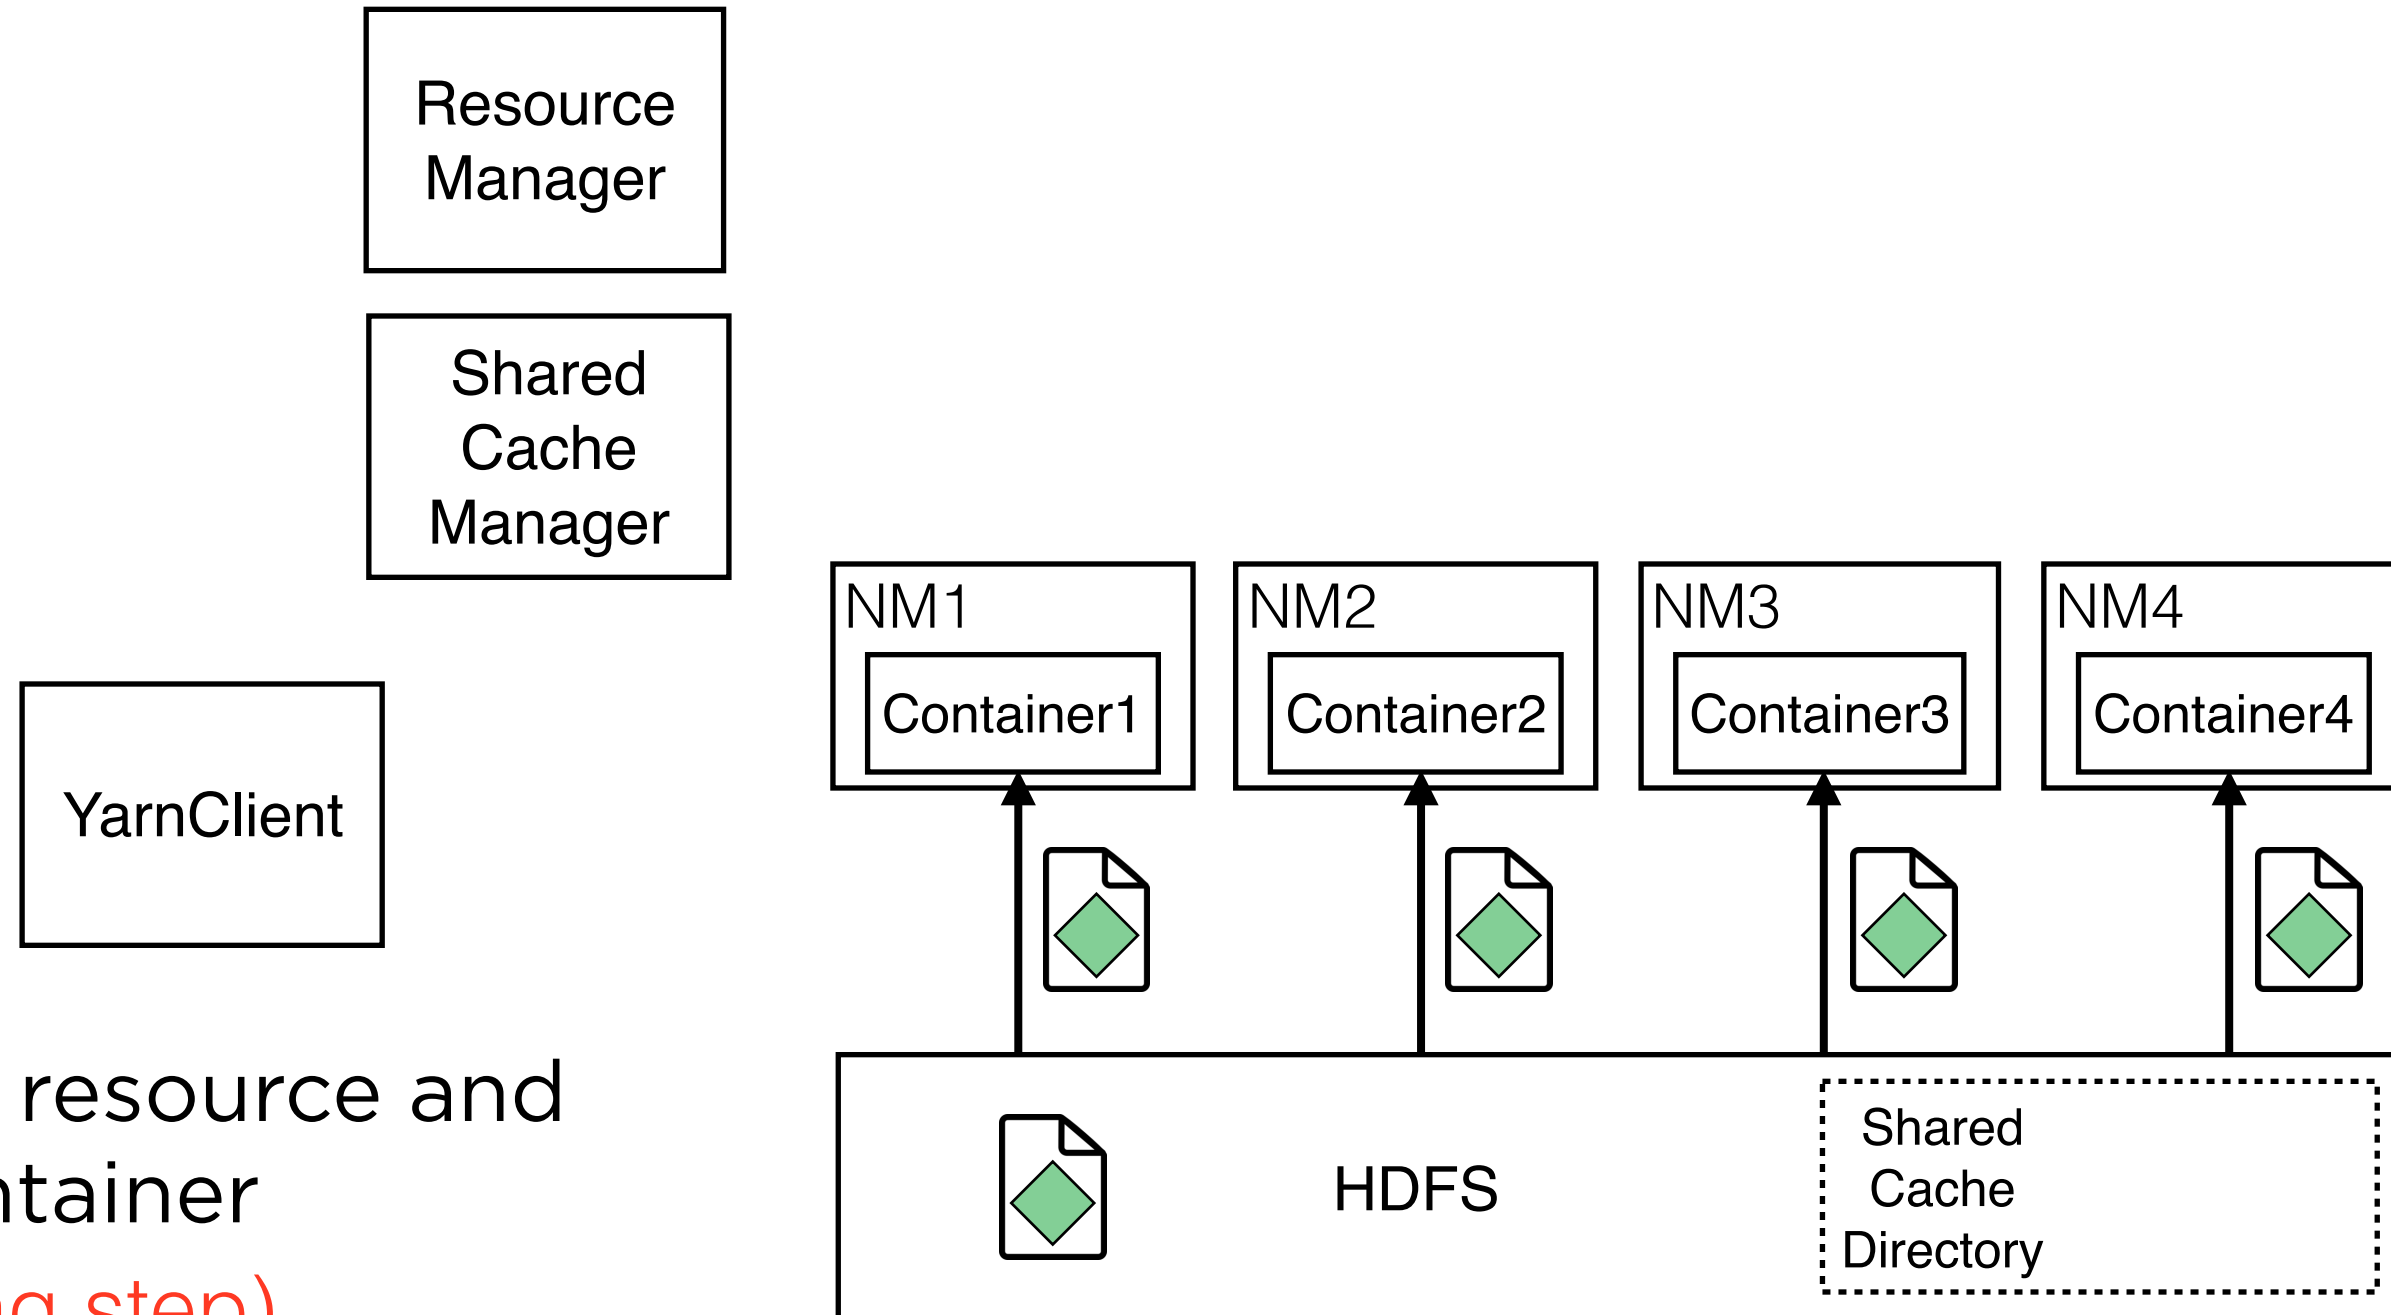
# ADDING TO THE CACHE



6. Schedule containers  
(Existing step)



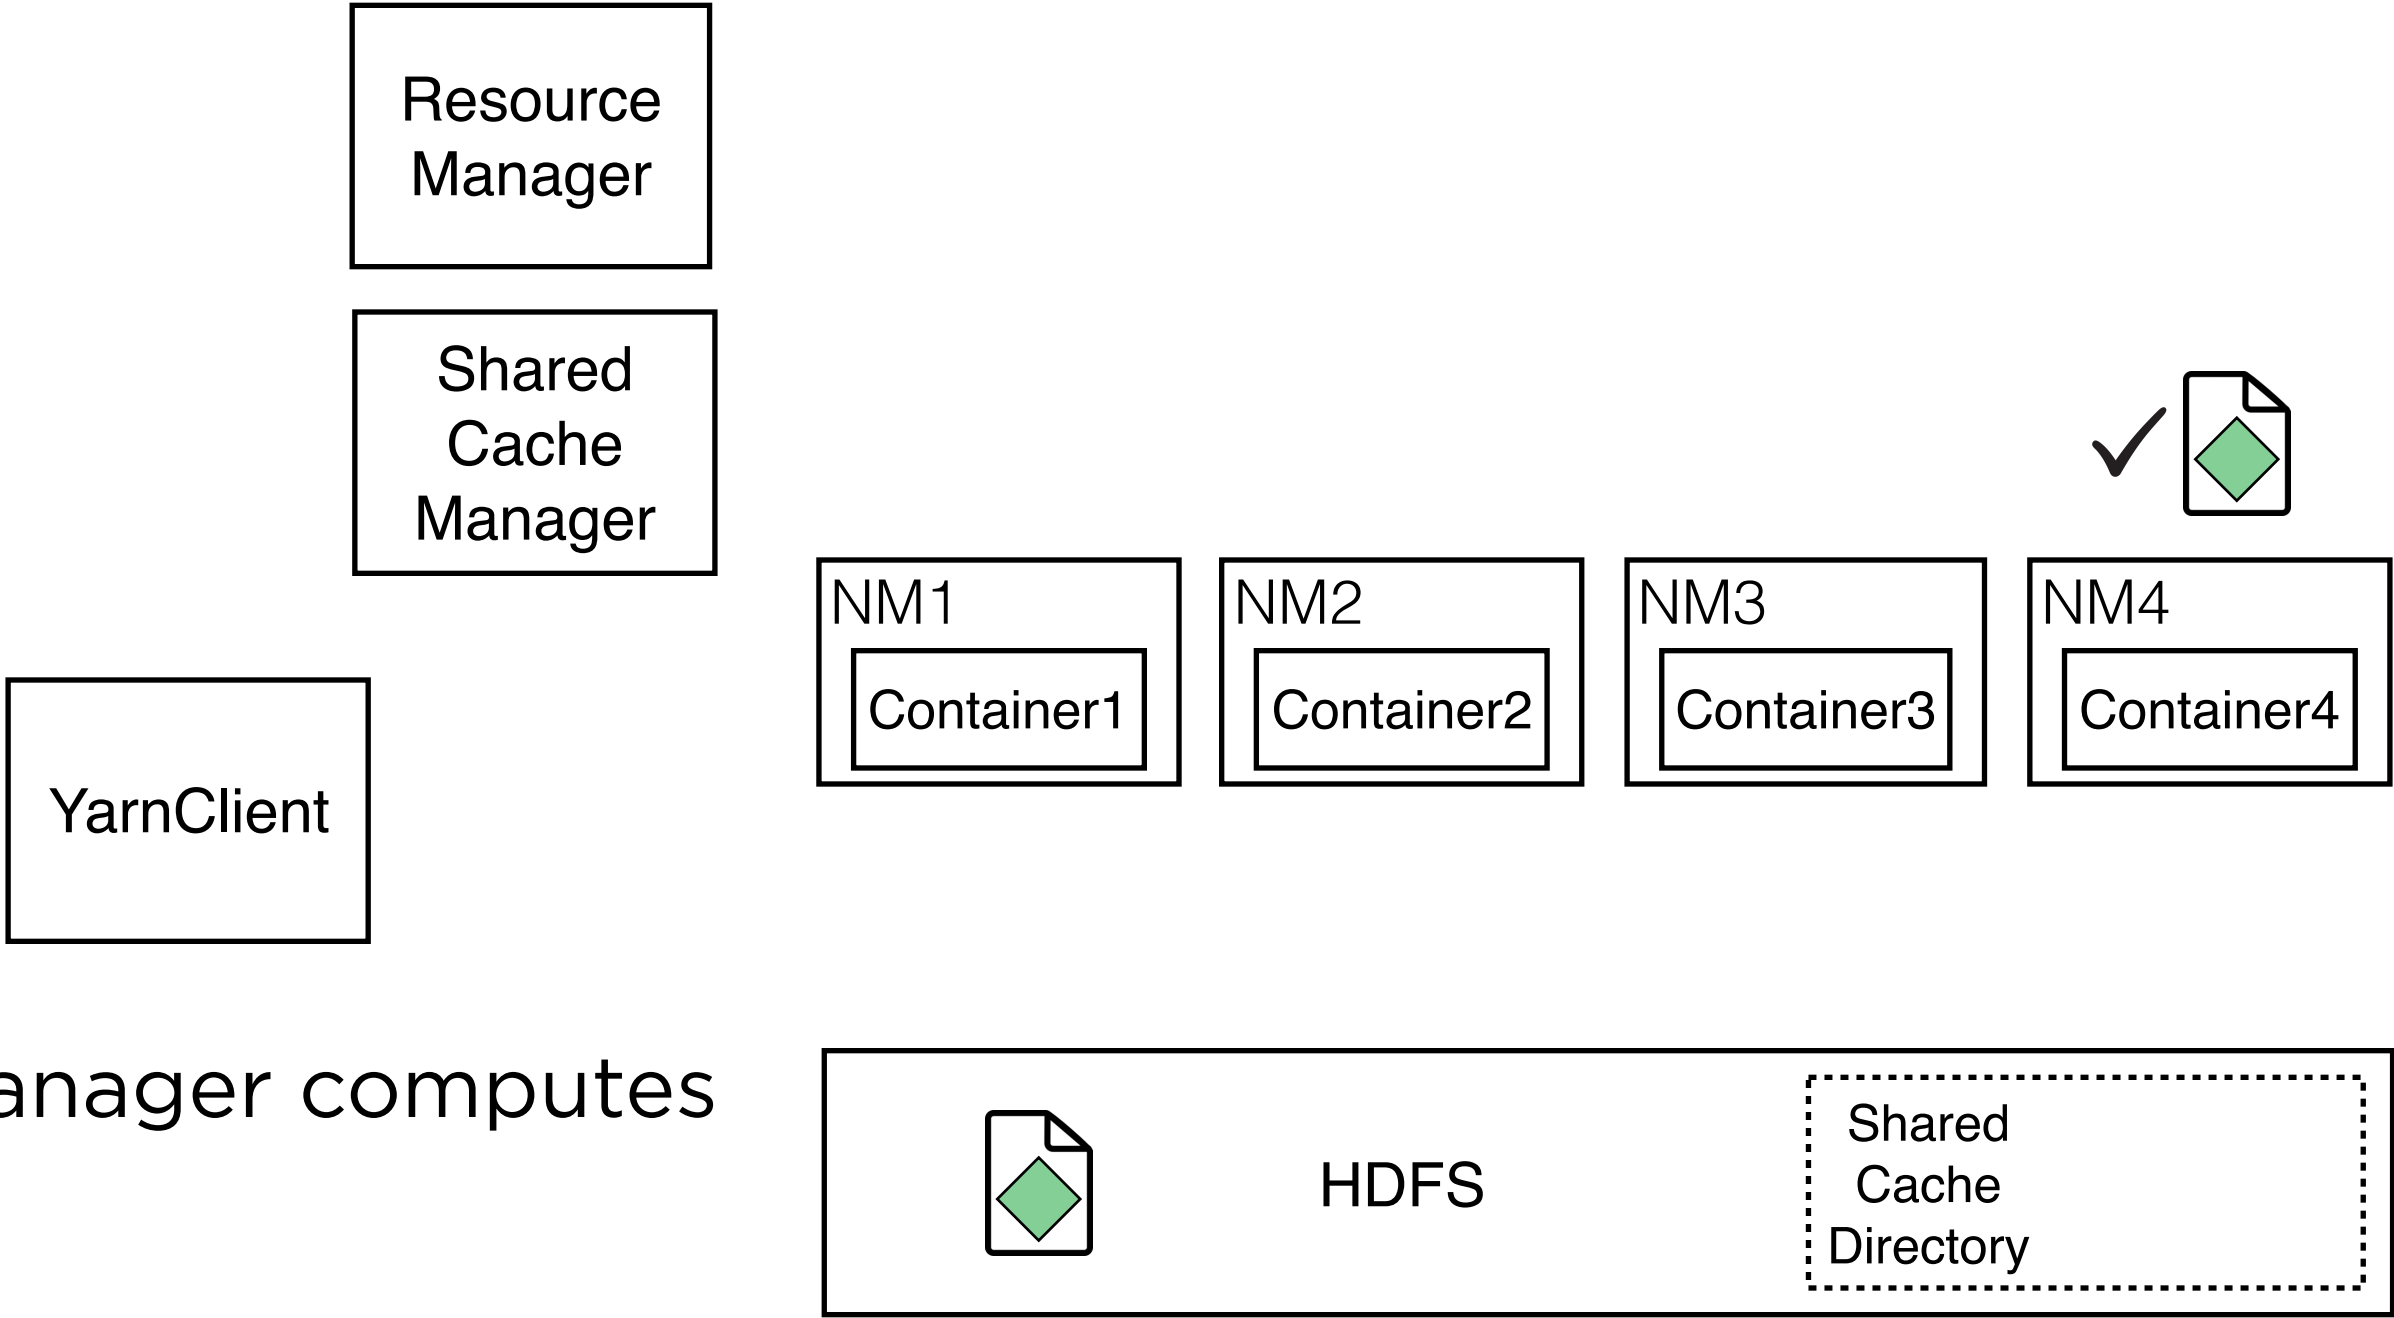
# ADDING TO THE CACHE



7. Localize resource and  
launch container  
(Existing step)



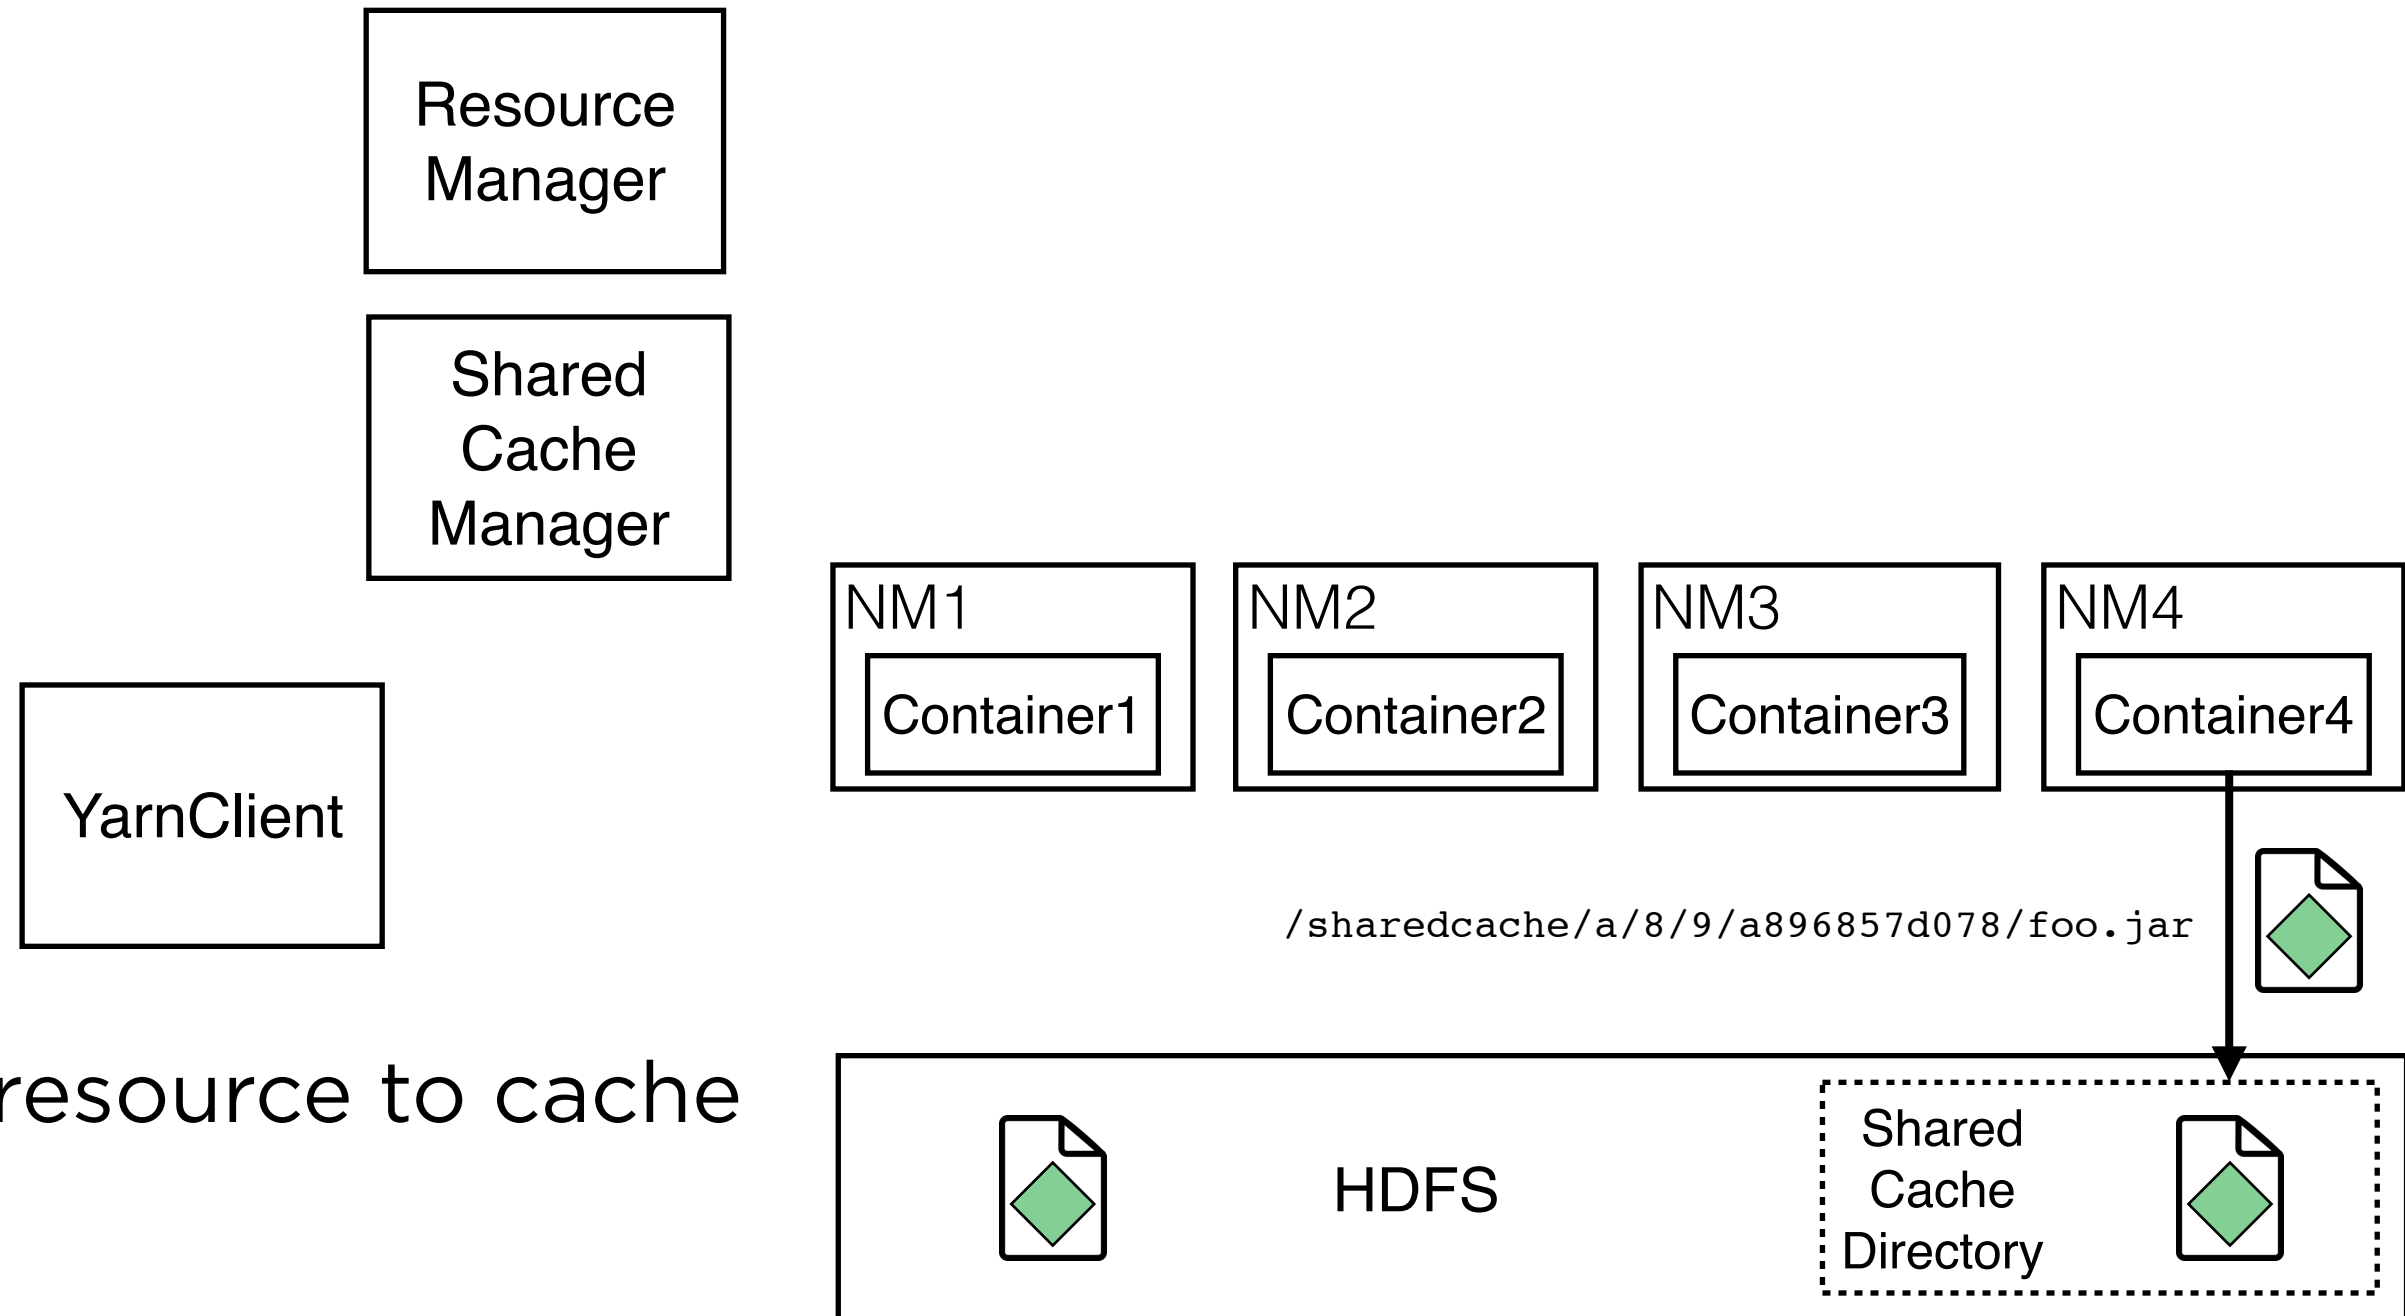
# ADDING TO THE CACHE



8. Node manager computes checksum




# ADDING TO THE CACHE



9. Upload resource to cache



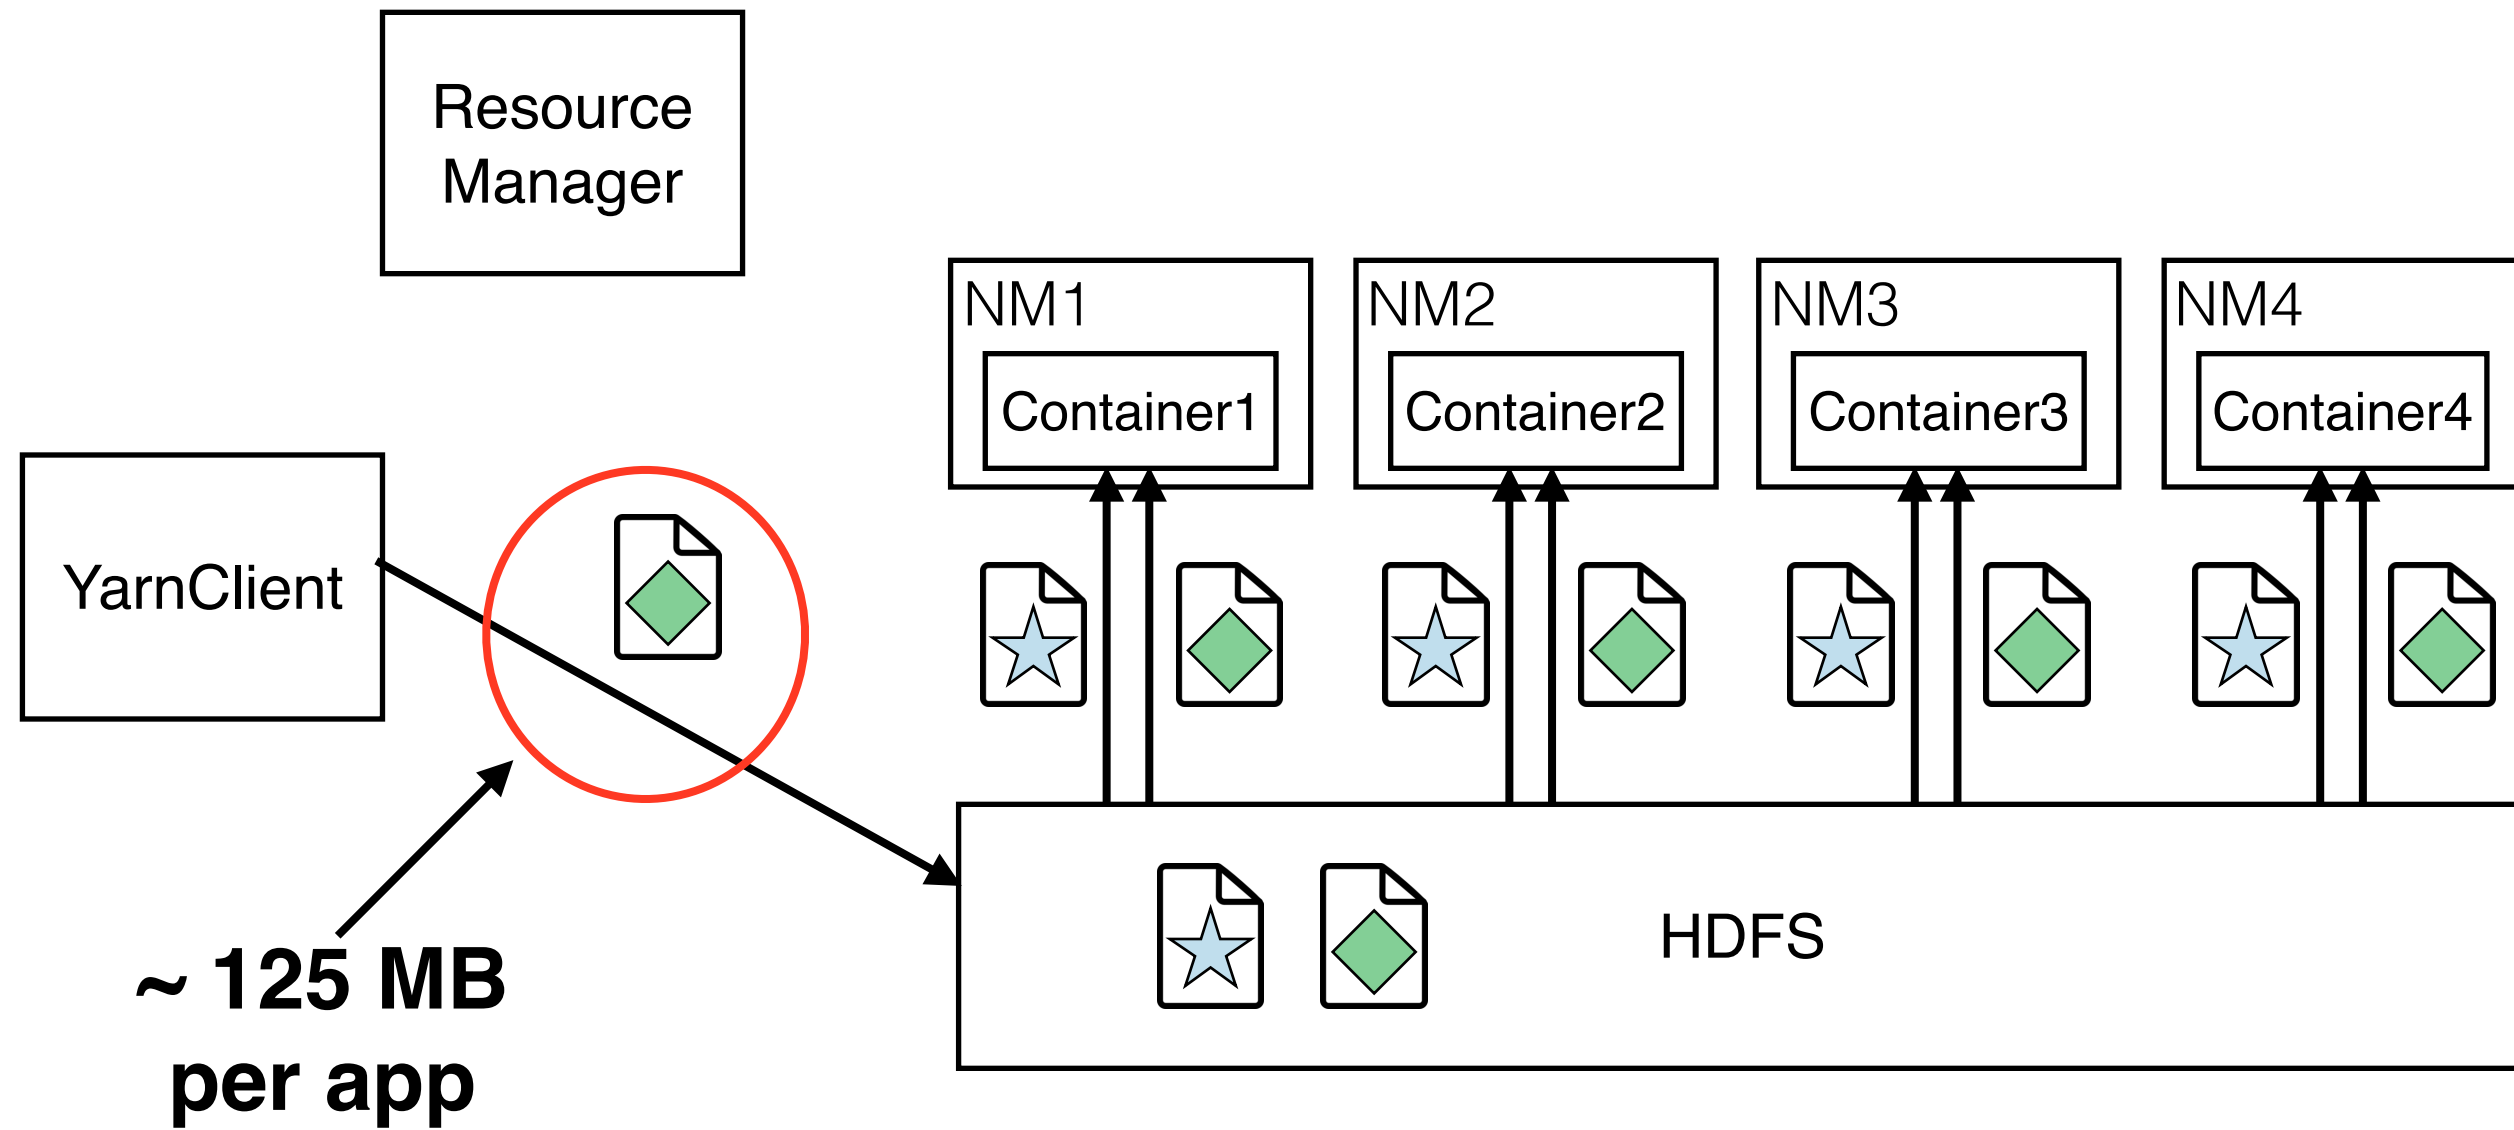
# ADDING TO THE CACHE



10. Notify SCM of new resource

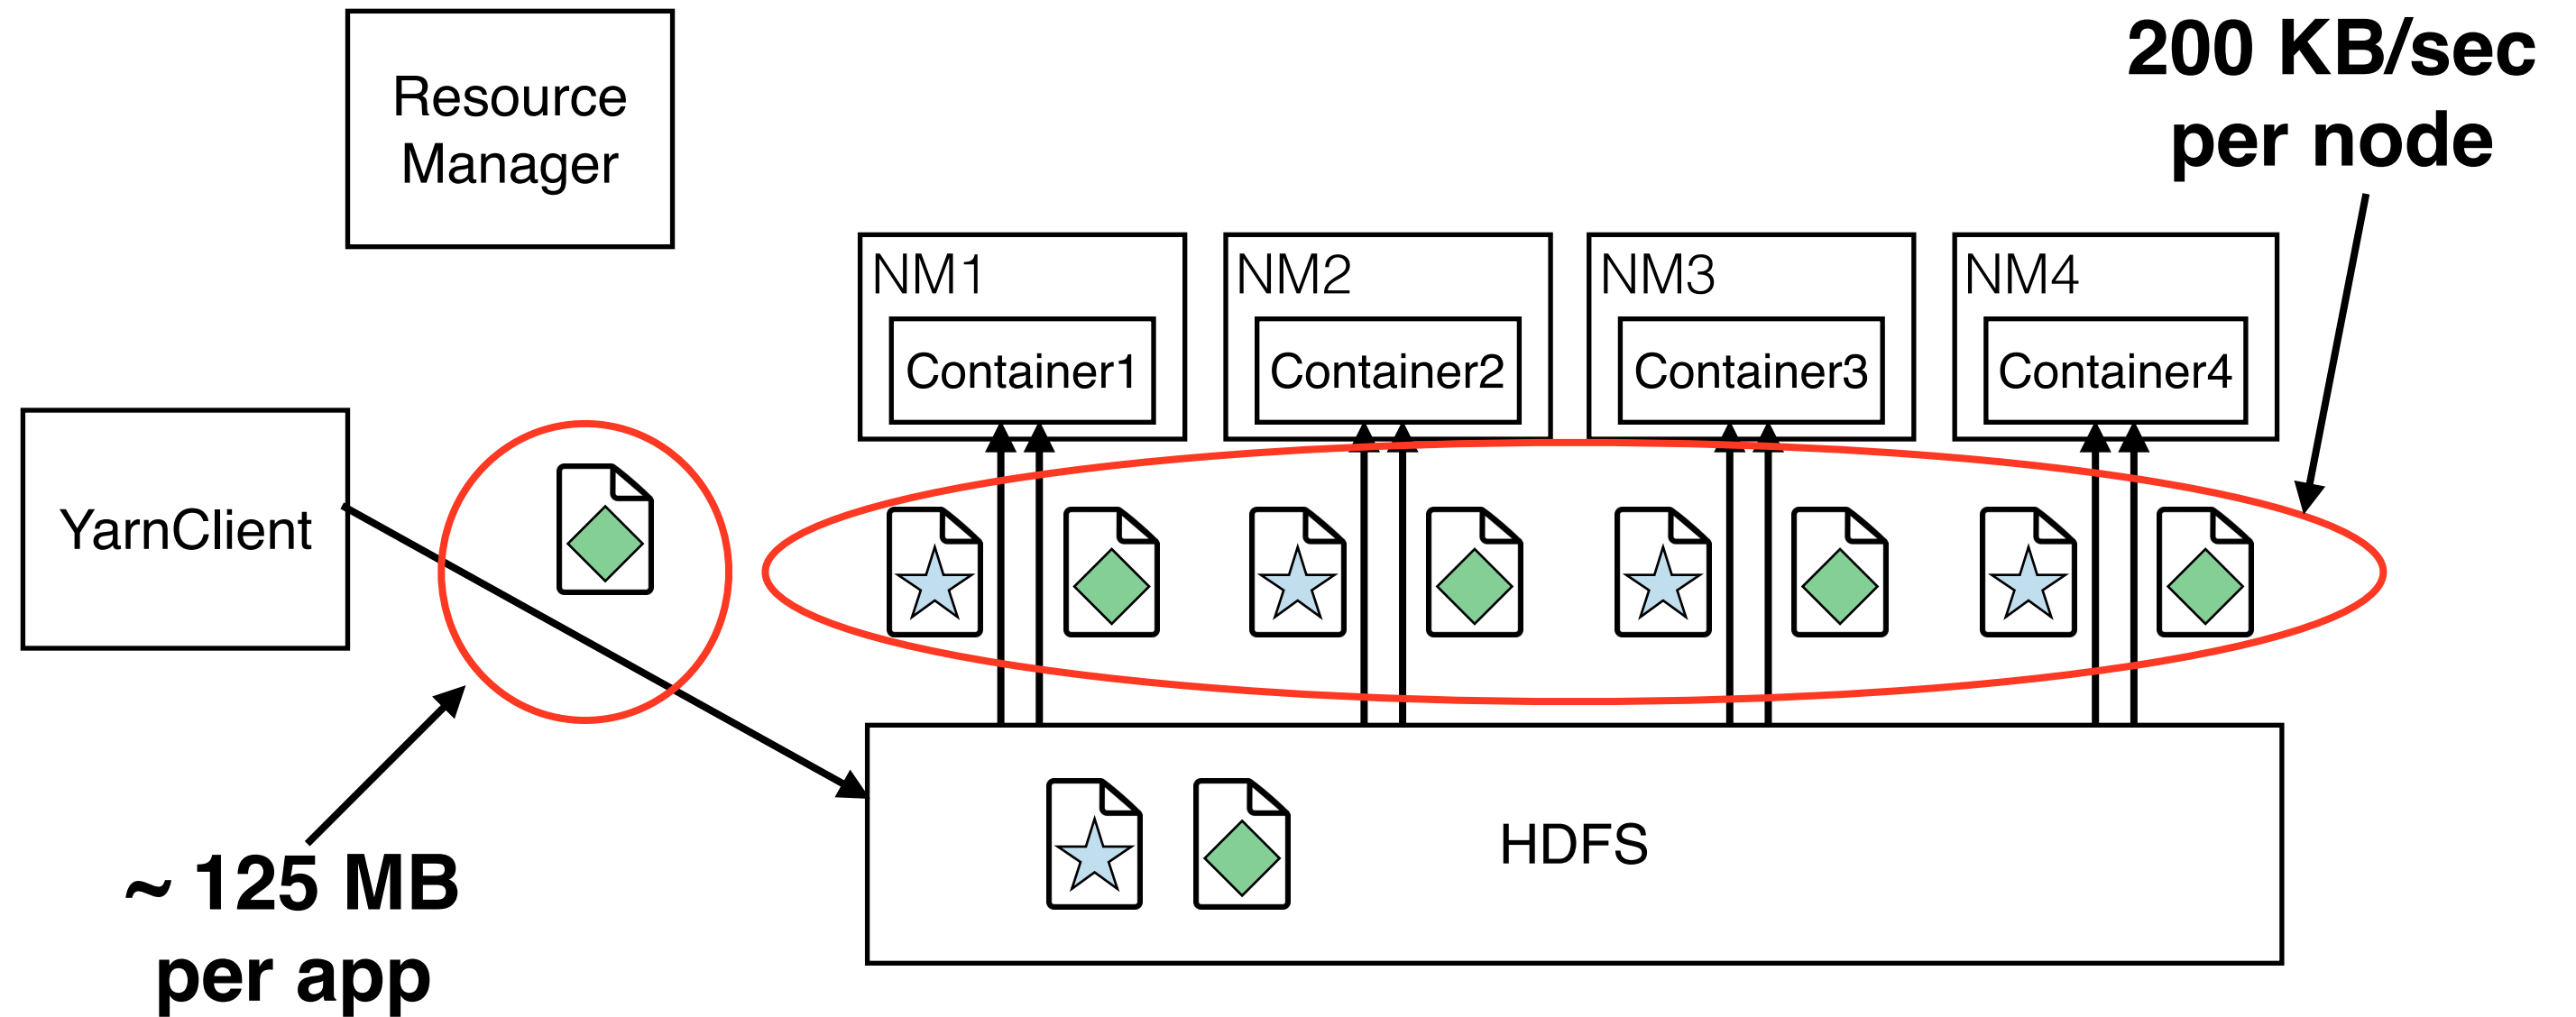


# DOES IT WORK?

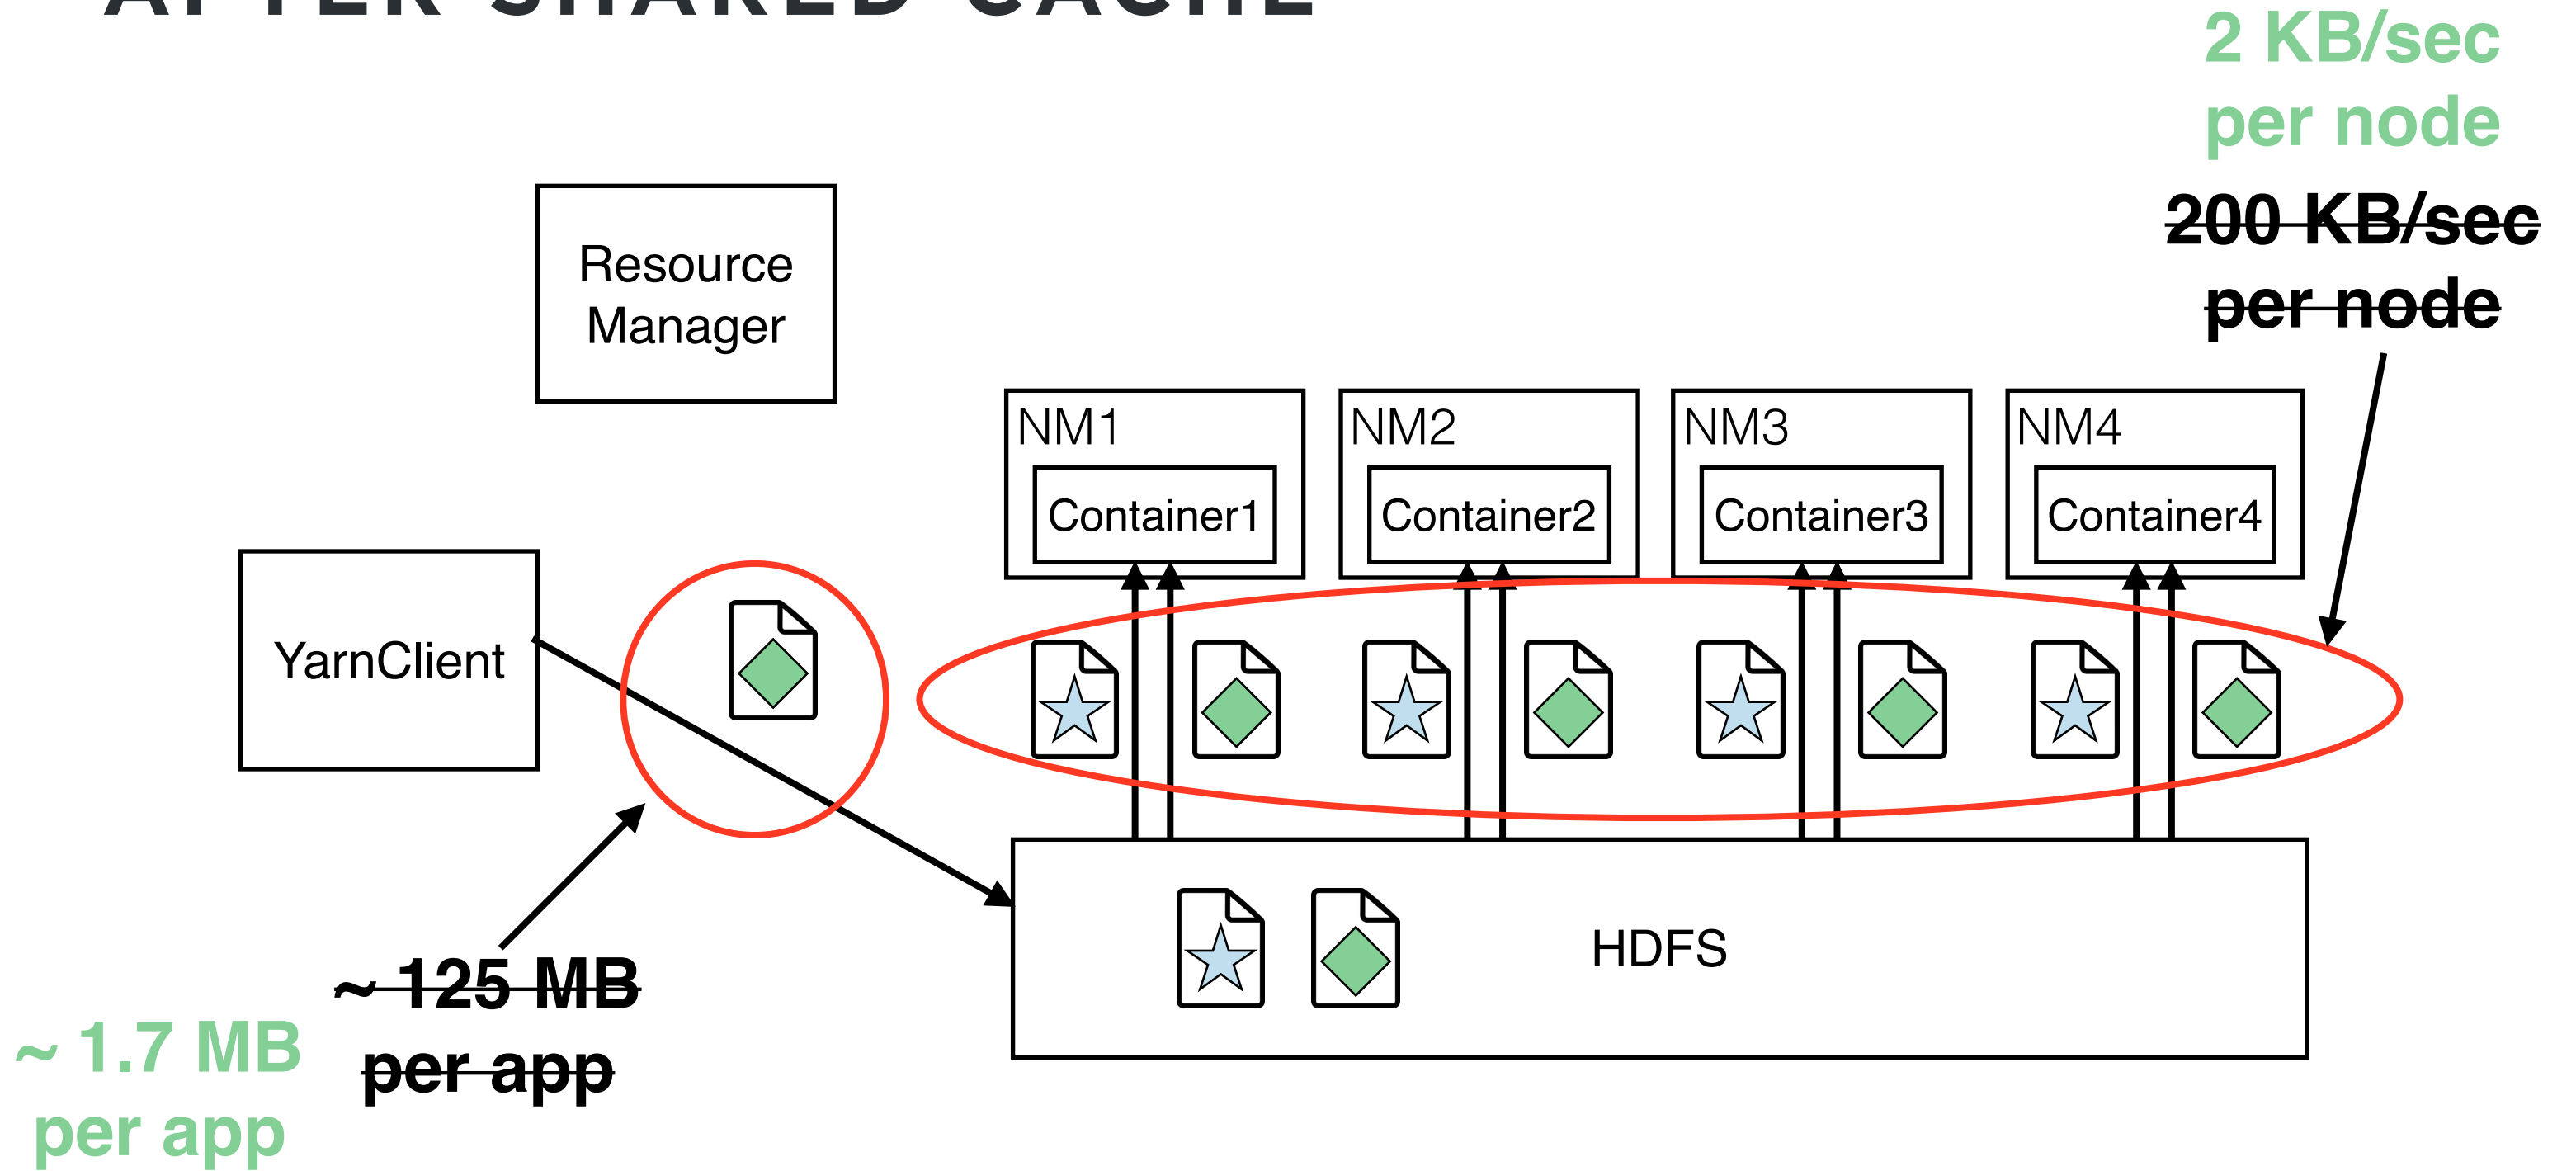




# DOES IT WORK?



# AFTER SHARED CACHE



# BEFORE SHARED CACHE

- Each application uploaded and localized (on average)
  - ~ 12 resources
  - ~ 125 MB total size
- Each container localized (on average)
  - ~ 6 resources
  - ~ 63 MB total size




# AFTER SHARED CACHE

- Each application uploaded and localized (on average)
  - ~ 0.16 resources
  - ~ 1.7 MB total size
- Each container localized (on average)
  - ~ 0.08 resources
  - ~ 840 KB total size
- Saving localization bandwidth by 99%!




# DOES IT WORK?

- Eliminates network bandwidth usage completely for a client that submits jobs constantly



# DOES IT WORK?

- Eliminates network bandwidth usage completely for a client that submits jobs constantly




# ANTI-PATTERNS THAT CAUSE CHURN

- Resource churn
  - Build process may introduce artificial change
  - Fat jars




# FAT JARS






# FAT JARS




# FAT JARS




# RESOURCE CHURN



# RESOURCE CHURN



# RESOURCE CHURN



# ANTI-PATTERNS THAT CAUSE CHURN

- Local cache churn
  - Local cache size too small for working set
  - Competing use of public cache (e.g. pig jar cache, cascading intermediate files)



# ADMIN TIPS

- Avoid “fat jars”
- Make your jars repeatable
- Set the local cache size appropriately
  - `yarn.nodemanager.localizer.cache.target-size-mb=65536`
- Increase public localizer thread pool size
  - `yarn.nodemanager.localizer.fetch.thread-count=12`
- Adjust the cleaner frequency to your usage pattern
  - `yarn.sharedcache.cleaner.period.minutes` (default: 1 day)



# DEV TIPS

- YARN Developer
  - Invoke `use` API to claim resources
  - Use the `LocalResource` API to add resources
- MapReduce Developer
  - Set `mapreduce.job.sharedcache.mode`
    - `jobjar,libjar,files,archives`





# ACKNOWLEDGEMENTS

- Code: Chris Trezzo, Sangjin Lee, Ming Ma
- Shepherd: Karthik Kambatla
- Design Review
  - Karthik Kambatla
  - Vinod Kumar Vavilapalli
  - Jason Lowe
  - Joep Rottinghuis



# Q&A - THANKS!

- Chris Trezzo - @ctrezzo

